

A SHORT INTRODUCTION TO

SnailTrail: Generalizing Critical Paths for Online Analysis of Distributed Dataflows*

strymon.systems.ethz.ch

*Moritz Hoffmann, Andrea Lattuada, John Liagouris, Vasiliki Kalavri,
Desislava Dimitrova, Sebastian Wicki, Zaheer Chothia, Timothy Roscoe*

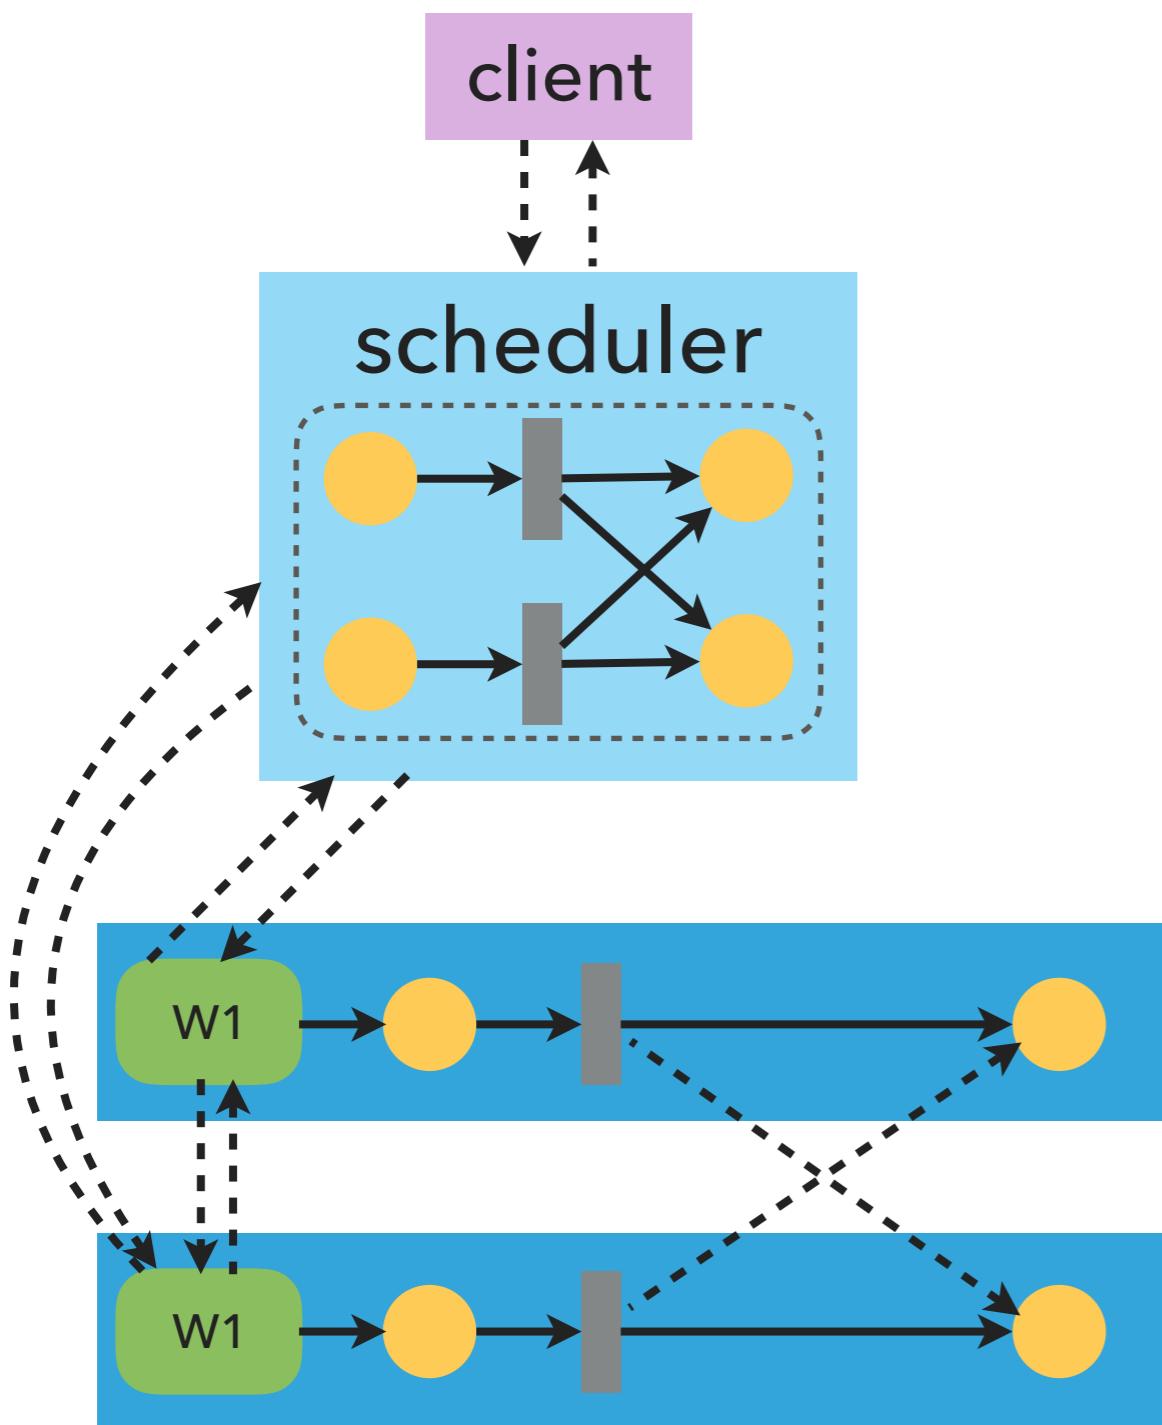
Systems Group, Department of Computer Science, ETH Zürich

firstname.lastname@inf.ethz.ch

Malte Sandstede (samalt@ethz.ch)
Slides (partly) stolen from Vasia Kalavri



ANALYZING DISTRIBUTED DATAFLOWS



Hard to troubleshoot...

- ▶ long-running, dynamic workloads
- ▶ many tasks, activities, operators, dependencies
- ▶ bottleneck causes are usually not isolated but span multiple processes

ANALYZING DISTRIBUTED DATAFLOWS

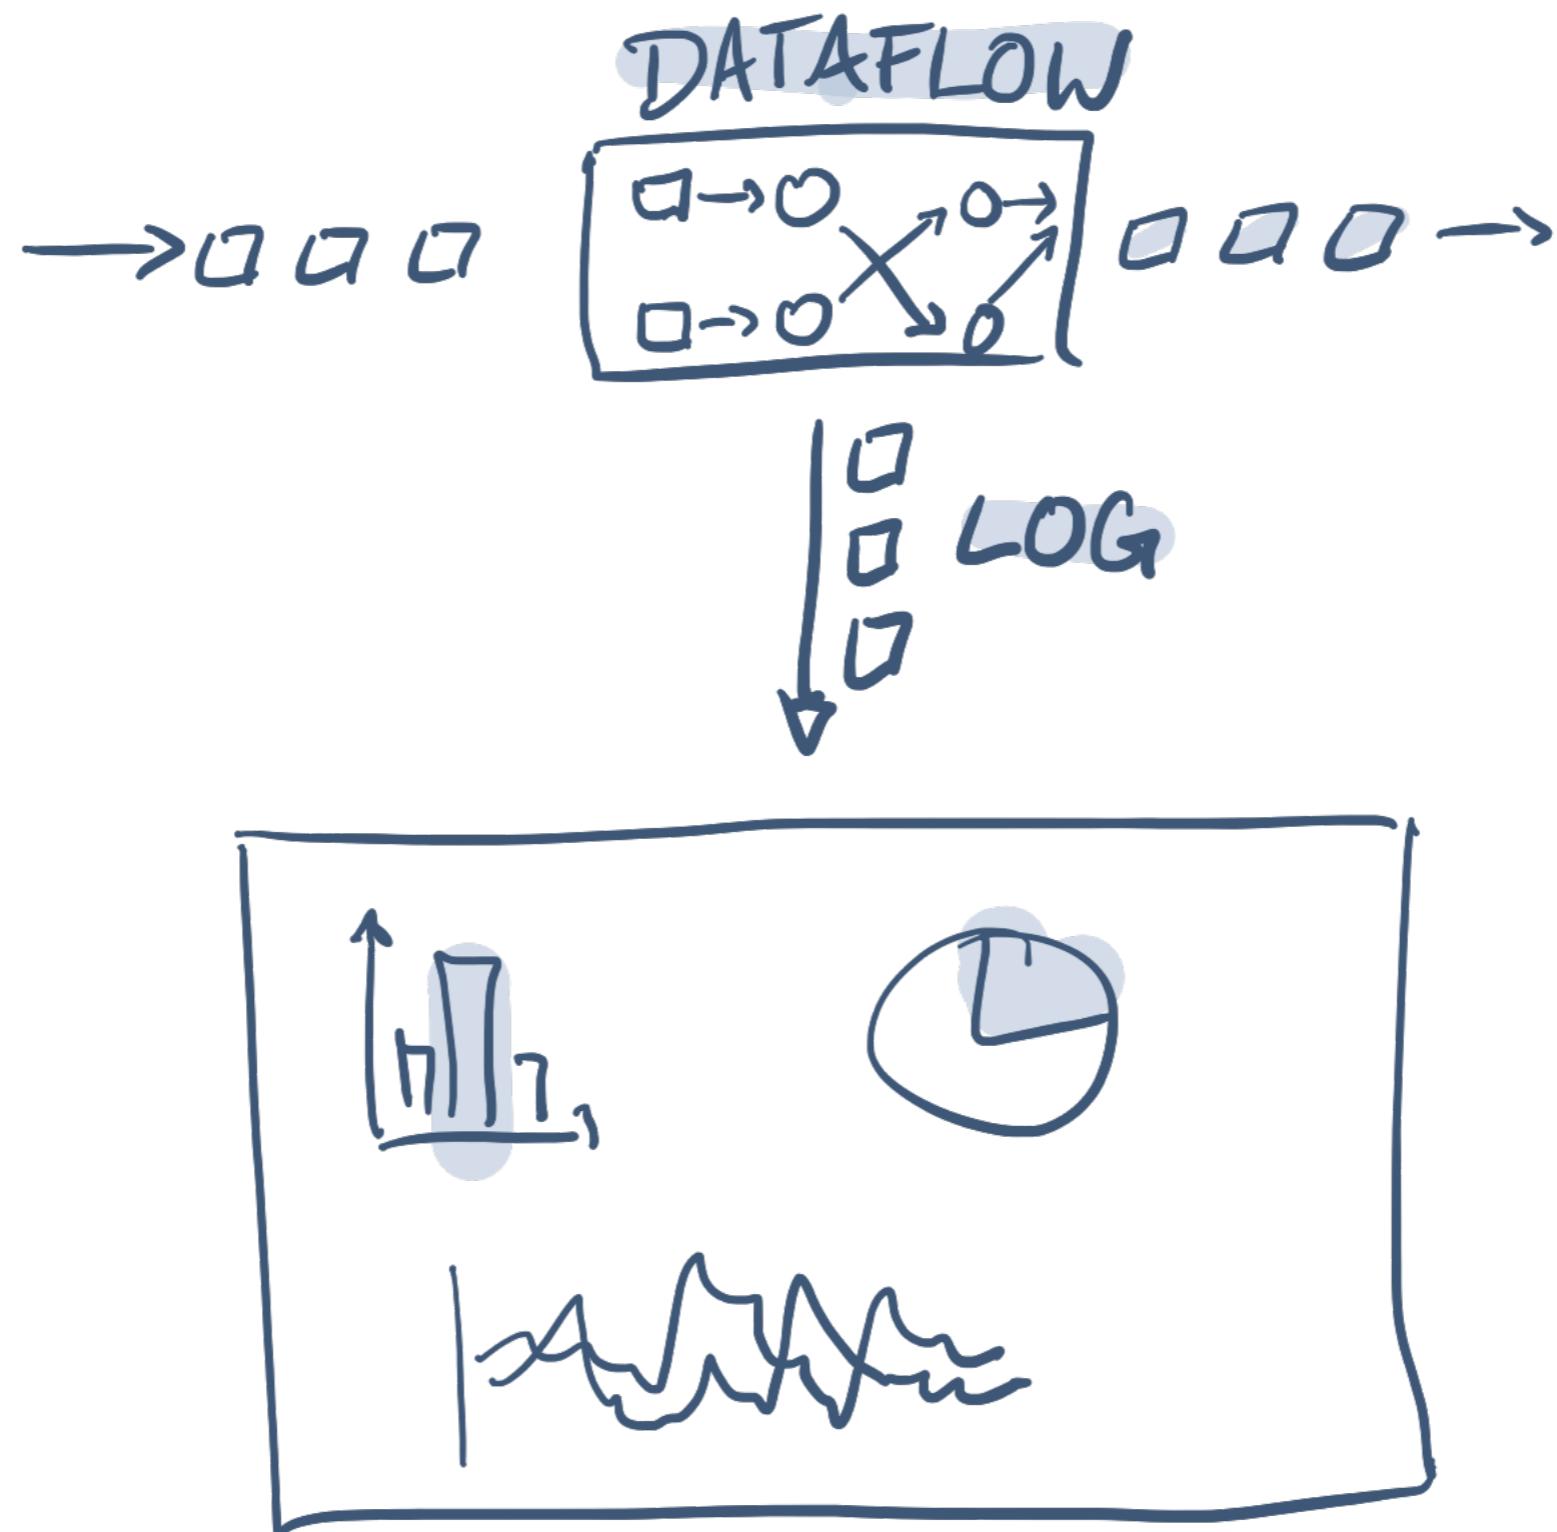
Hard to troubleshoot...

...but very useful (duh!)

- ▶ long-running, dynamic workloads
- ▶ many tasks, activities, operators, dependencies
- ▶ bottleneck causes are usually not isolated but span multiple processes
- ▶ Provenance & Explainability
- ▶ Performance Tuning
- ▶ Debugging
- ▶ SLAs

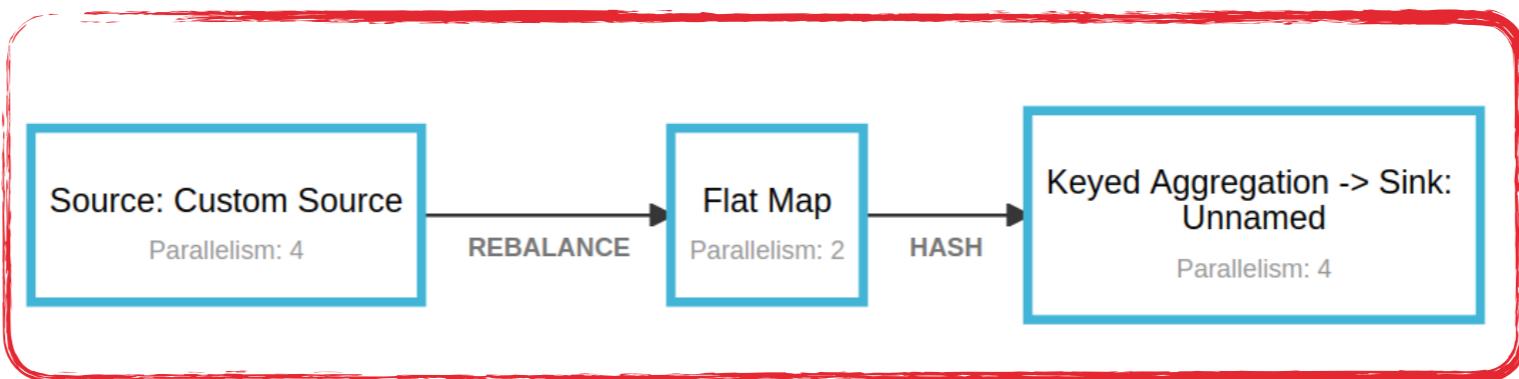
CONVENTIONAL PERFORMANCE METRICS

CONVENTIONAL PERFORMANCE METRICS



CONVENTIONAL PERFORMANCE METRICS

Dataflow graph



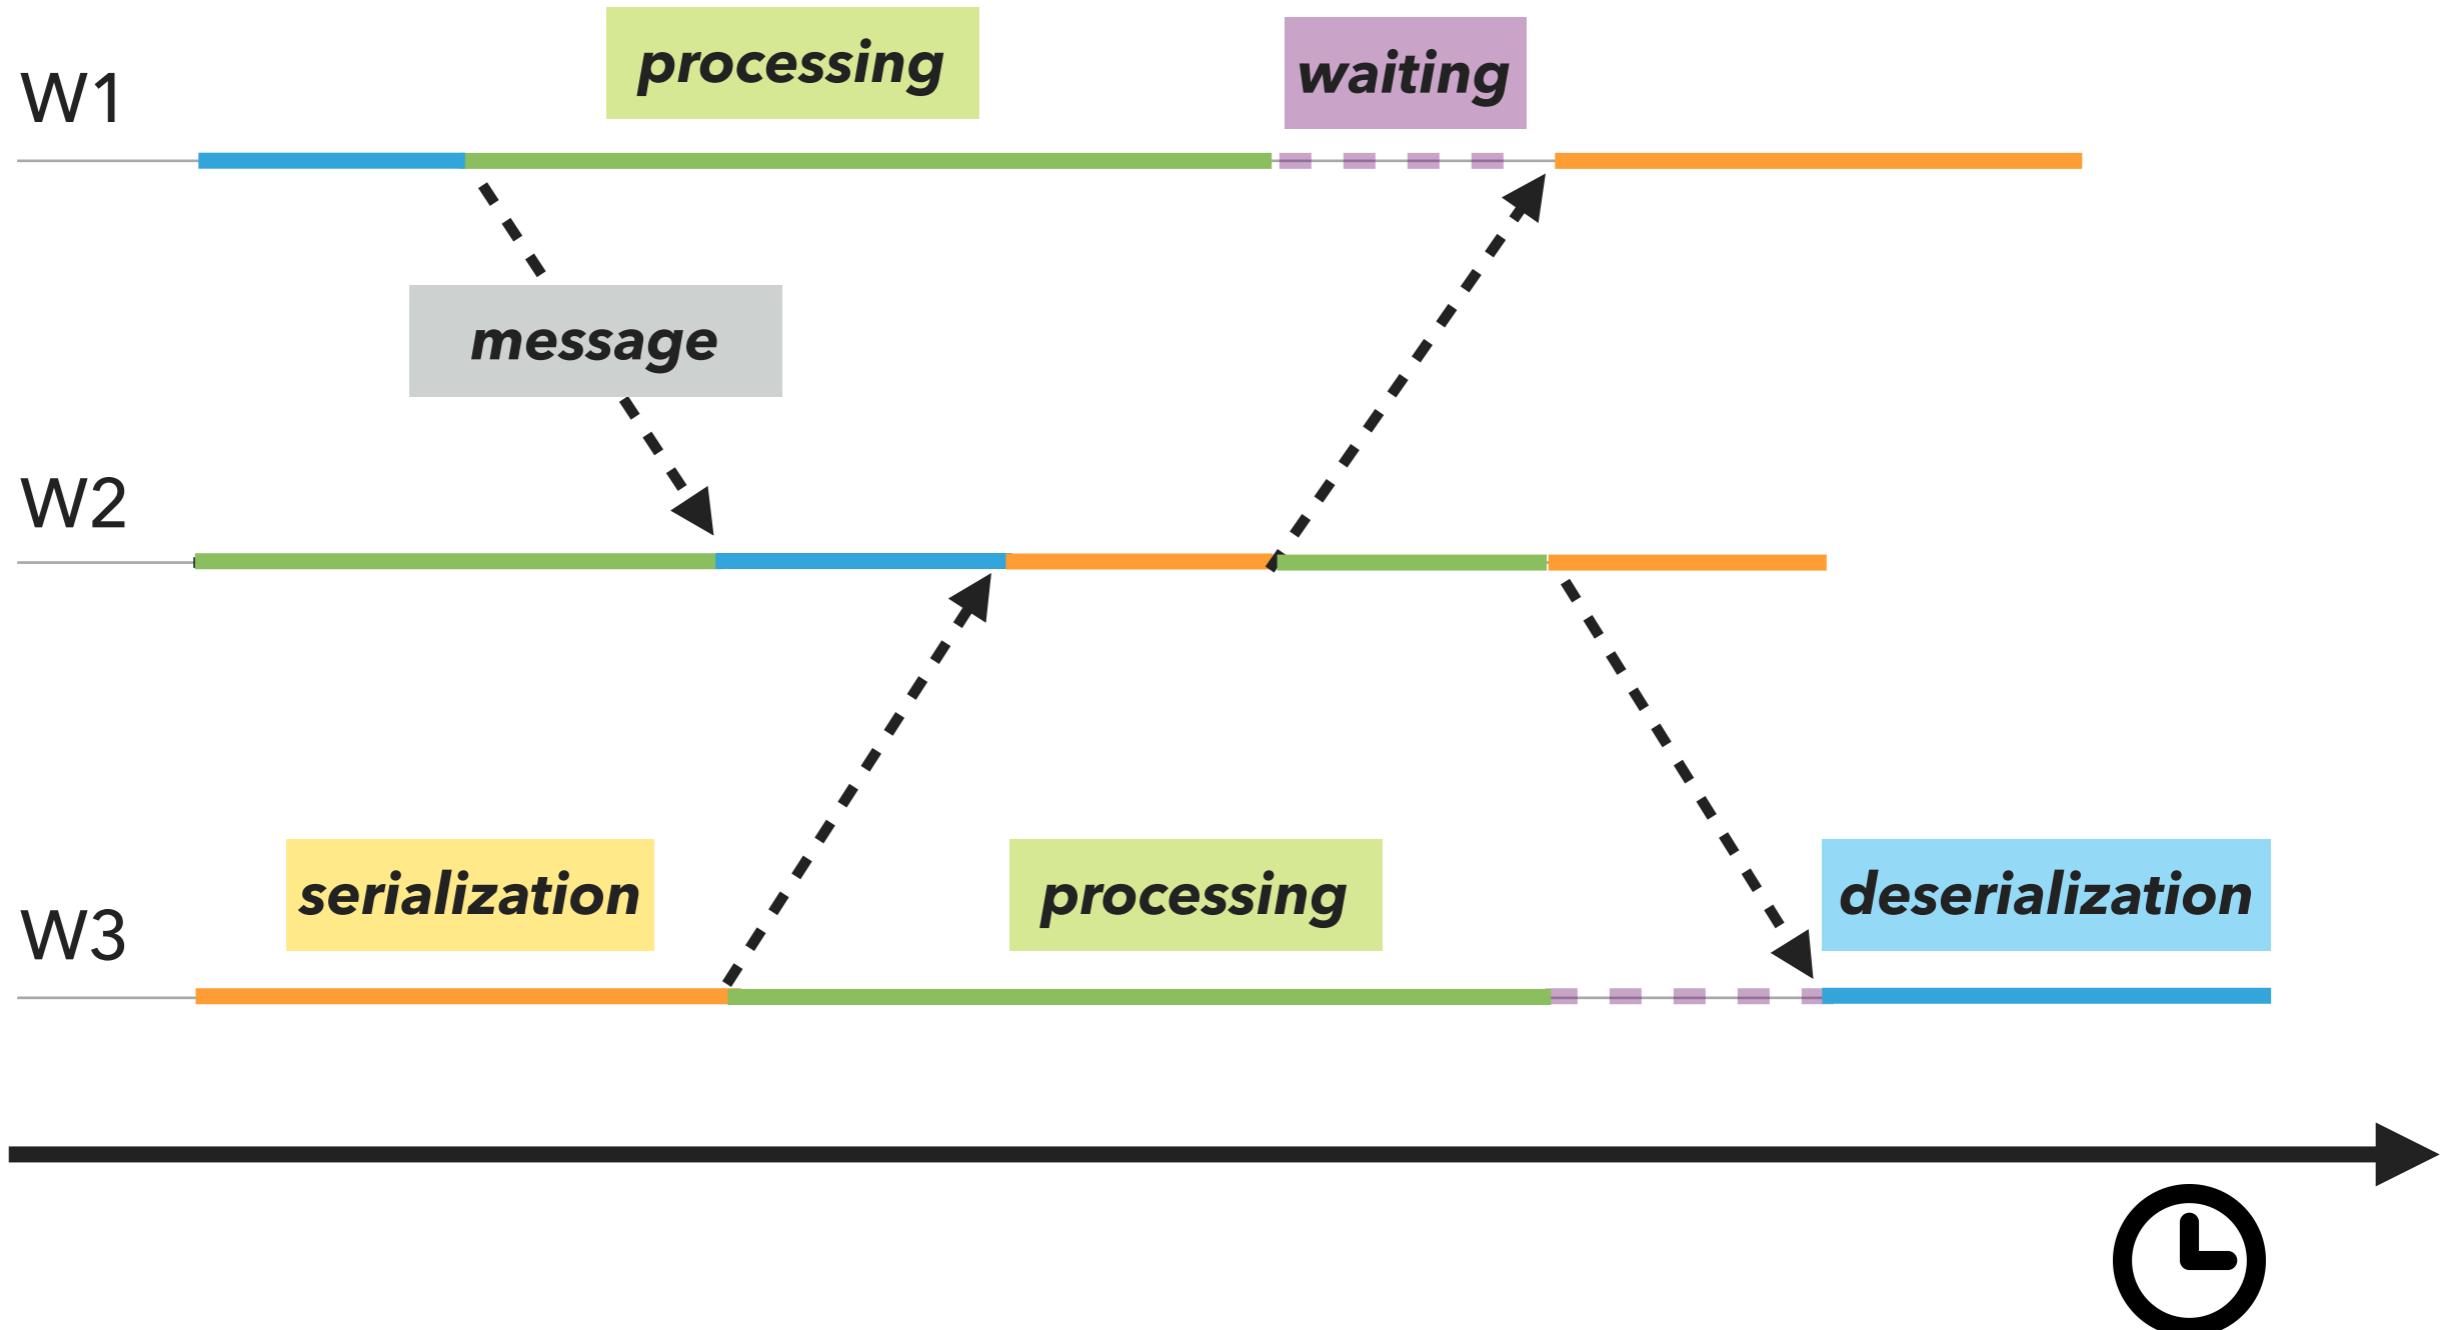
Custom
aggregate metrics

Subtasks	TaskManagers	Metrics	Accumulators	Checkpoints	Back Pressure							
Start Time	End Time	Duration	Name									
2017-10-19, 17:45:43	2017-10-19, 17:47:42	1m 58s	Source: Custom Source									
2017-10-19, 17:45:43	2017-10-19, 17:47:42	1m 58s	Flat Map									
Start Time	End Time	Duration	Bytes received	Records received								
2017-10-19, 17:45:43		1m 58s	1.26 GB	6,401,851								
2017-10-19, 17:45:43		1m 58s	1.27 GB	6,442,643								
2017-10-19, 17:45:43	2017-10-19, 17:47:42	1m 58s	Keyed Aggregation -> Sink: Unnamed									
Bytes received	Records received	Bytes sent	Records sent	Parallelism	Tasks							
0 B	0	2.56 GB	12,942,429	4	<table border="1"><tr><td>0</td><td>0</td><td>4</td><td>0</td><td>0</td><td>0</td><td>0</td></tr></table> RUNNING	0	0	4	0	0	0	0
0	0	4	0	0	0	0						
2.54 GB	12,844,494	4.59 GB	254,656,192	2	<table border="1"><tr><td>0</td><td>0</td><td>2</td><td>0</td><td>0</td><td>0</td><td>0</td></tr></table> RUNNING	0	0	2	0	0	0	0
0	0	2	0	0	0	0						
Bytes sent	Records sent	Attempt	Host	Status								
2.29 GB	126,924,569	1	kalamari:43402	RUNNING								
2.30 GB	127,731,623	1	kalamari:43402	RUNNING								
4.59 GB	254,646,032	0 B	0	4	<table border="1"><tr><td>0</td><td>0</td><td>4</td><td>0</td><td>0</td><td>0</td><td>0</td></tr></table> RUNNING	0	0	4	0	0	0	0
0	0	4	0	0	0	0						

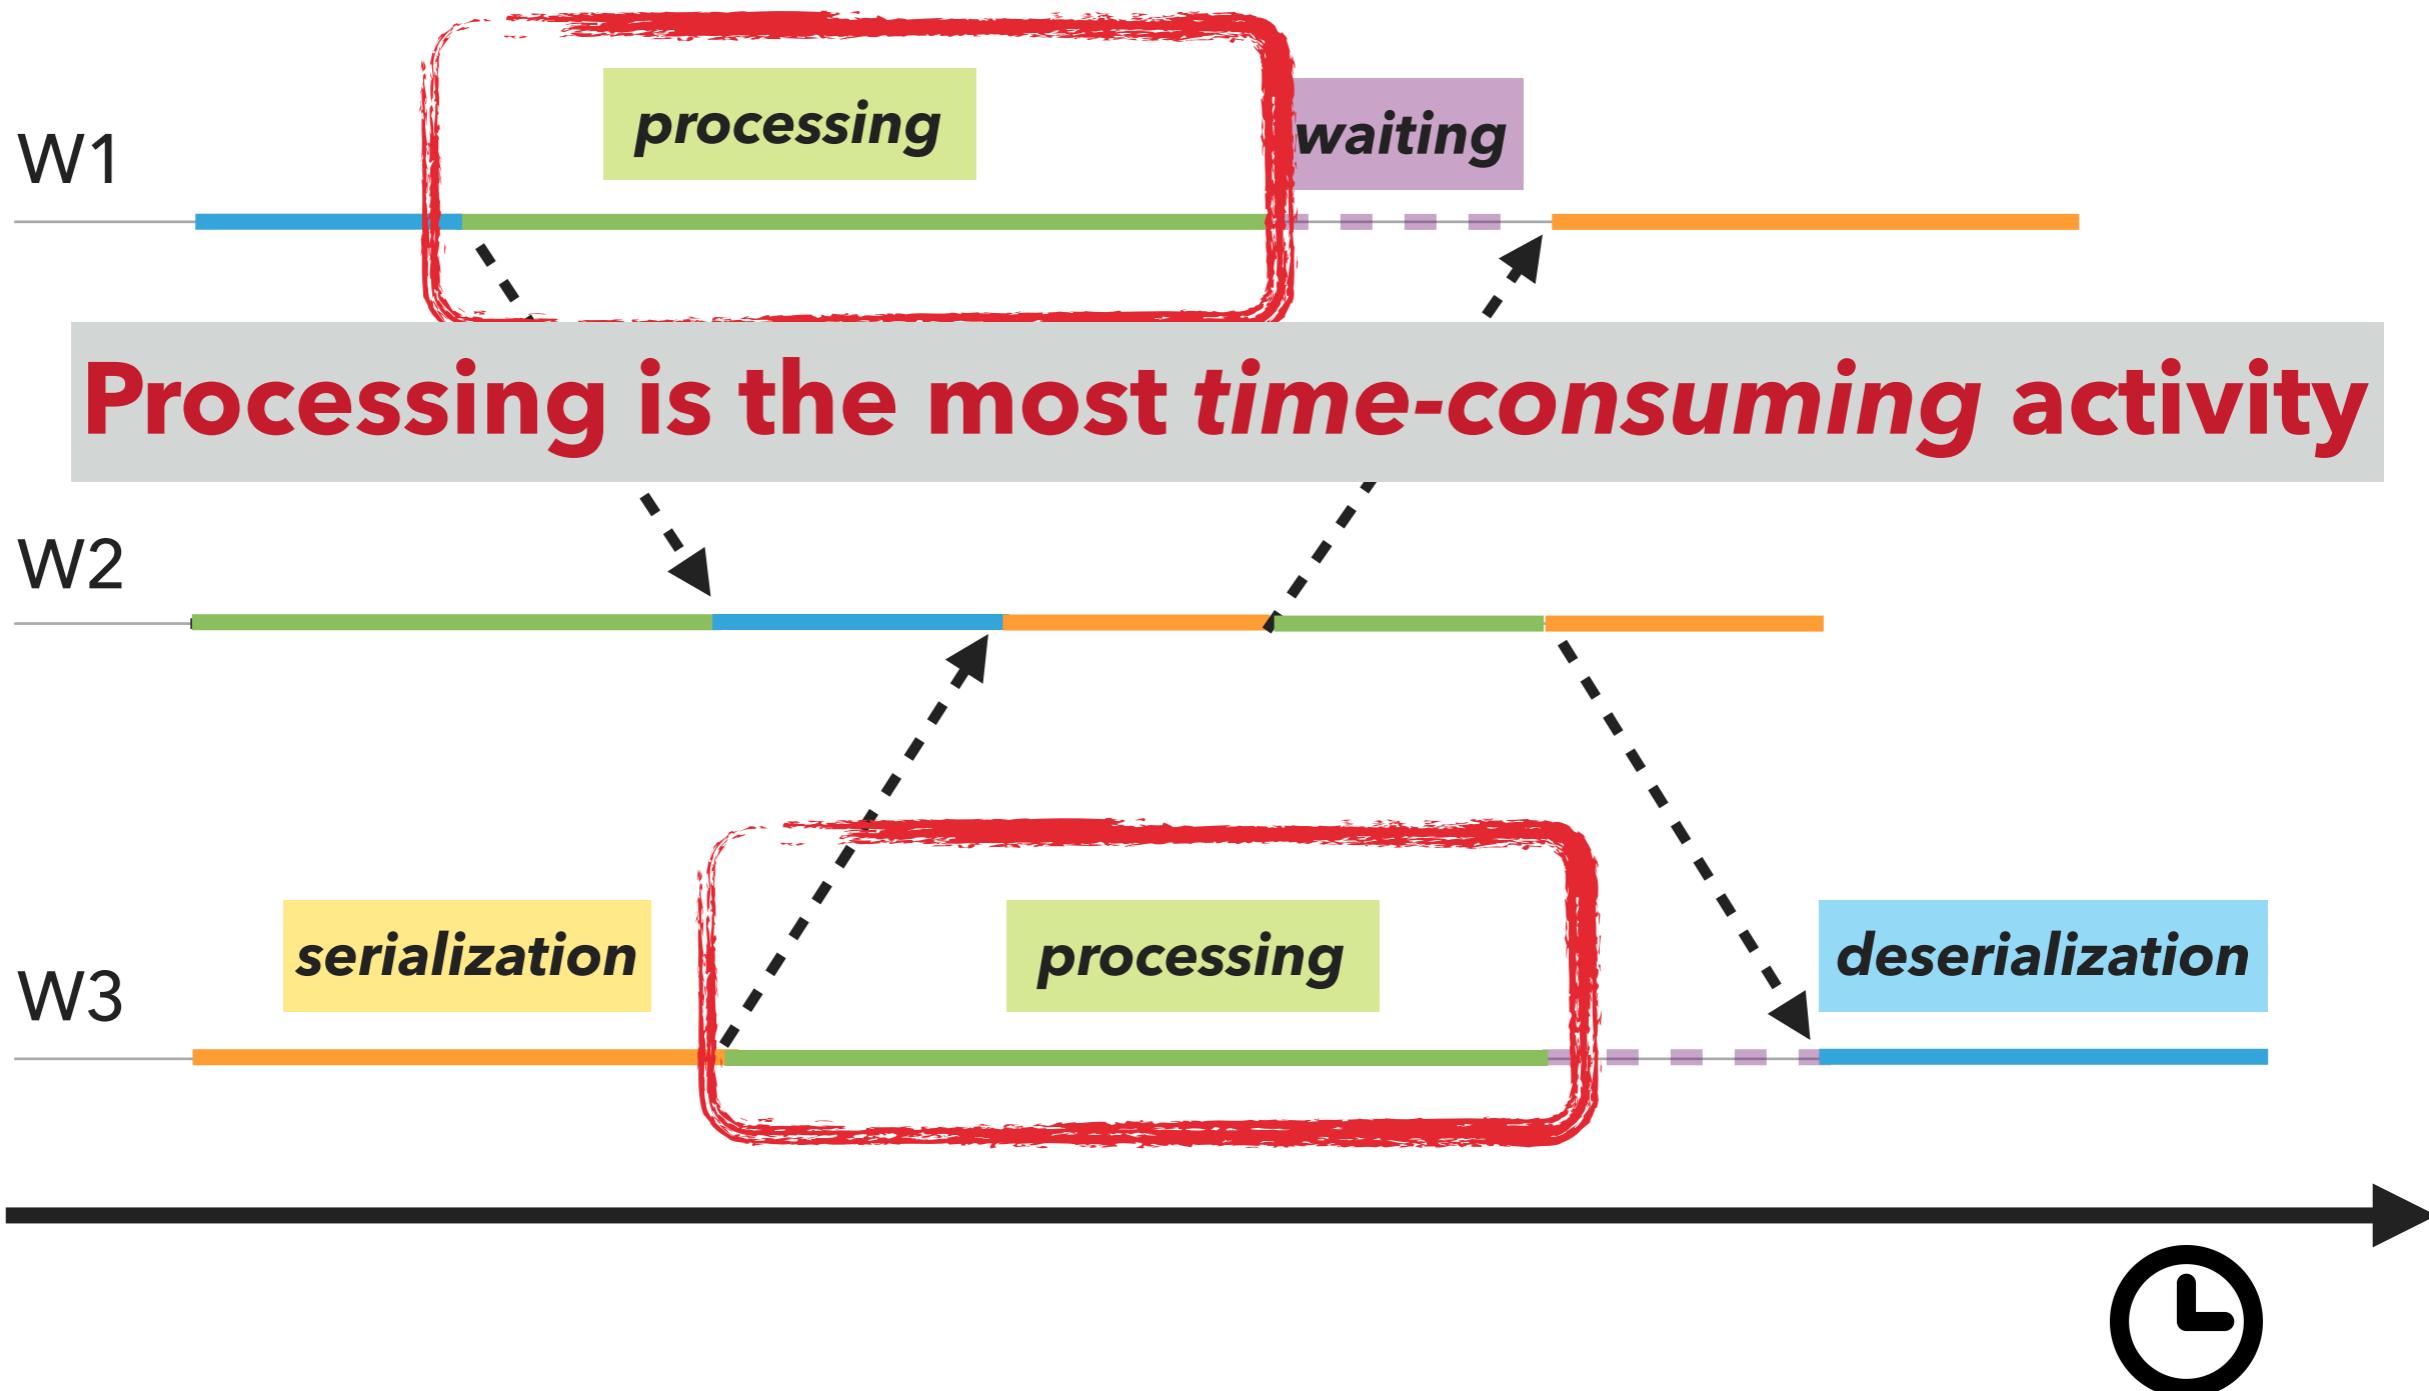
Duration

Aggregate data exchange

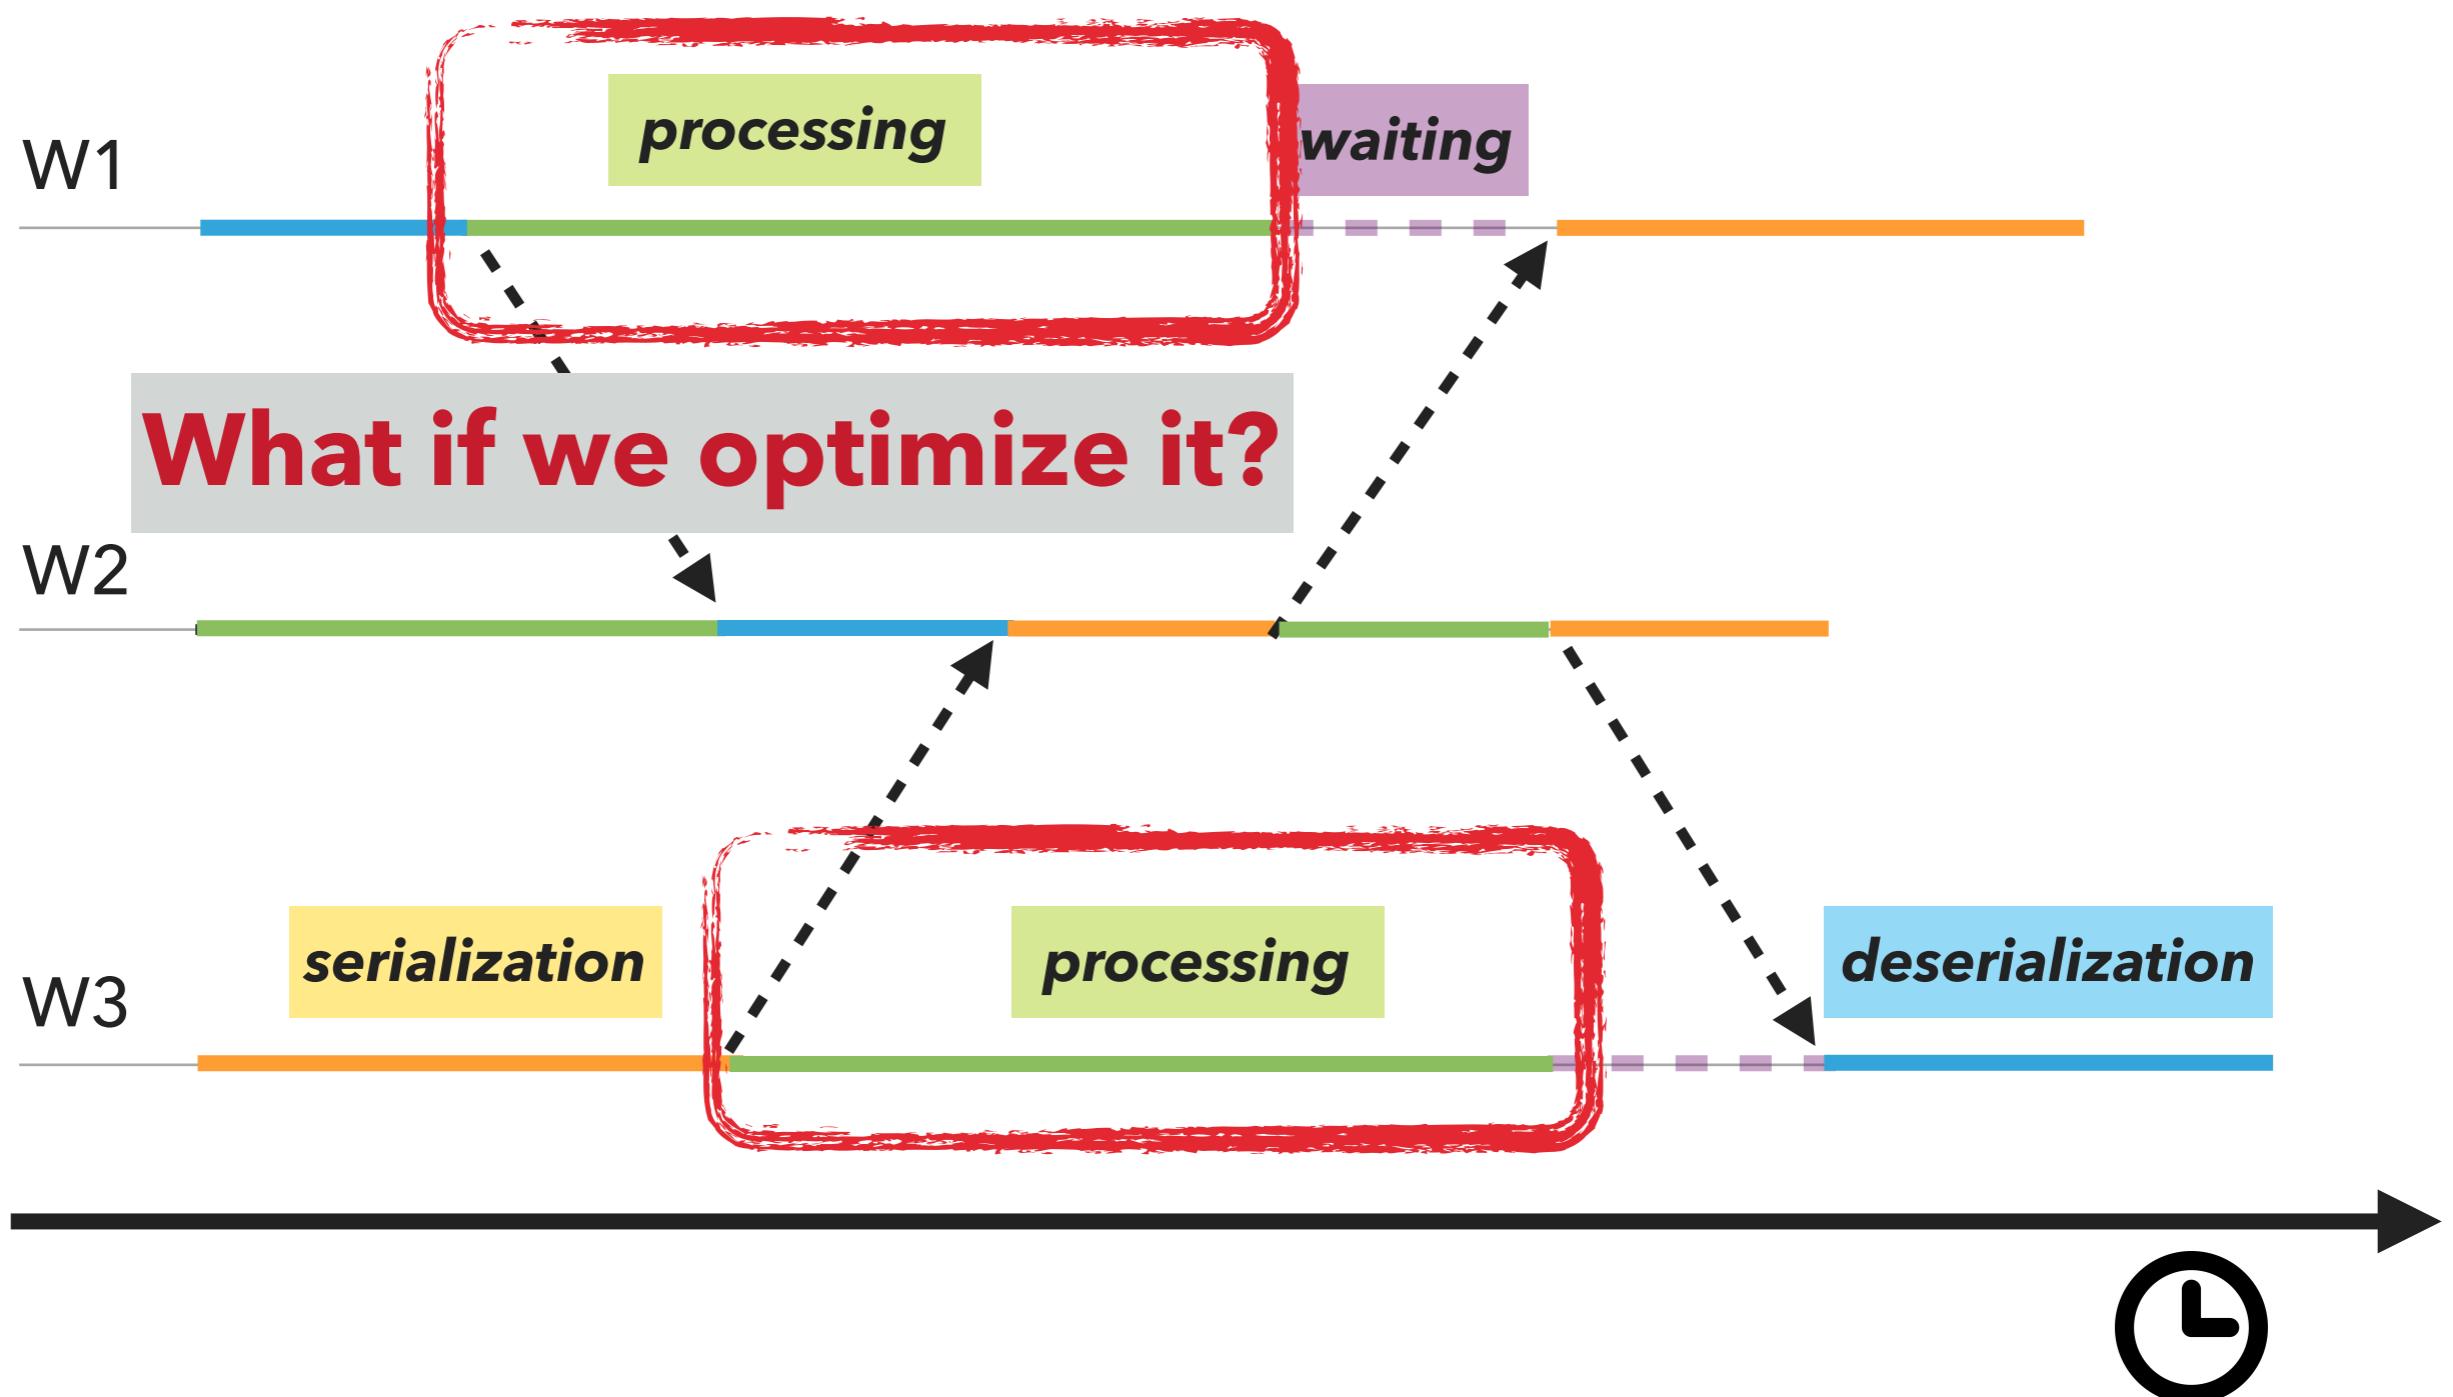
CONVENTIONAL: A BOTTLENECK IN A PARALLEL EXECUTION?



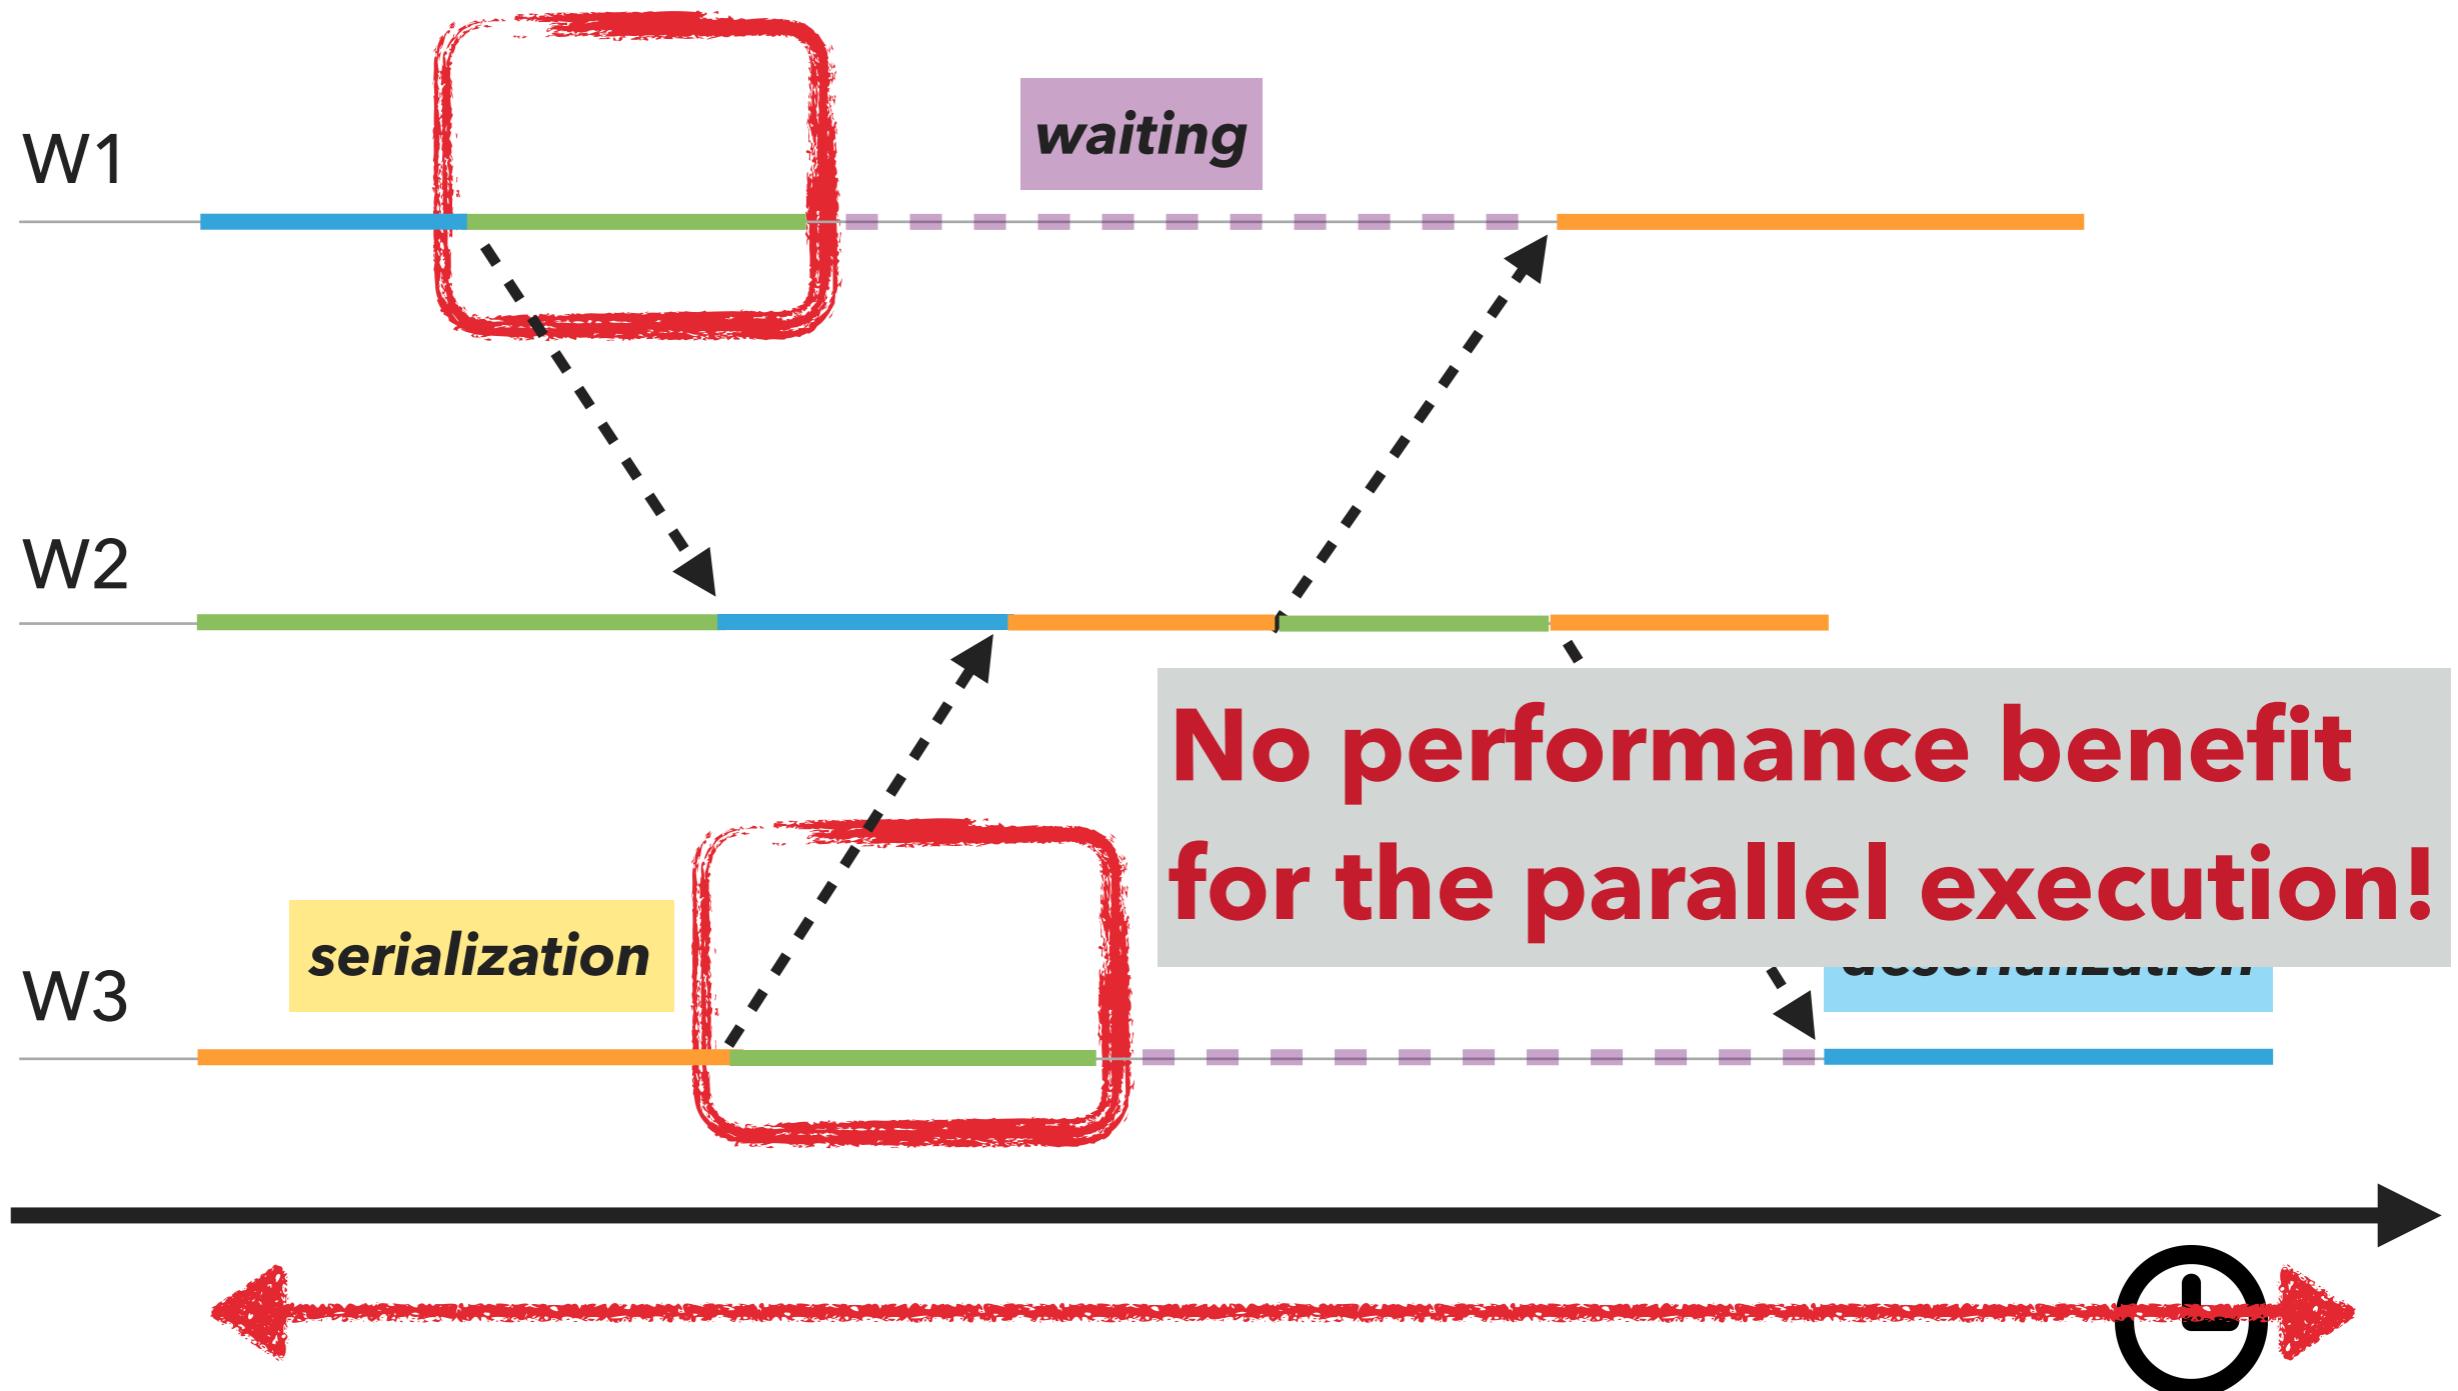
CONVENTIONAL: A BOTTLENECK IN A PARALLEL EXECUTION?



CONVENTIONAL: A BOTTLENECK IN A PARALLEL EXECUTION?



CONVENTIONAL: A BOTTLENECK IN A PARALLEL EXECUTION?

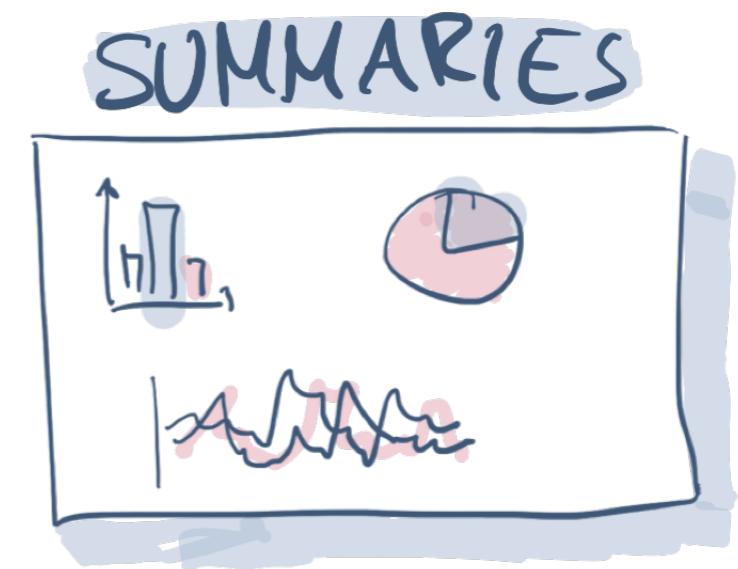
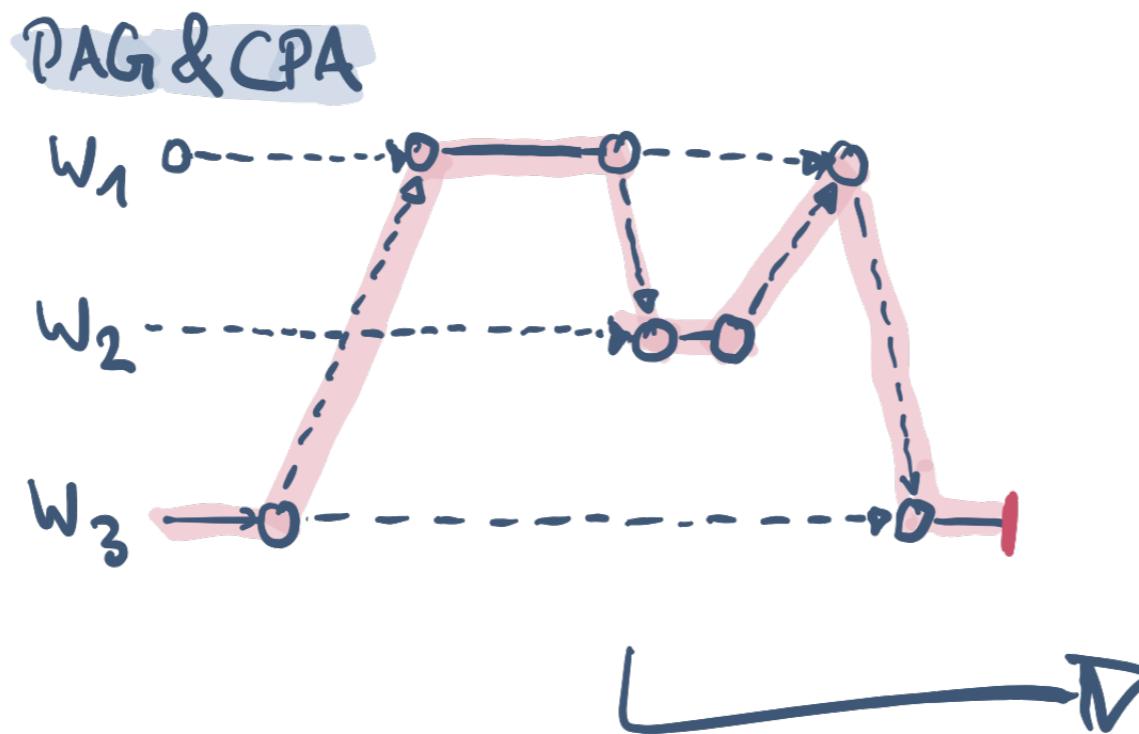
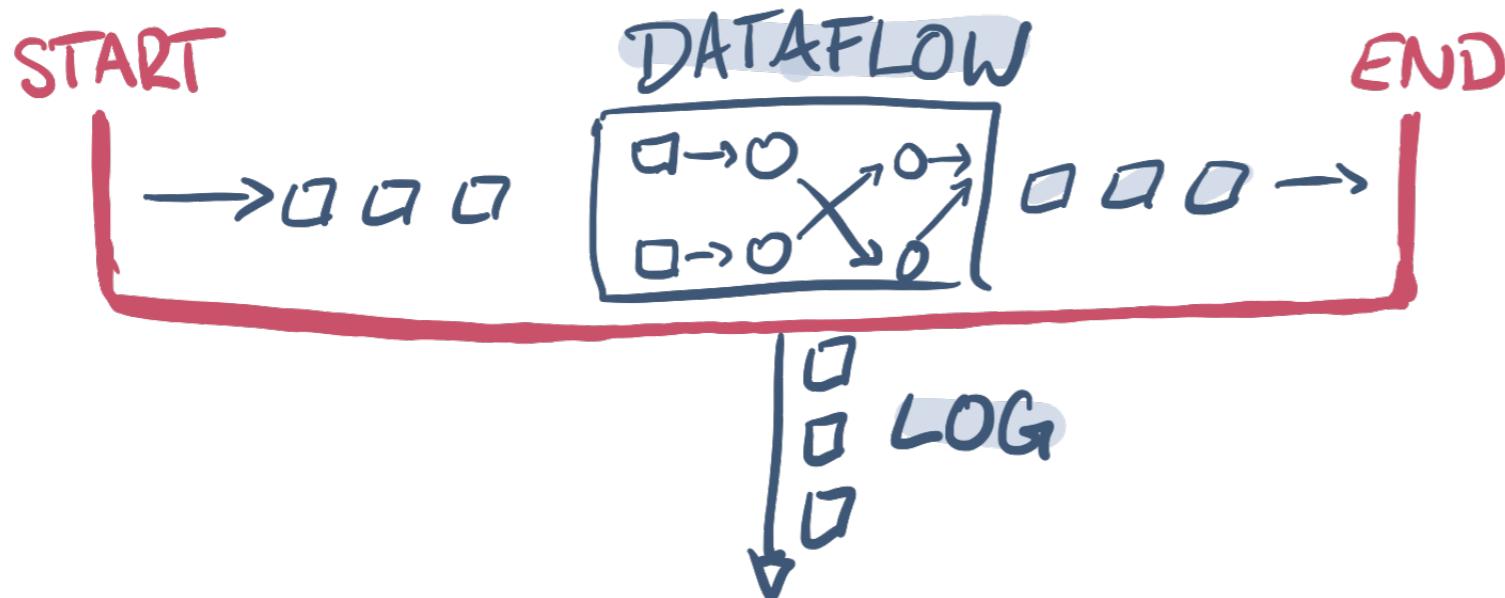


**CONVENTIONAL
PROFILING CAN
BE MISLEADING**

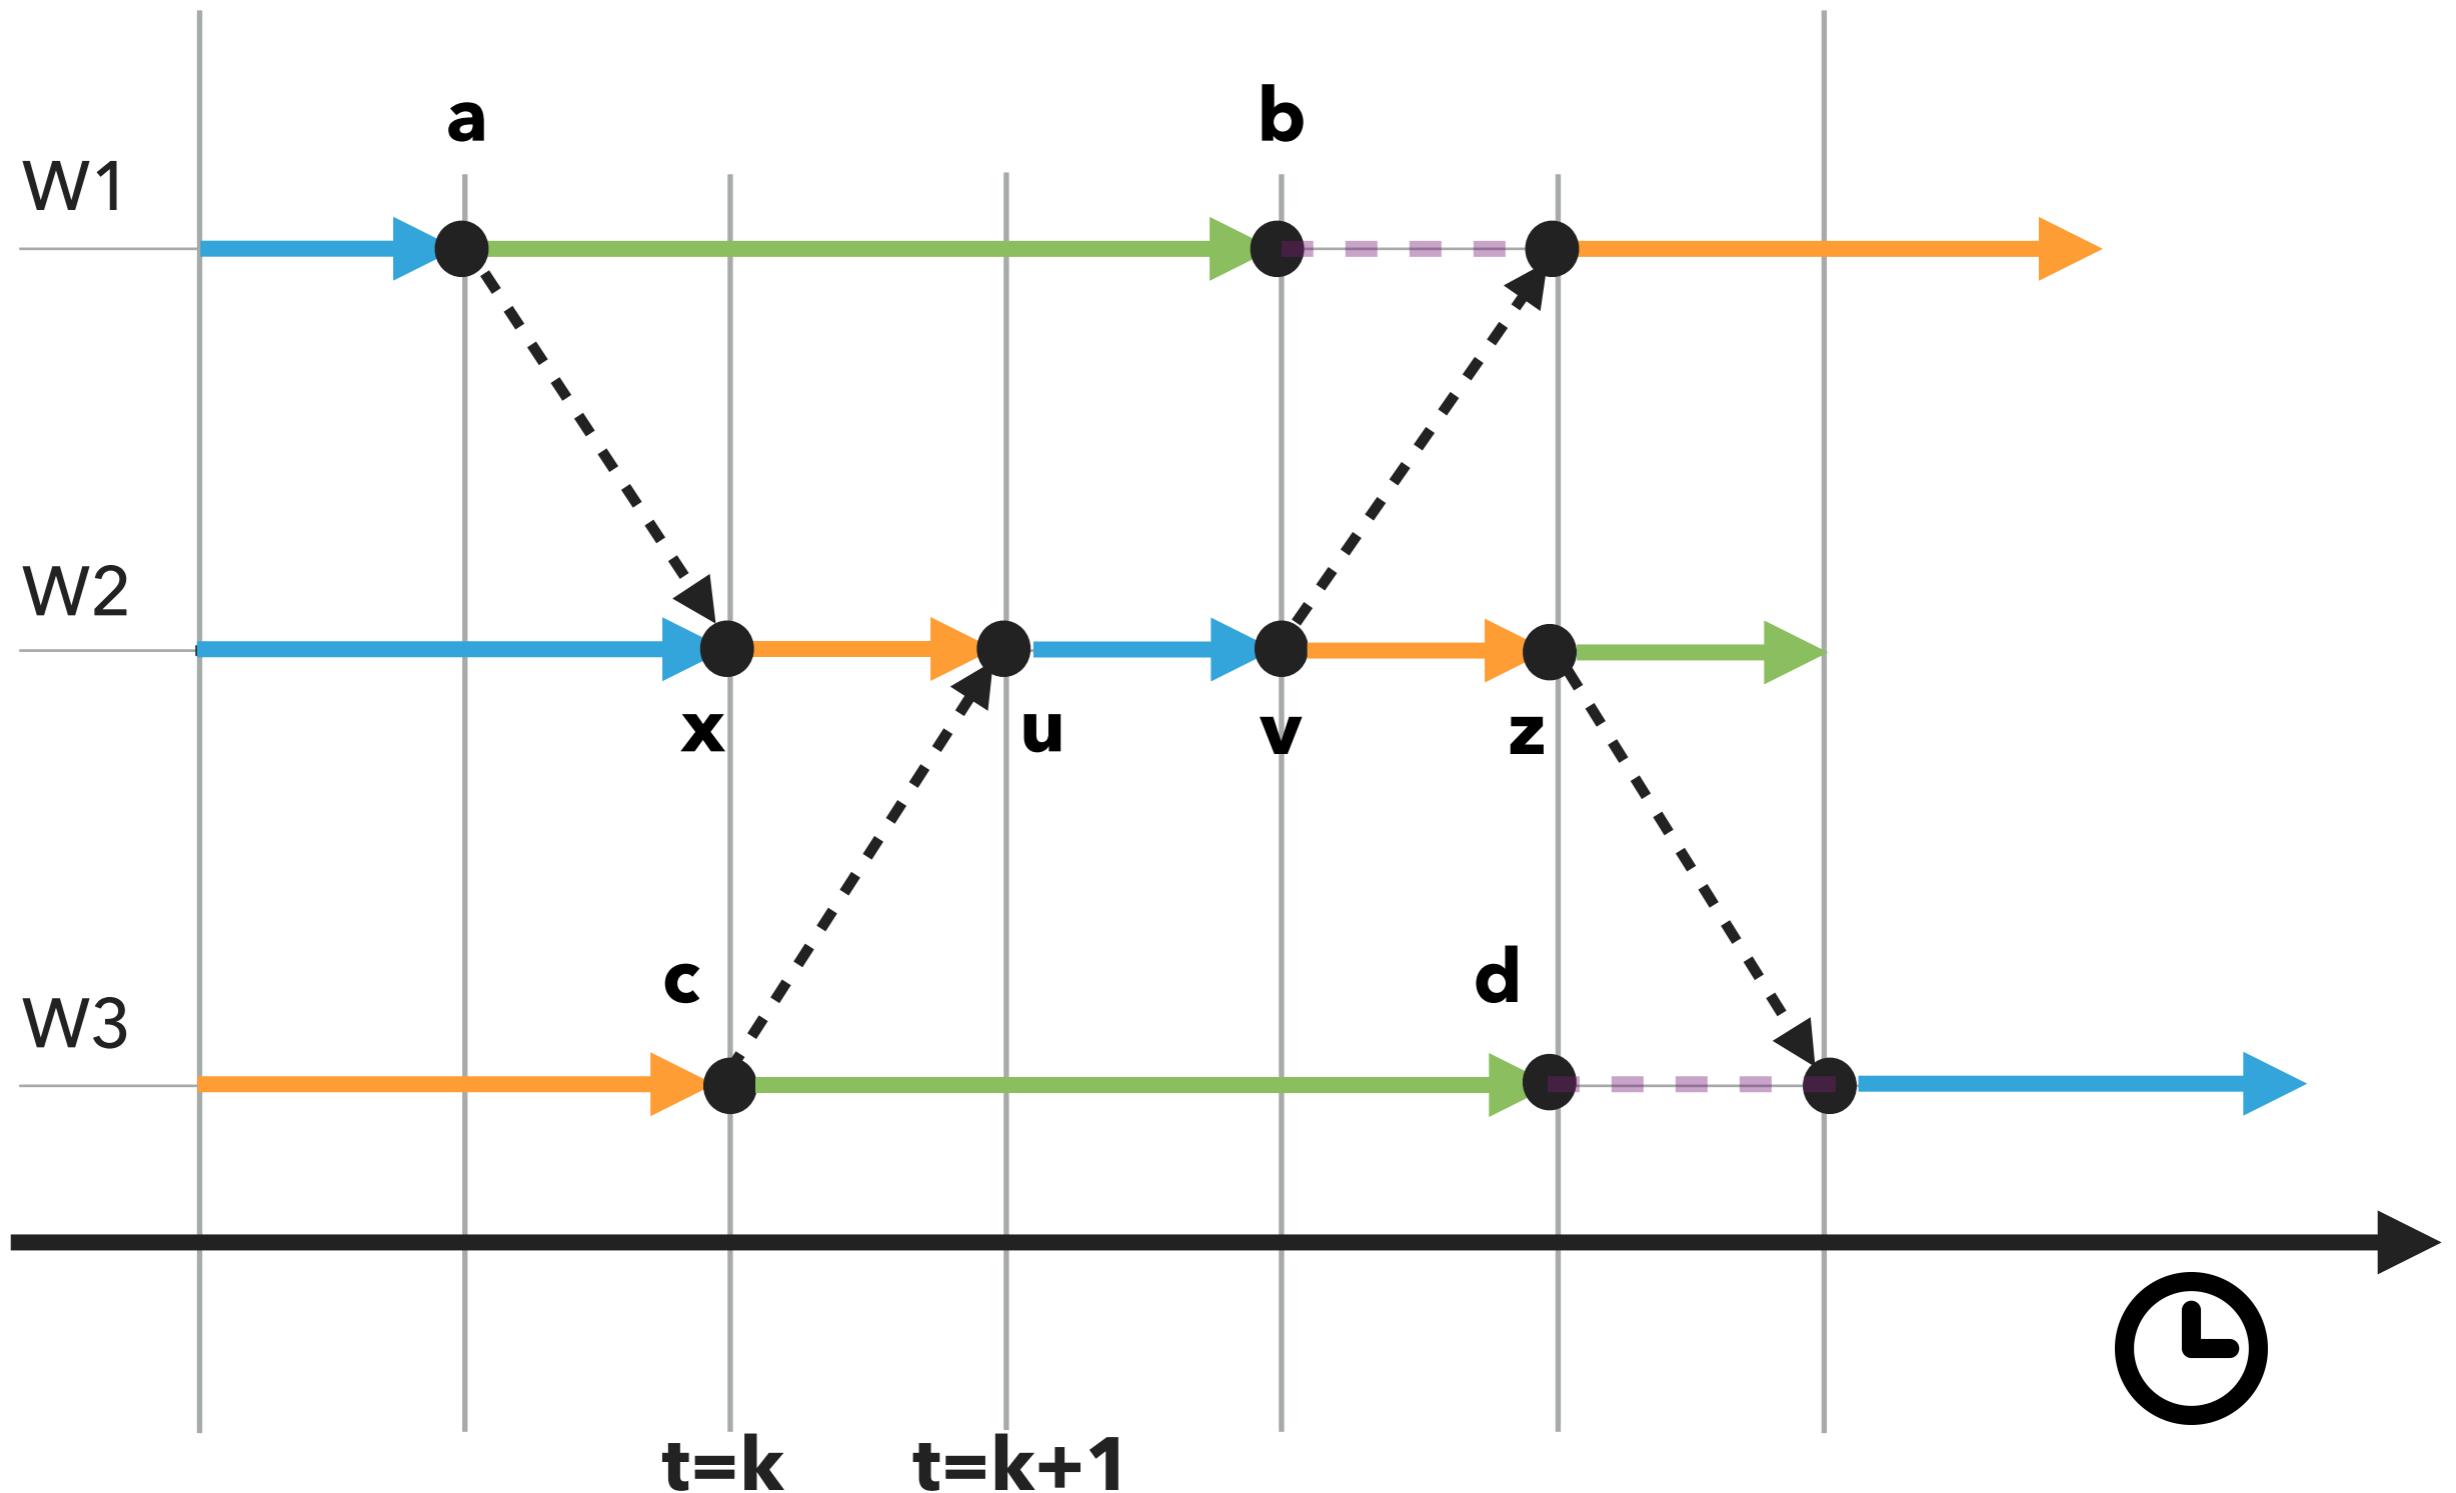


CRITICAL PATH ANALYSIS

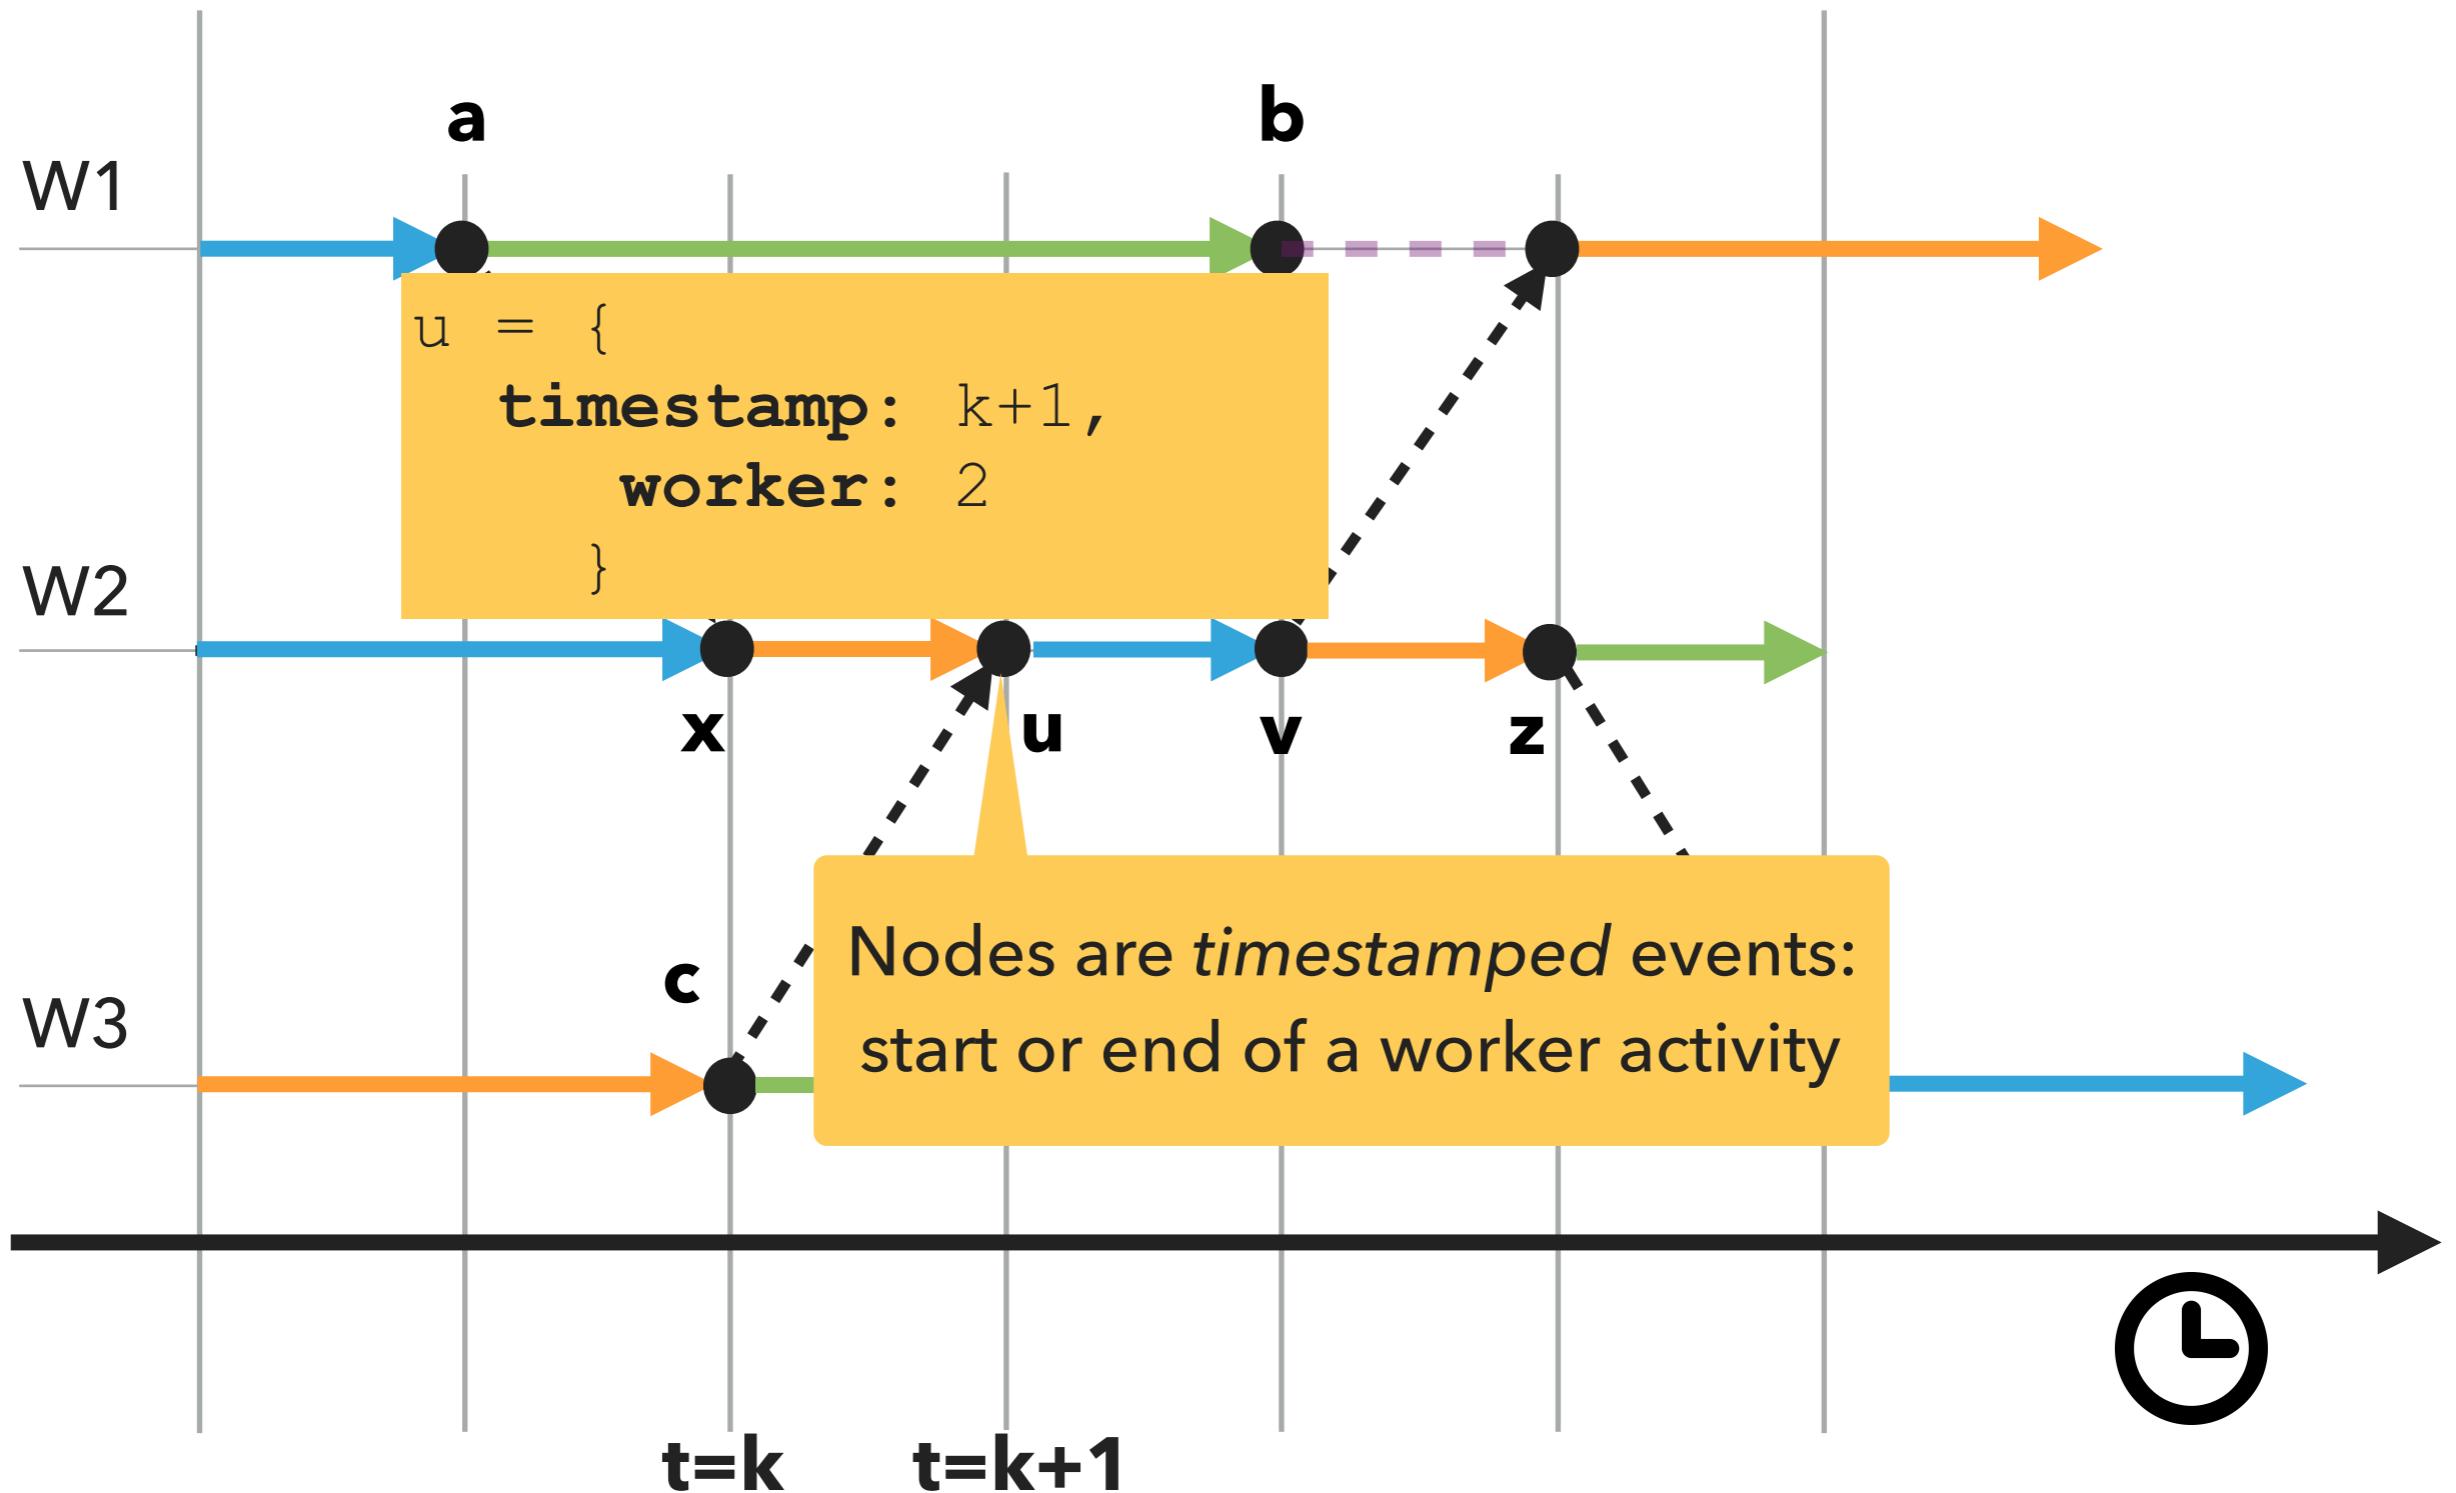
CRITICAL PATH ANALYSIS



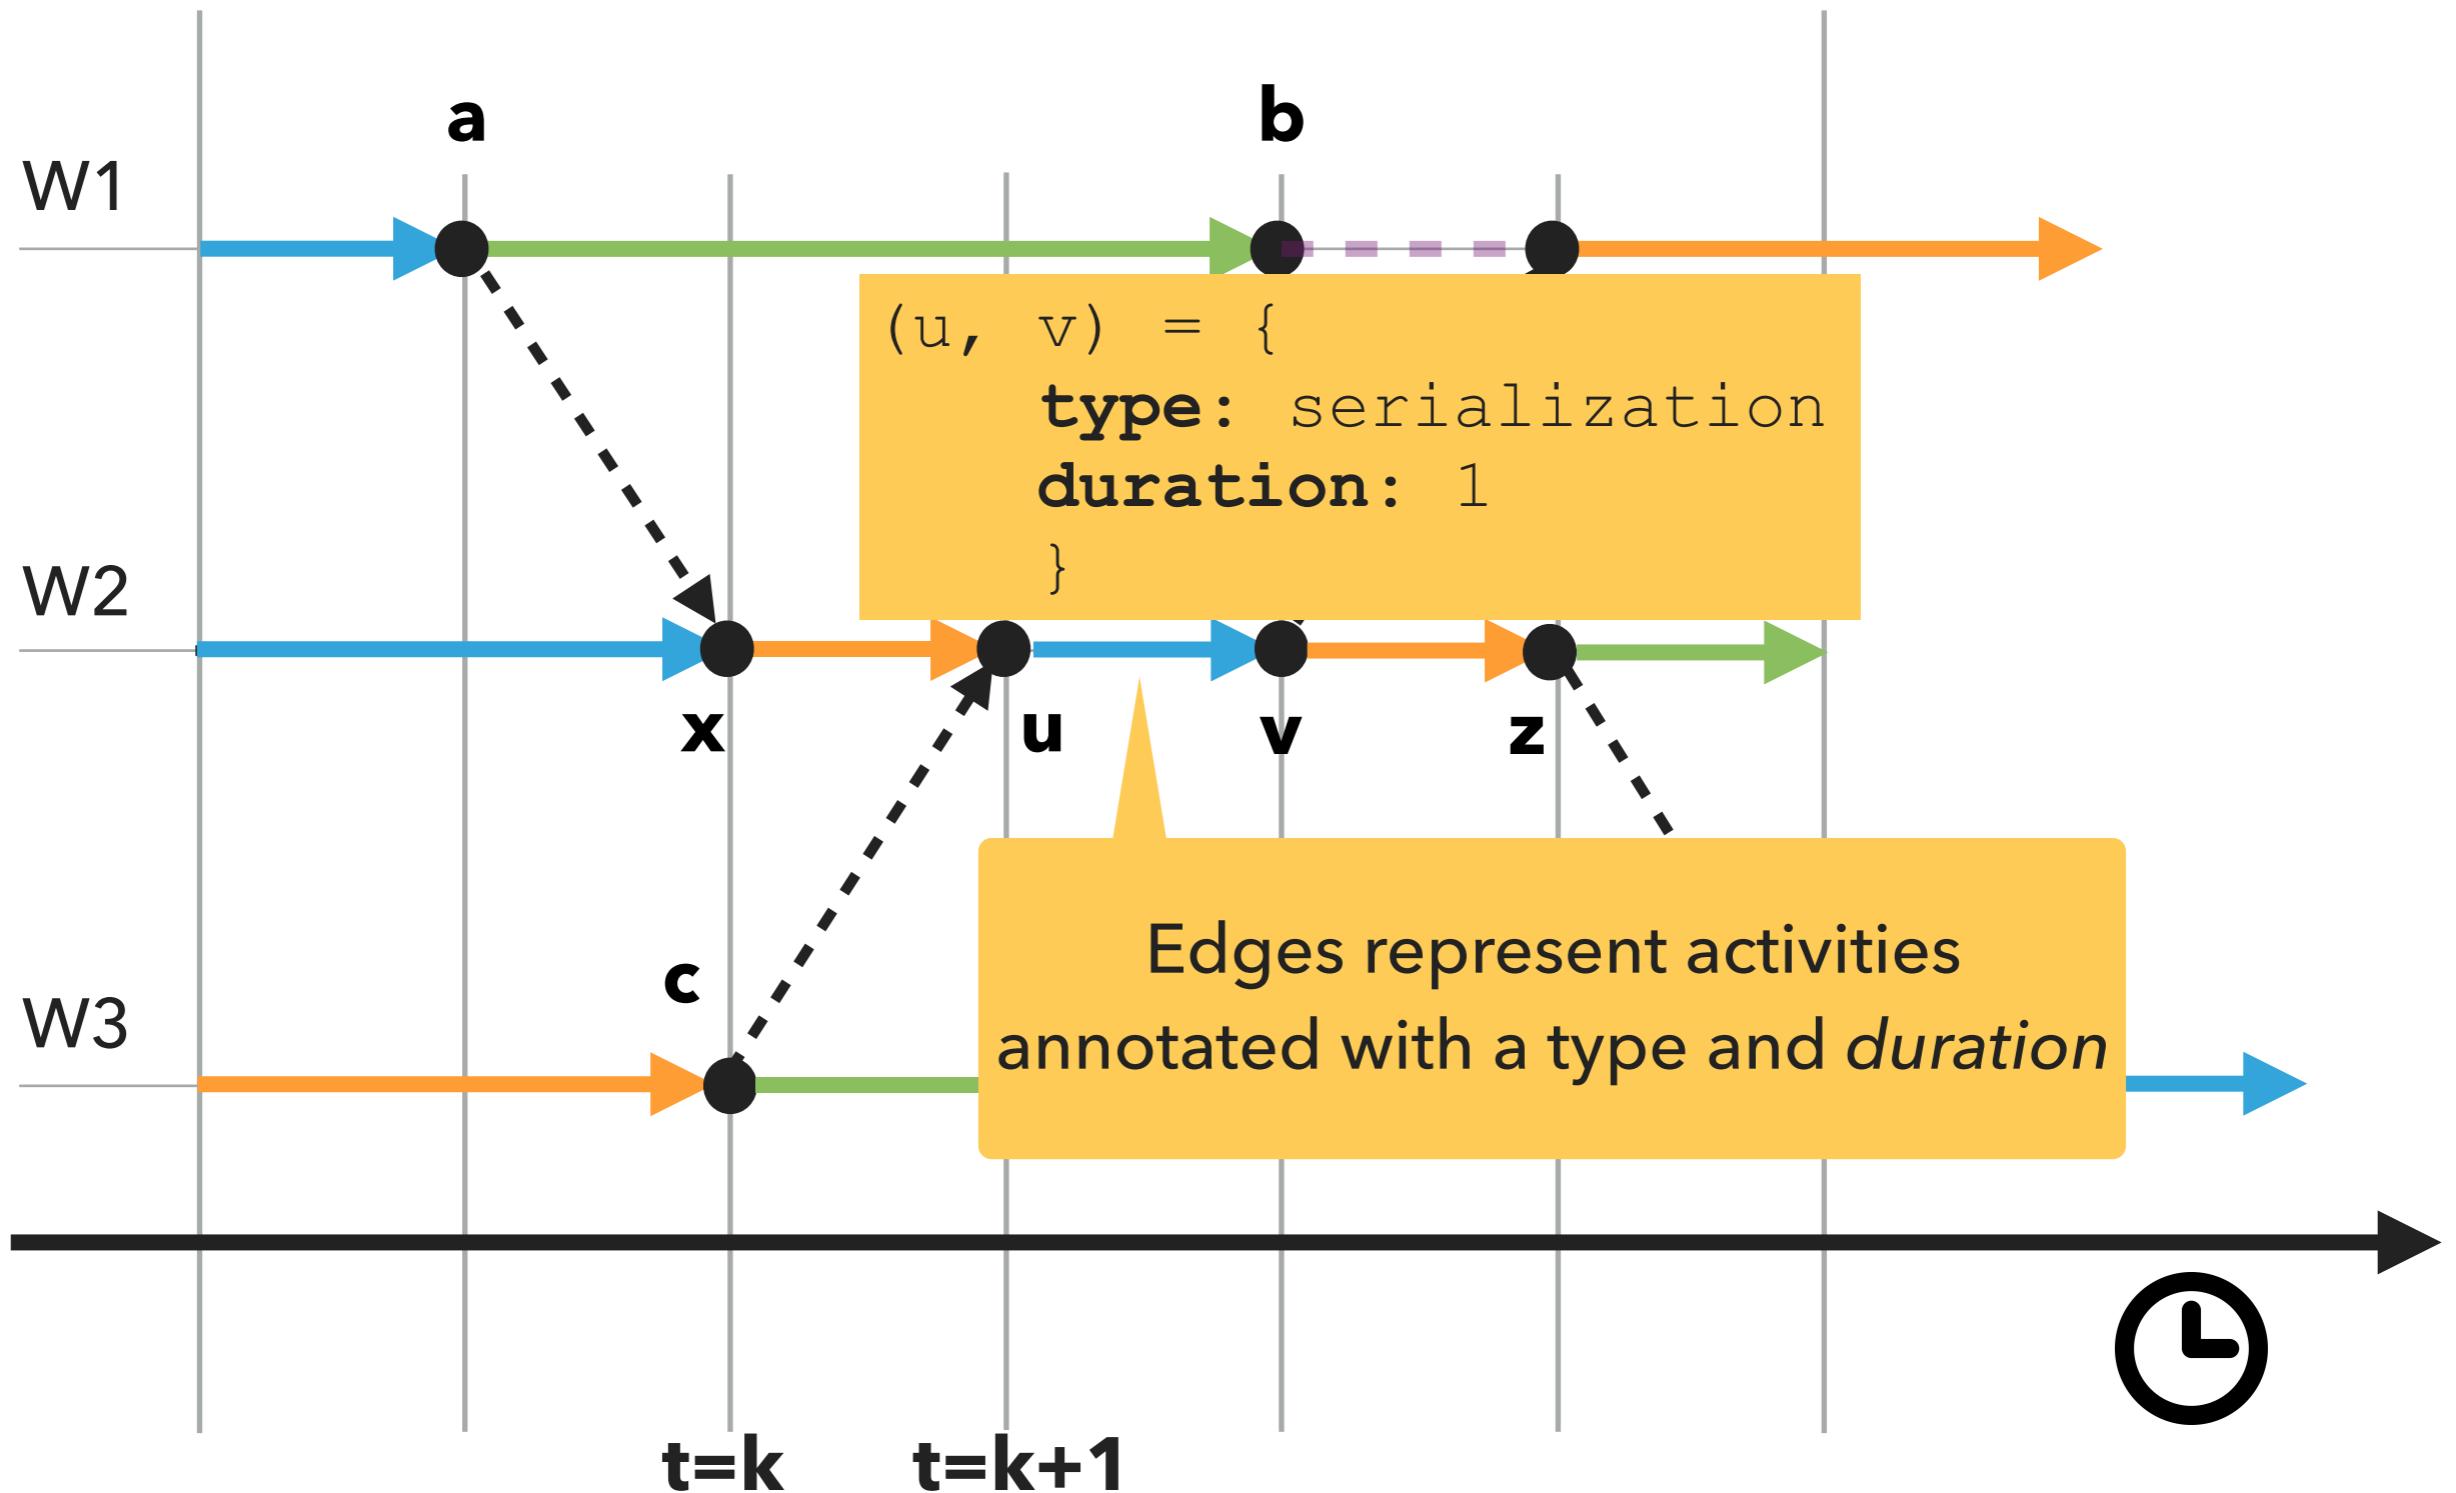
THE PROGRAM ACTIVITY GRAPH (PAG)



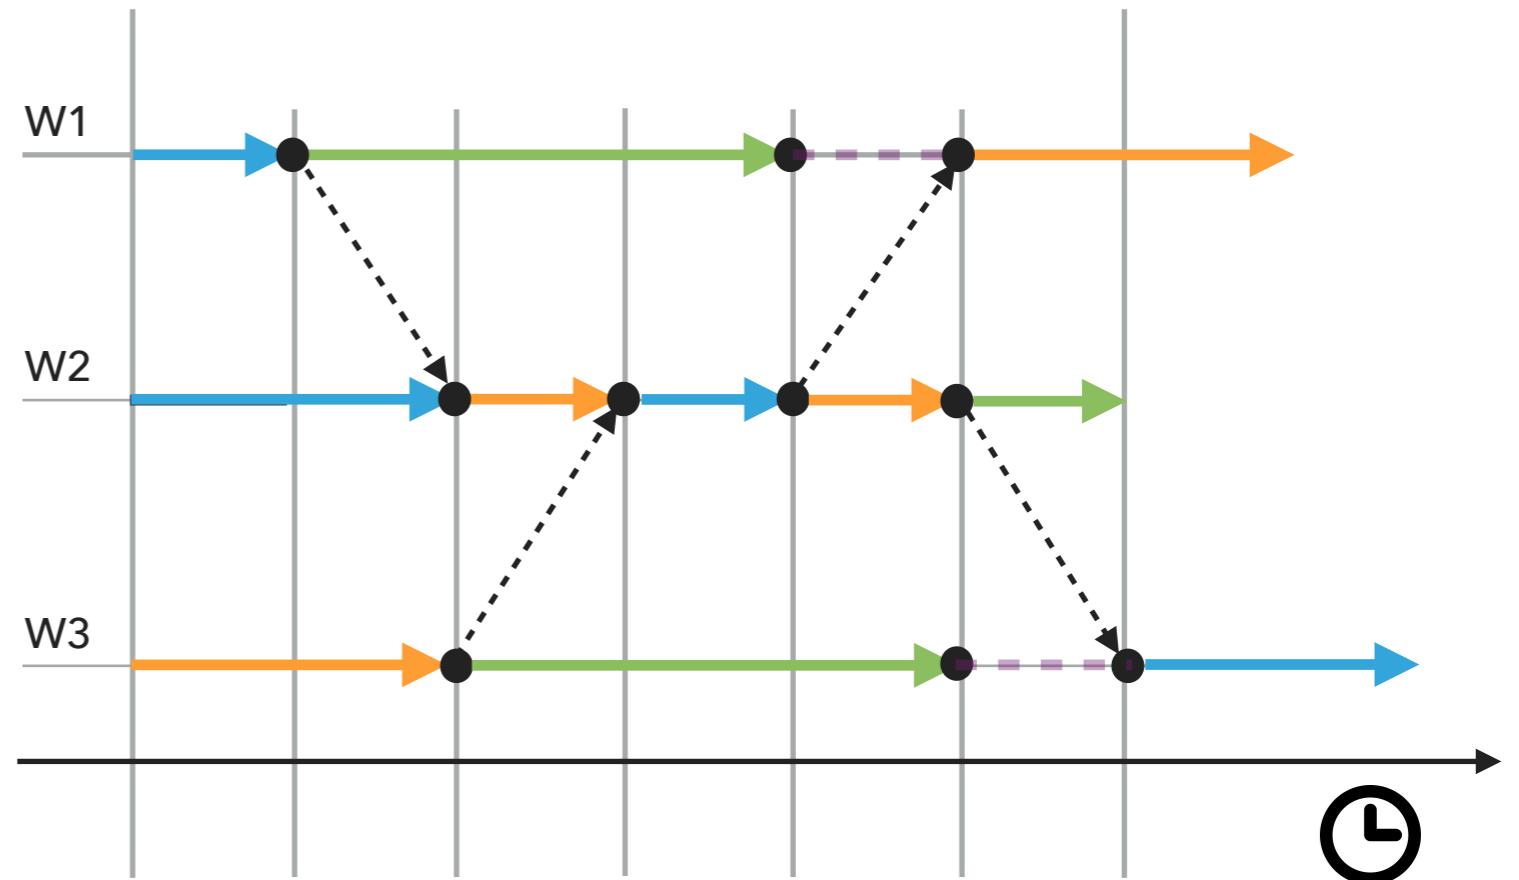
THE PROGRAM ACTIVITY GRAPH (PAG)



THE PROGRAM ACTIVITY GRAPH (PAG)

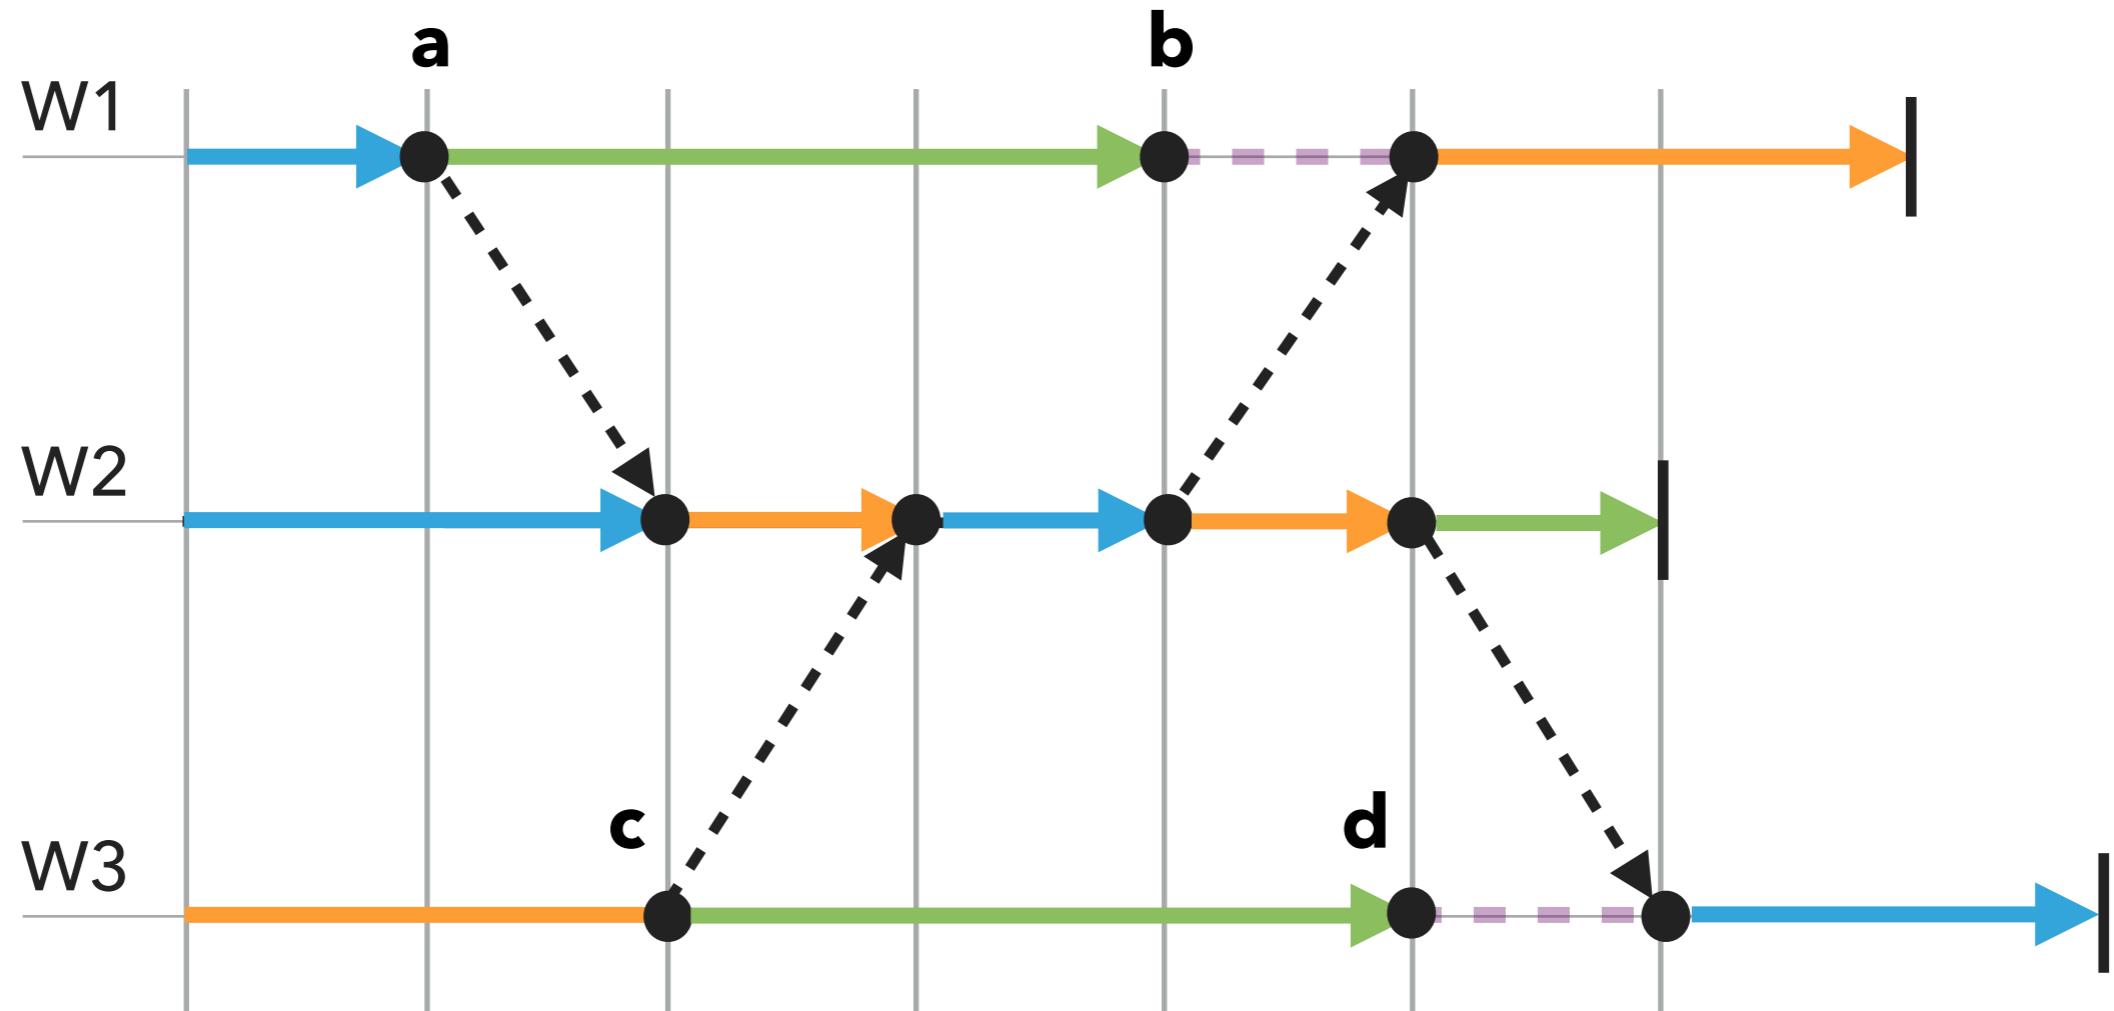


The Program Activity
Graph captures
computational
dependencies among
parallel workers



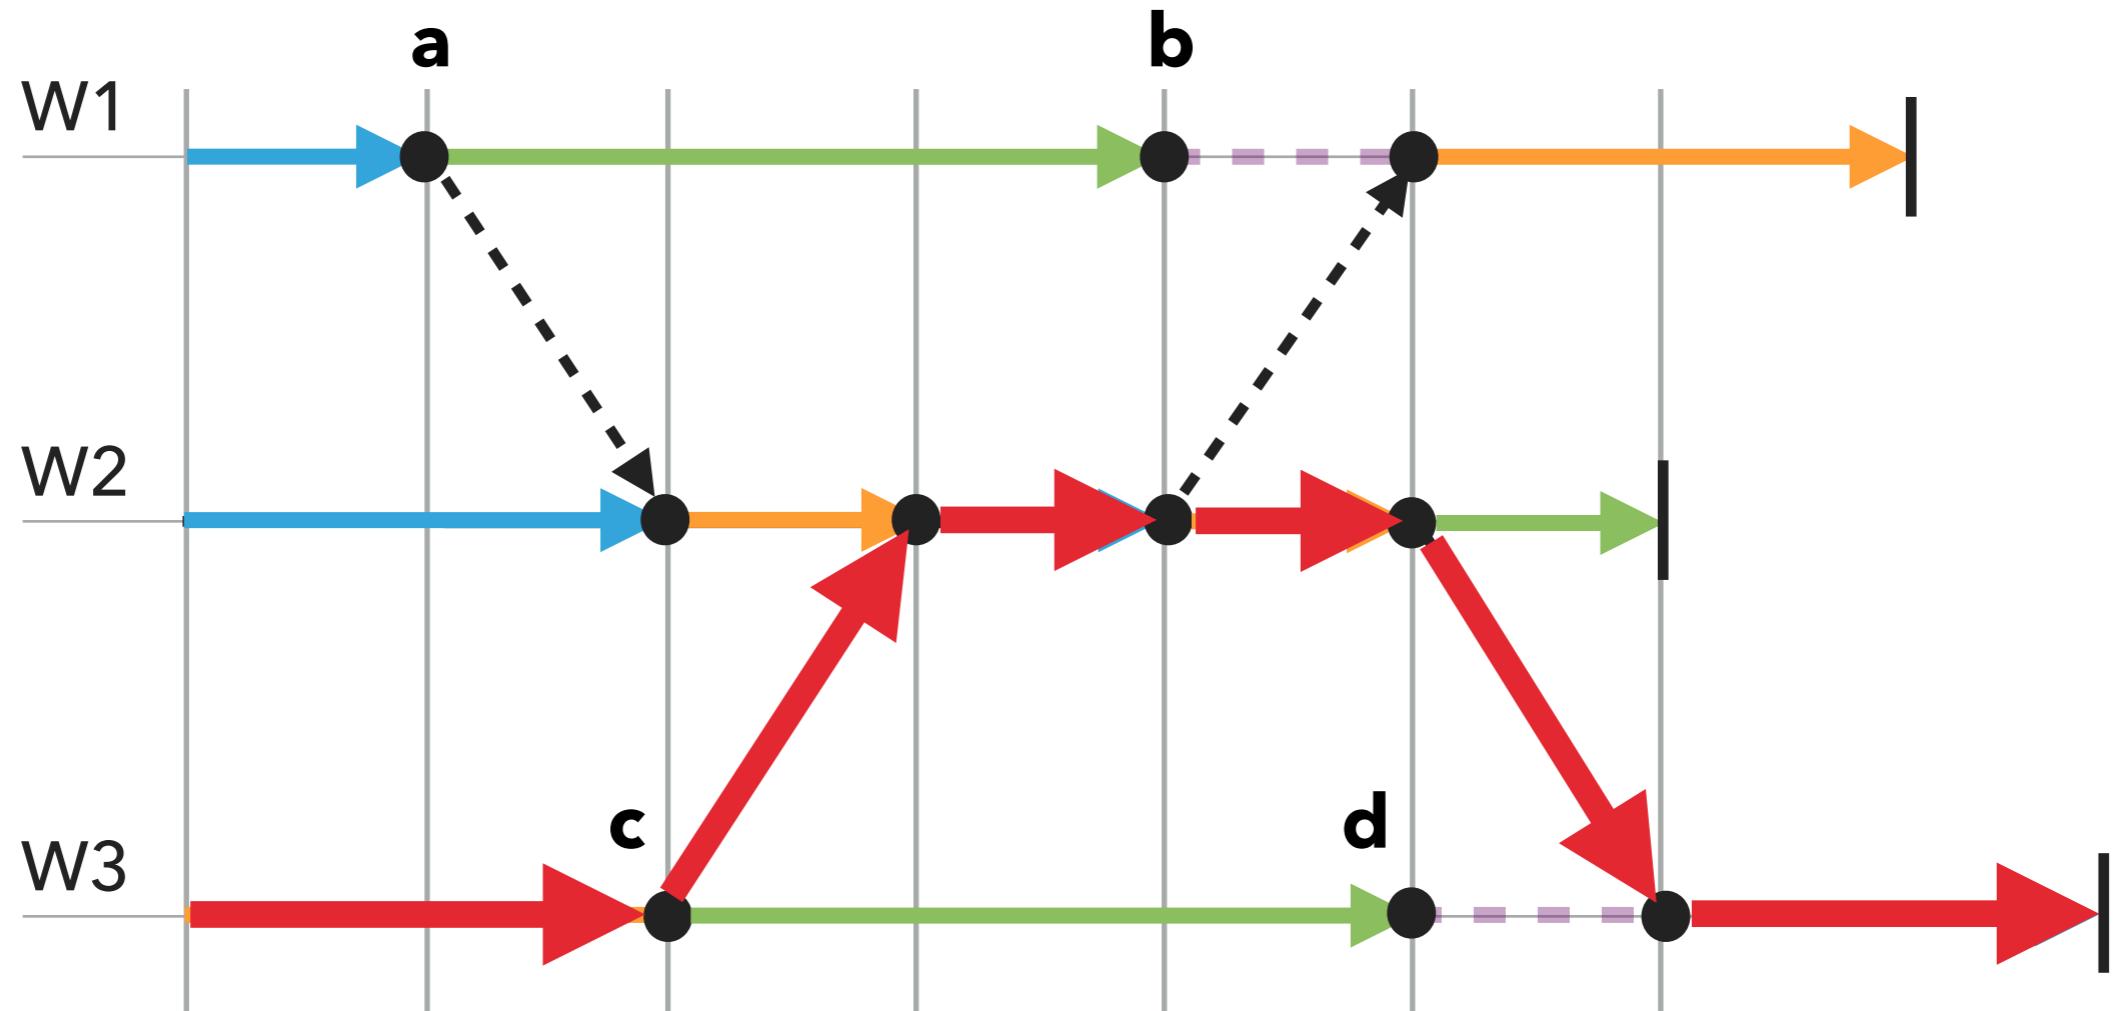
- ▶ Which activities **delay** the overall execution?
- ▶ i.e. which activities lie on the **critical path** of execution?

CRITICAL PATH: A BOTTLENECK IN A PARALLEL EXECUTION!



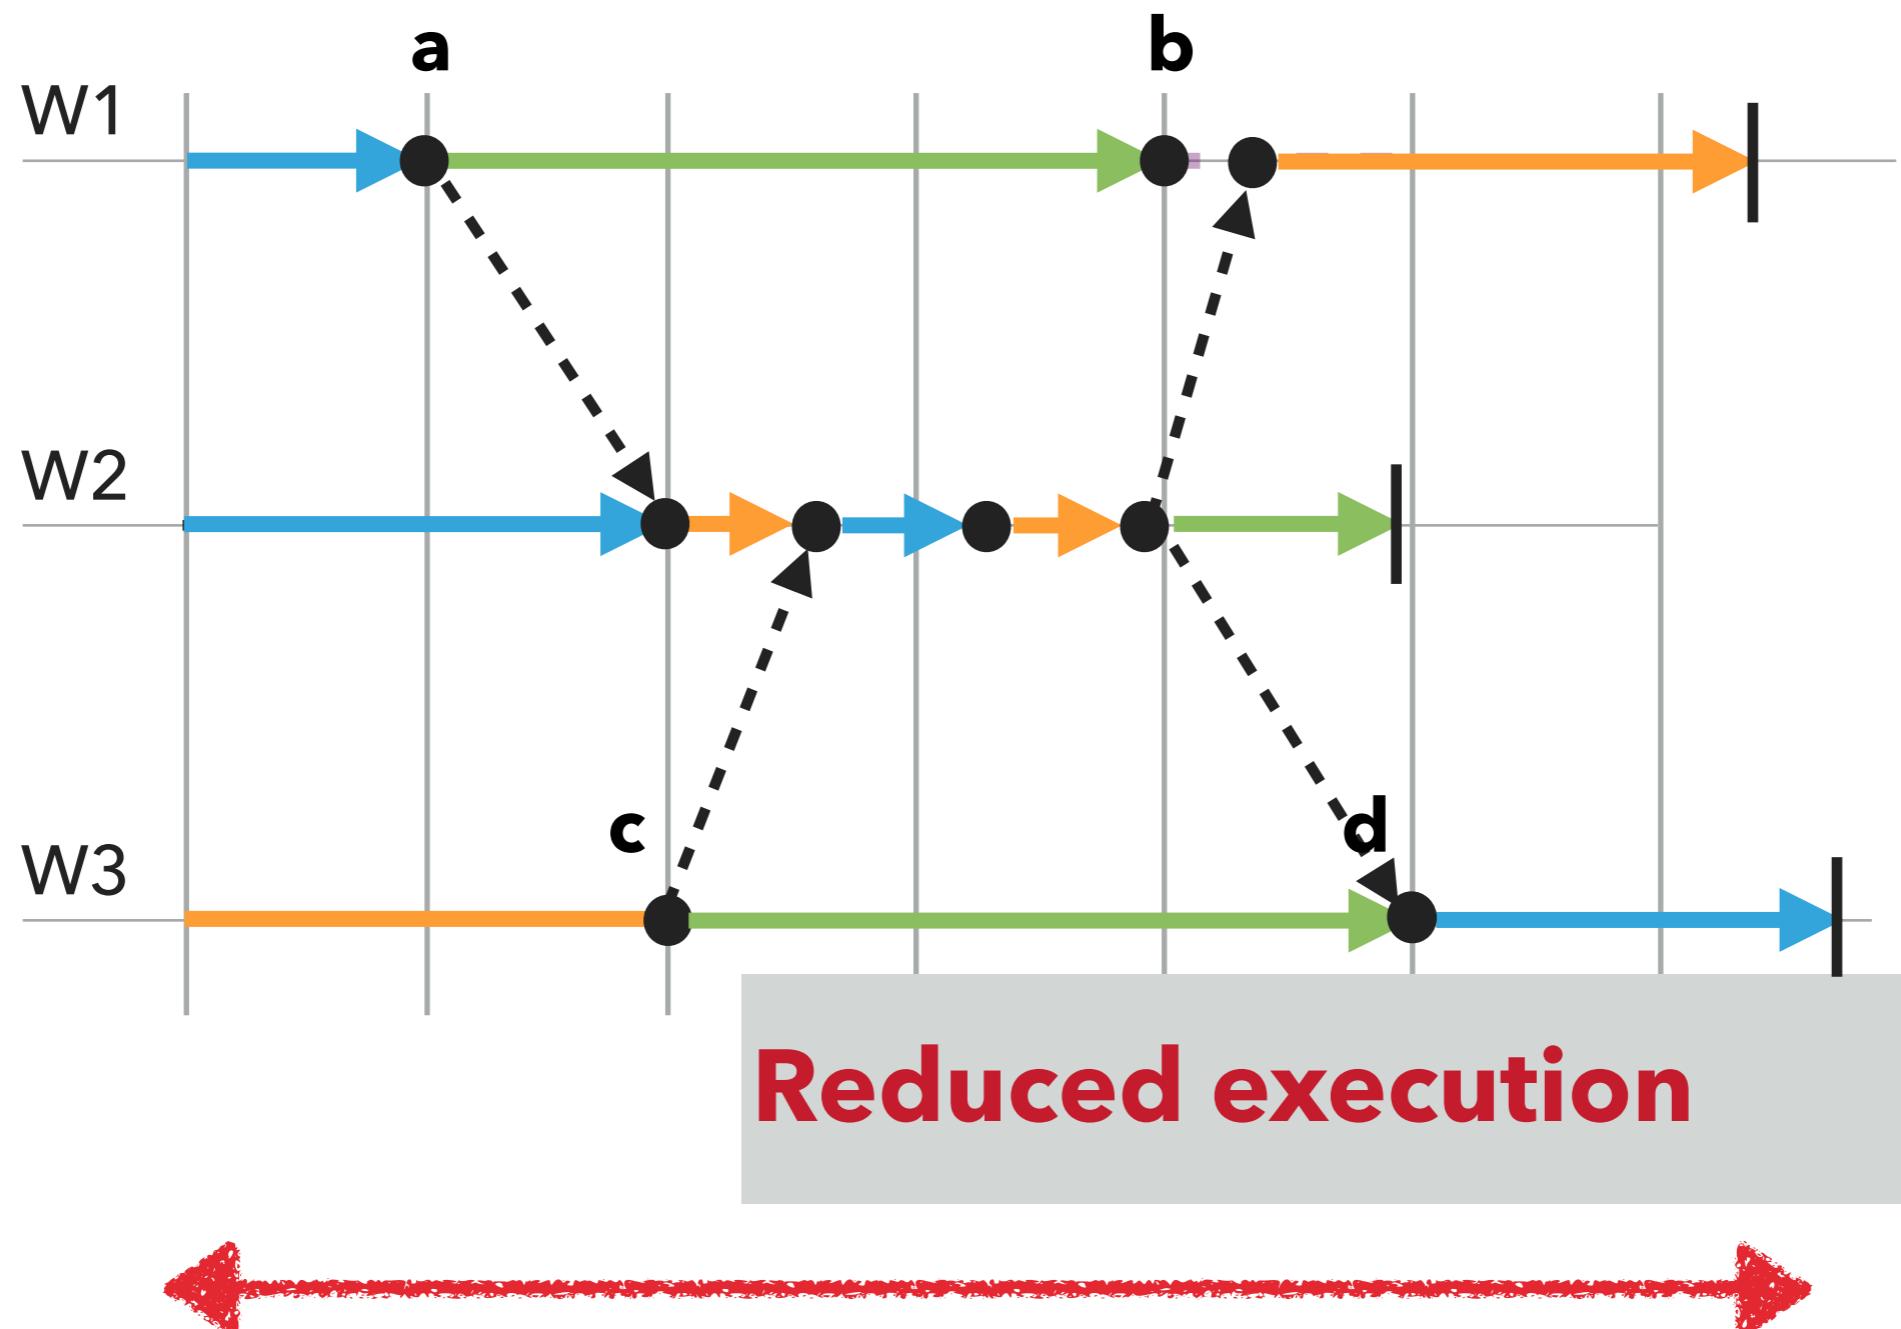
The **longest path** in the execution history
(not considering waiting activities)

CRITICAL PATH: A BOTTLENECK IN A PARALLEL EXECUTION!



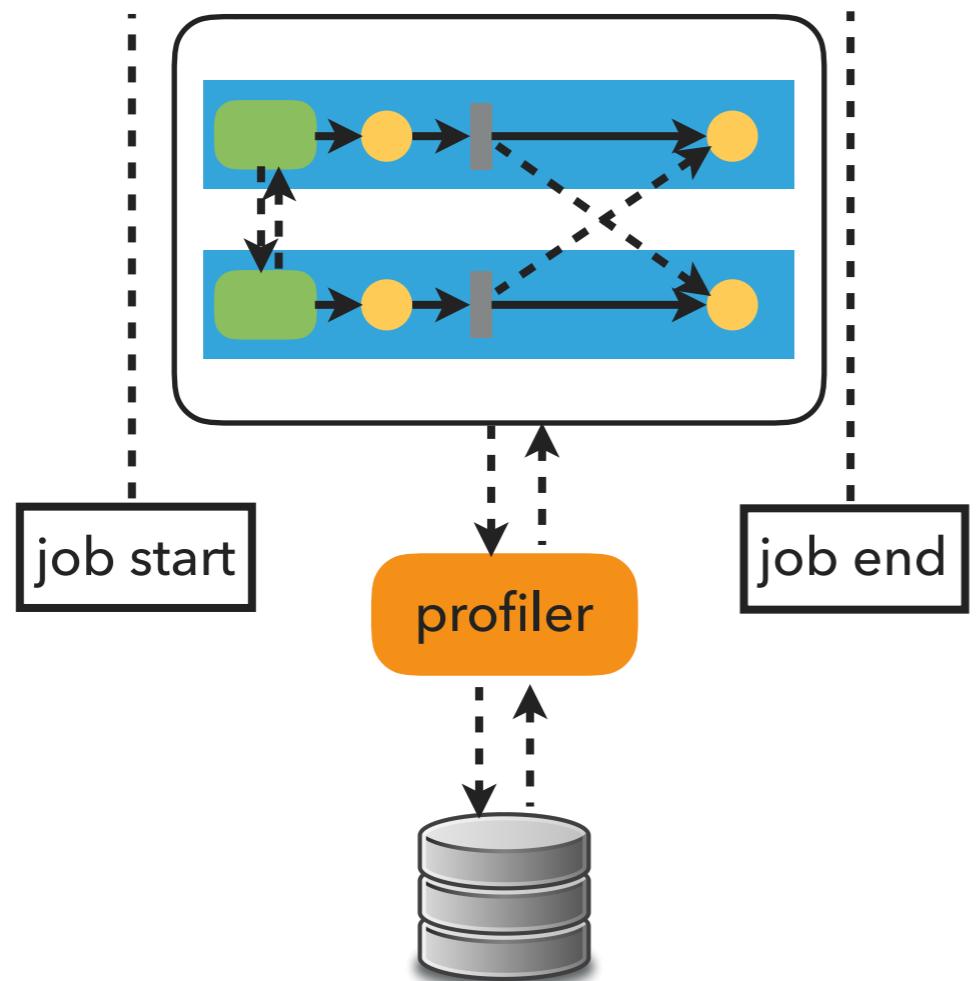
The **longest path** in the execution history
(not considering waiting activities)

CRITICAL PATH: A BOTTLENECK IN A PARALLEL EXECUTION!



POST-MORTEM CRITICAL PATH ANALYSIS IS EASY

1. Collect traces during execution



2. Analyze traces offline



Streaming applications run **continuously** with potentially **unbounded** input!

- ▶ There might be no “job end”
- ▶ The PAG and critical path are continuously evolving
- ▶ Stale profiling information is not useful



How do we compute the critical path online?

SnailTrail: Generalizing Critical Paths for Online Analysis of Distributed Dataflows*

strymon.systems.ethz.ch

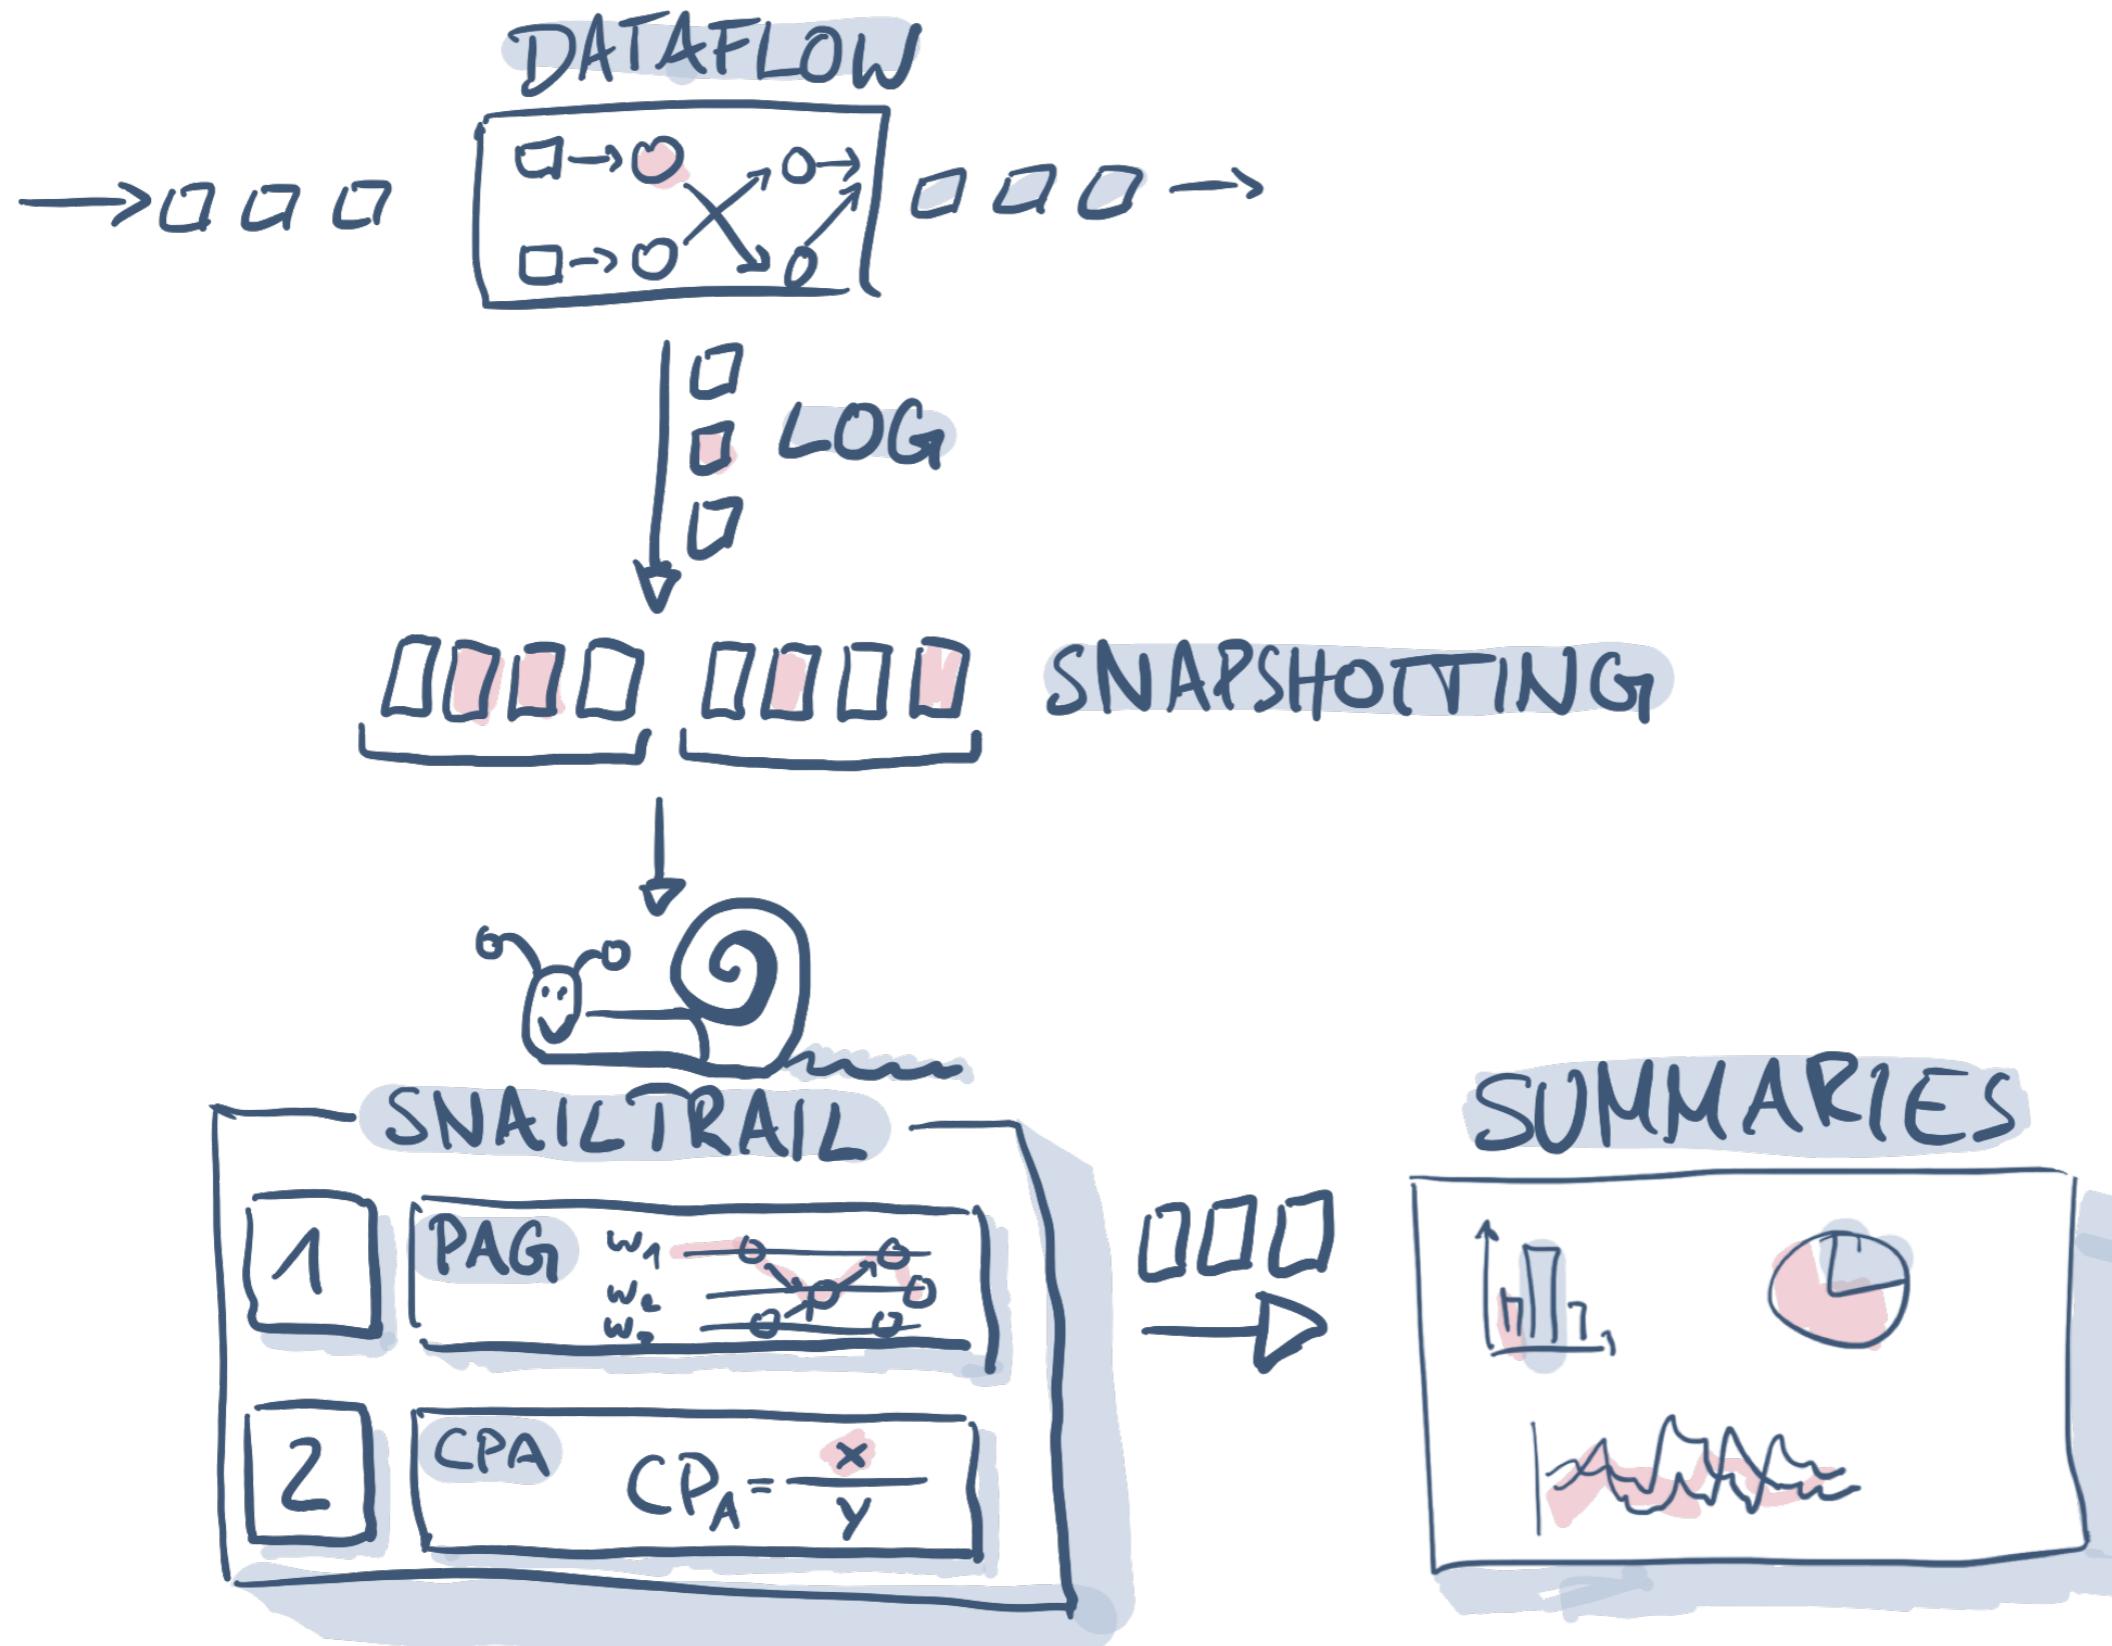
*Moritz Hoffmann, Andrea Lattuada, John Liagouris, Vasiliki Kalavri,
Desislava Dimitrova, Sebastian Wicki, Zaheer Chothia, Timothy Roscoe*

Systems Group, Department of Computer Science, ETH Zürich

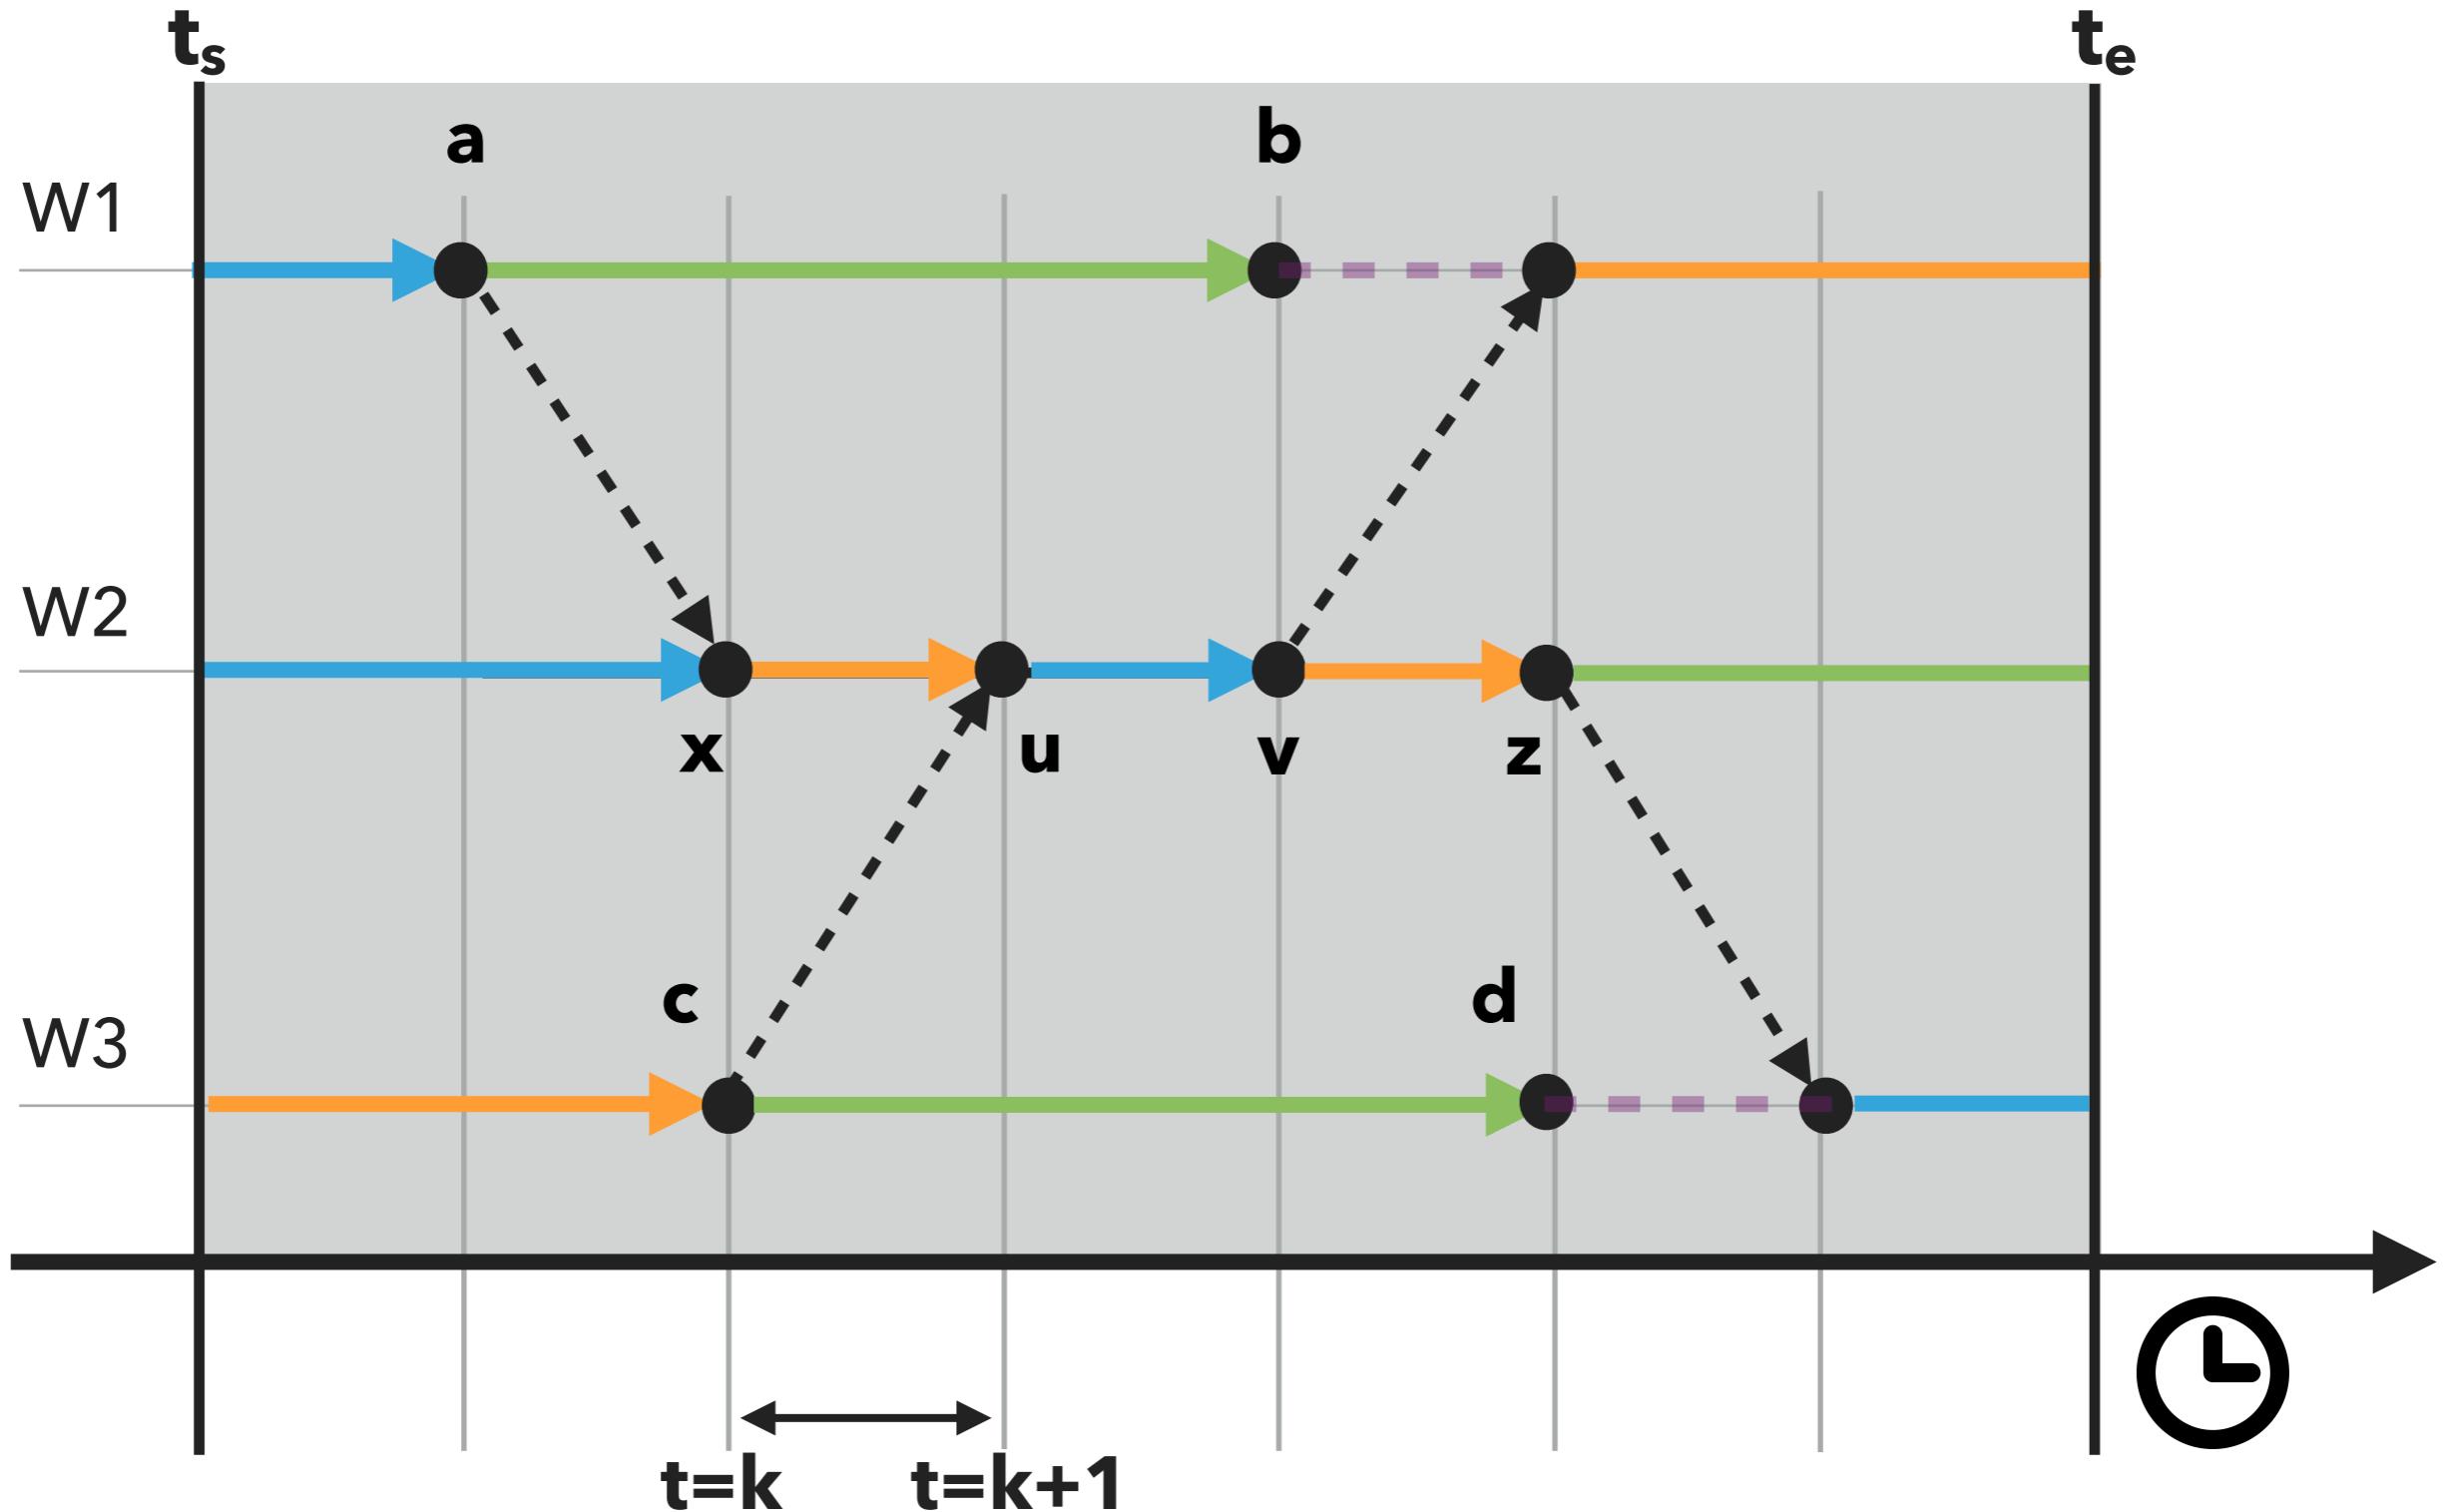
firstname.lastname@inf.ethz.ch

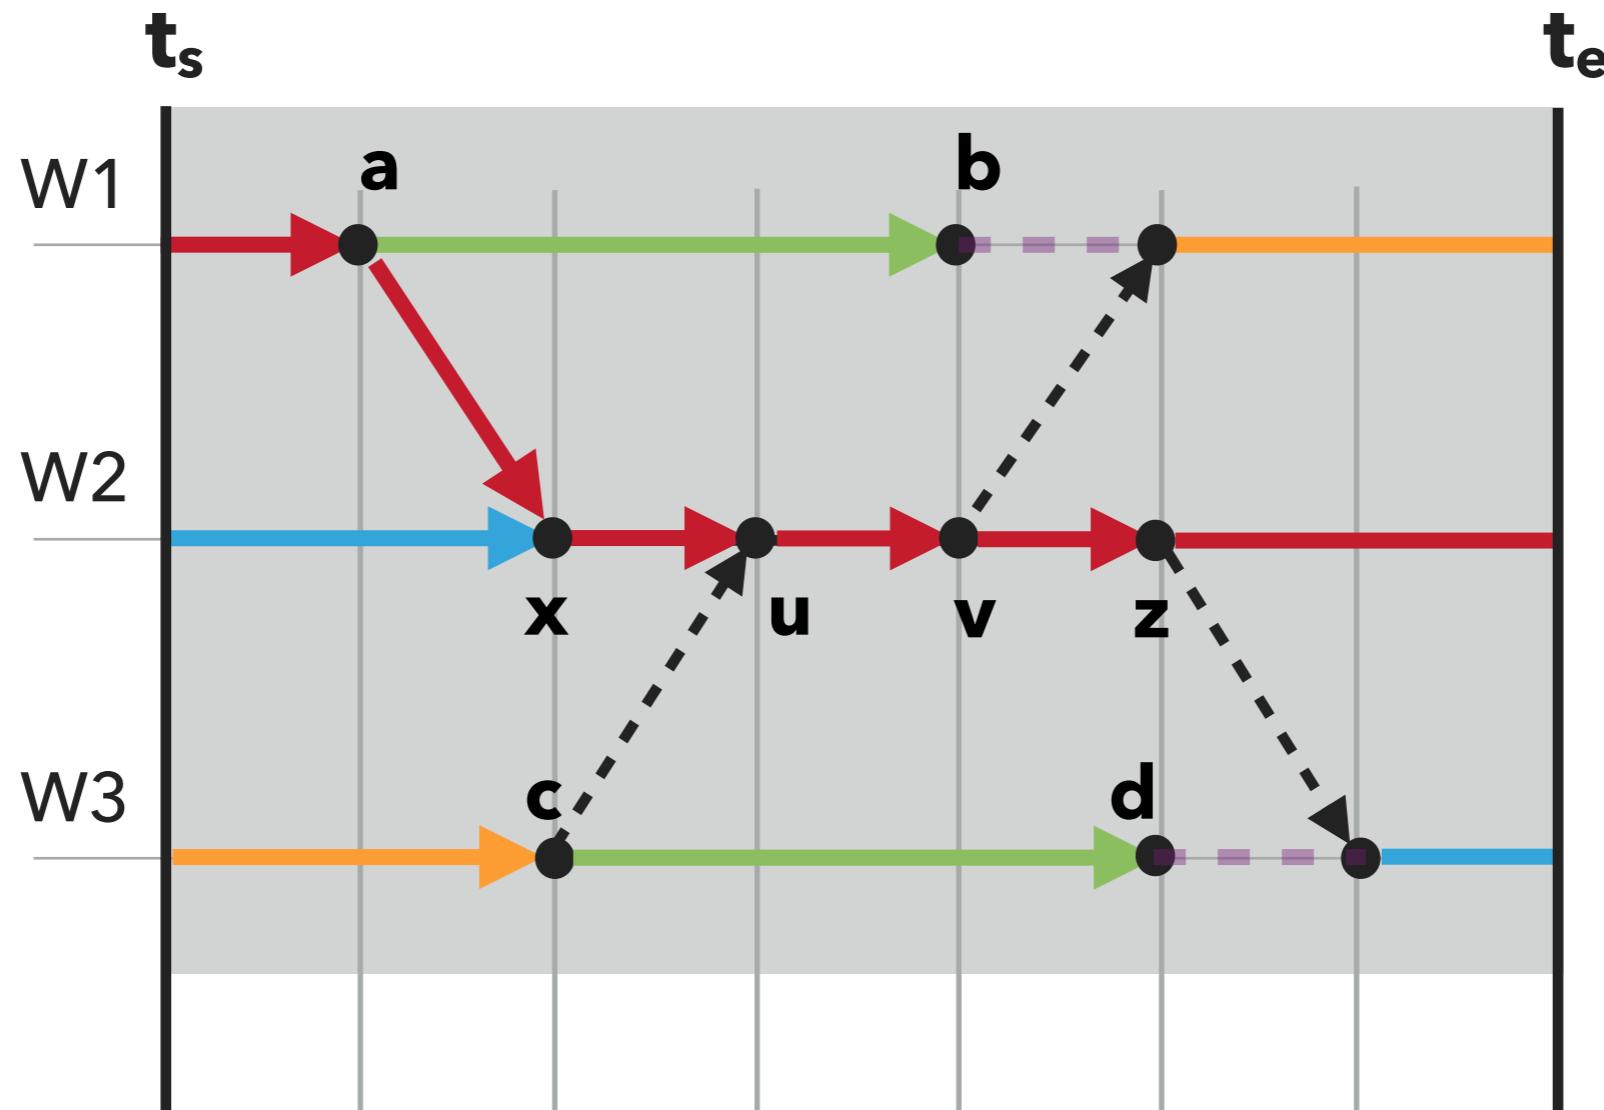
ONLINE CRITICAL PATH ANALYSIS

ONLINE ANALYSIS OF TRACE SNAPSHOTS

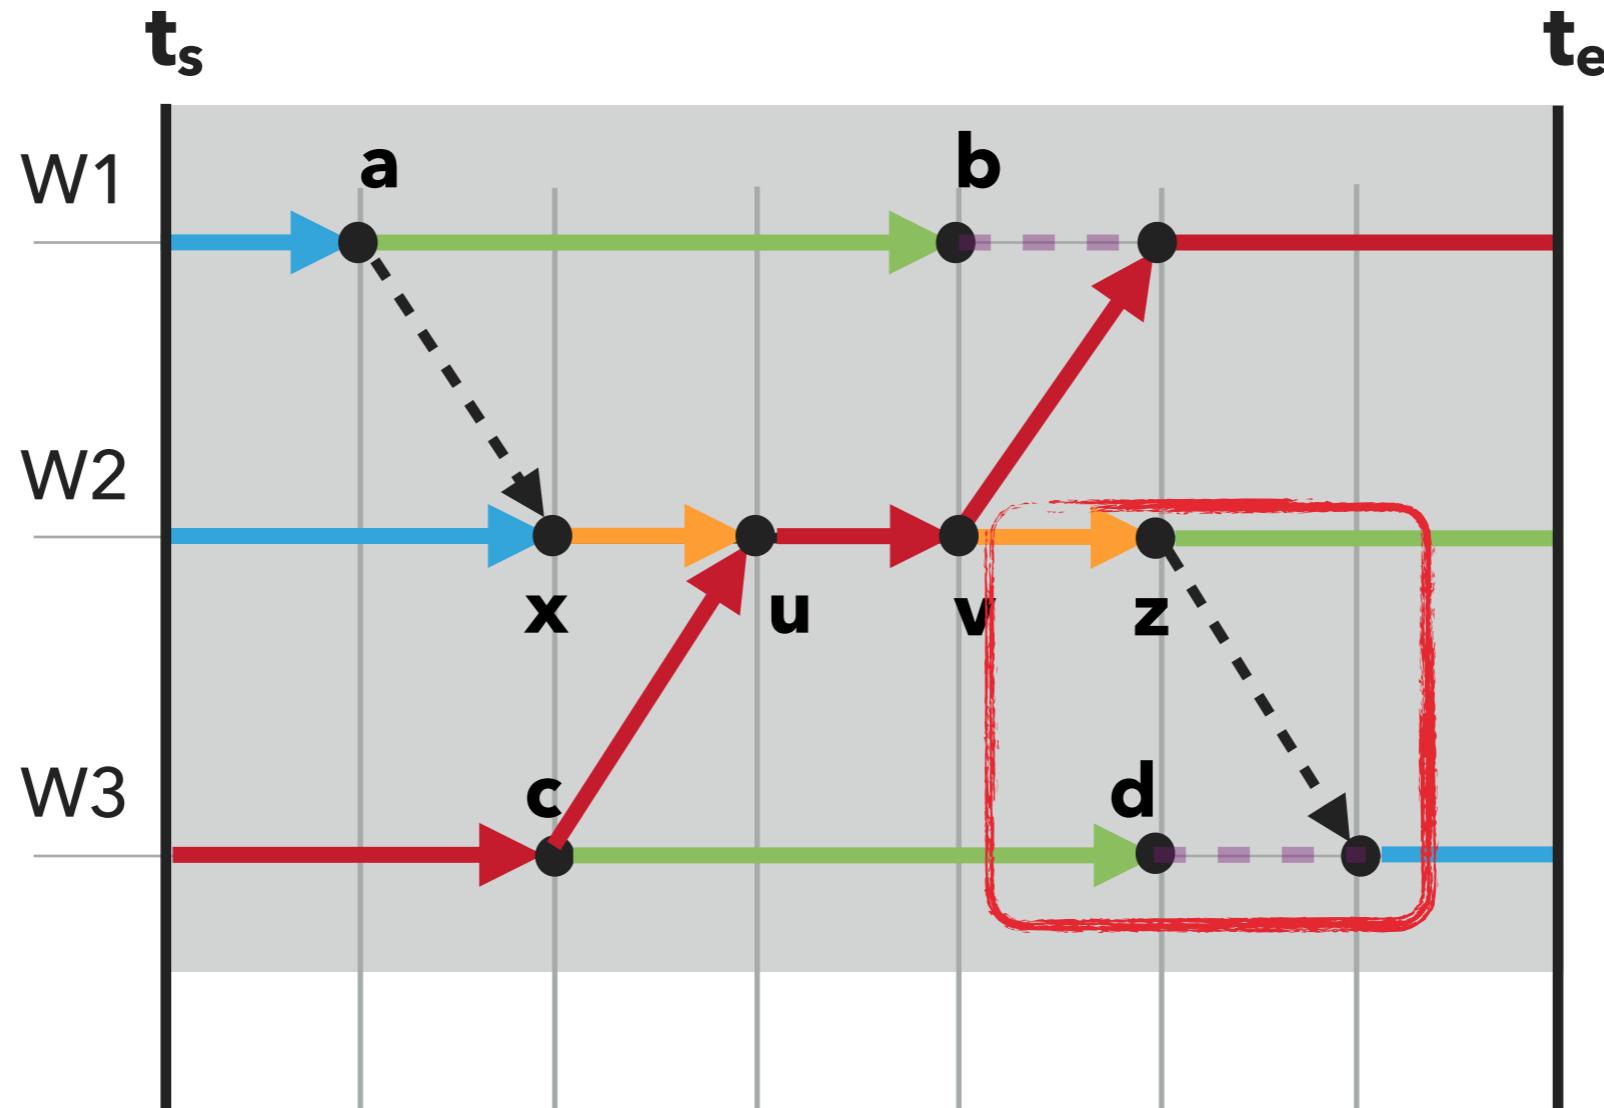


PROGRAM ACTIVITY GRAPH SNAPSHOT

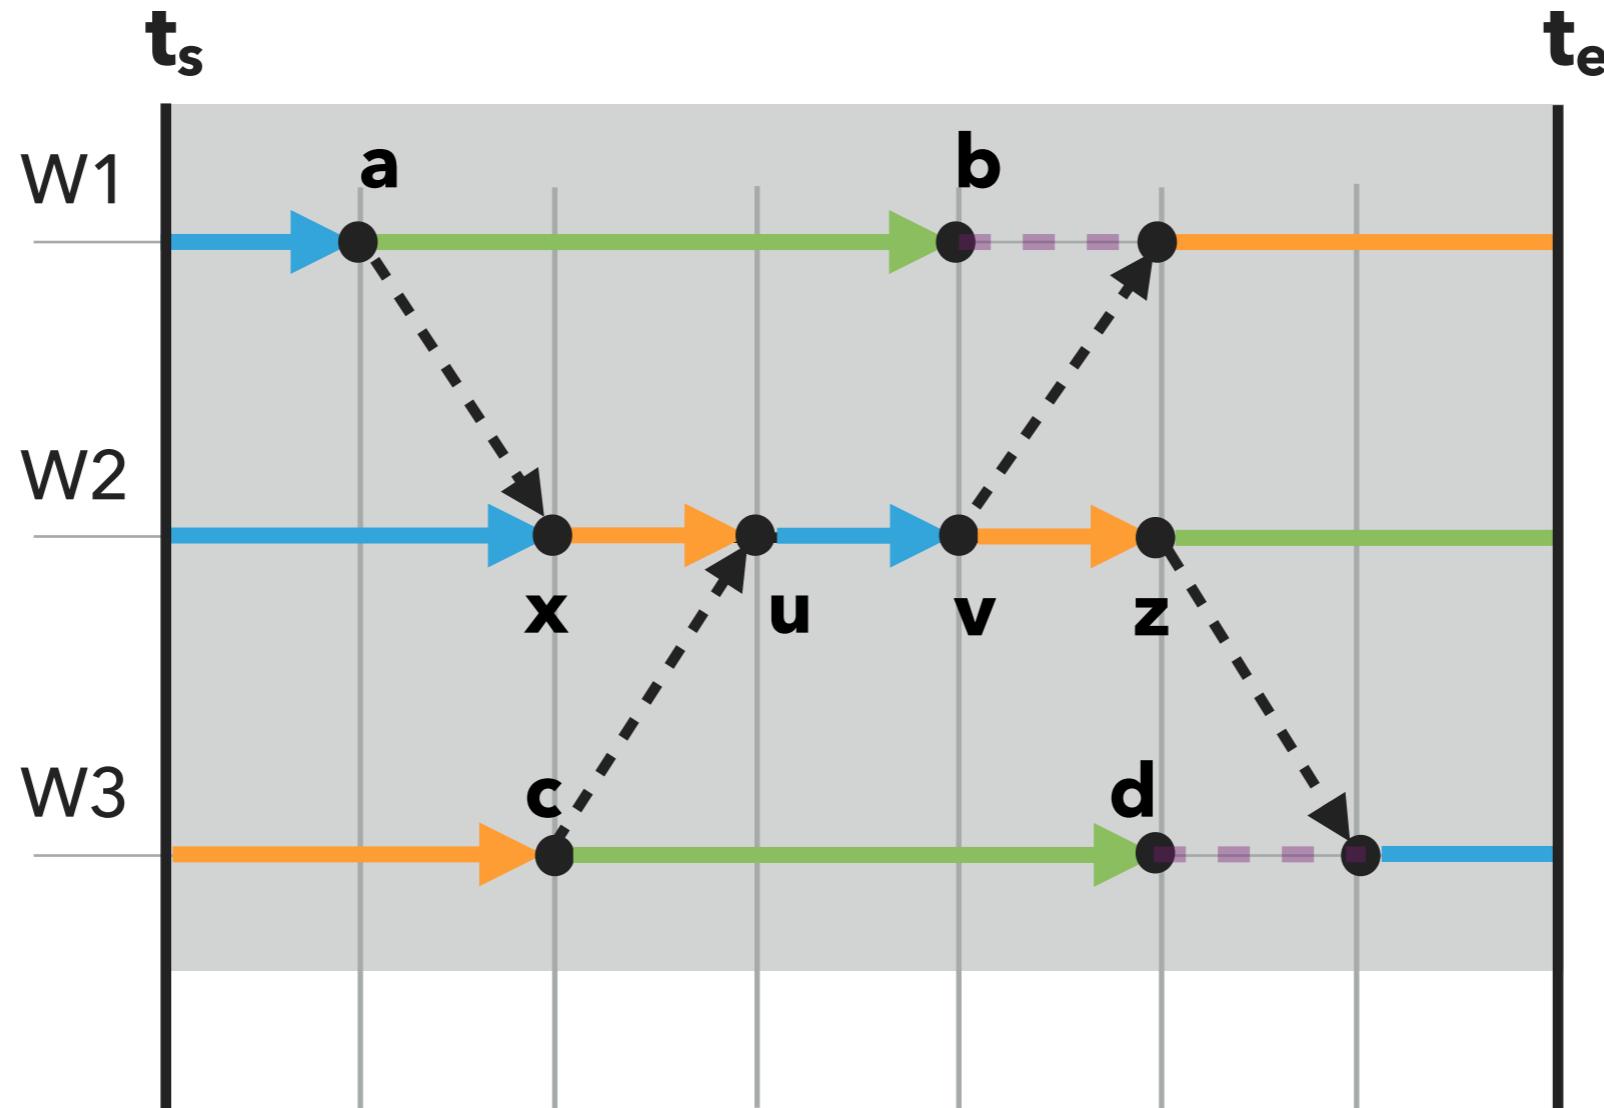




- ▶ All paths have the same length: $t_e - t_s$

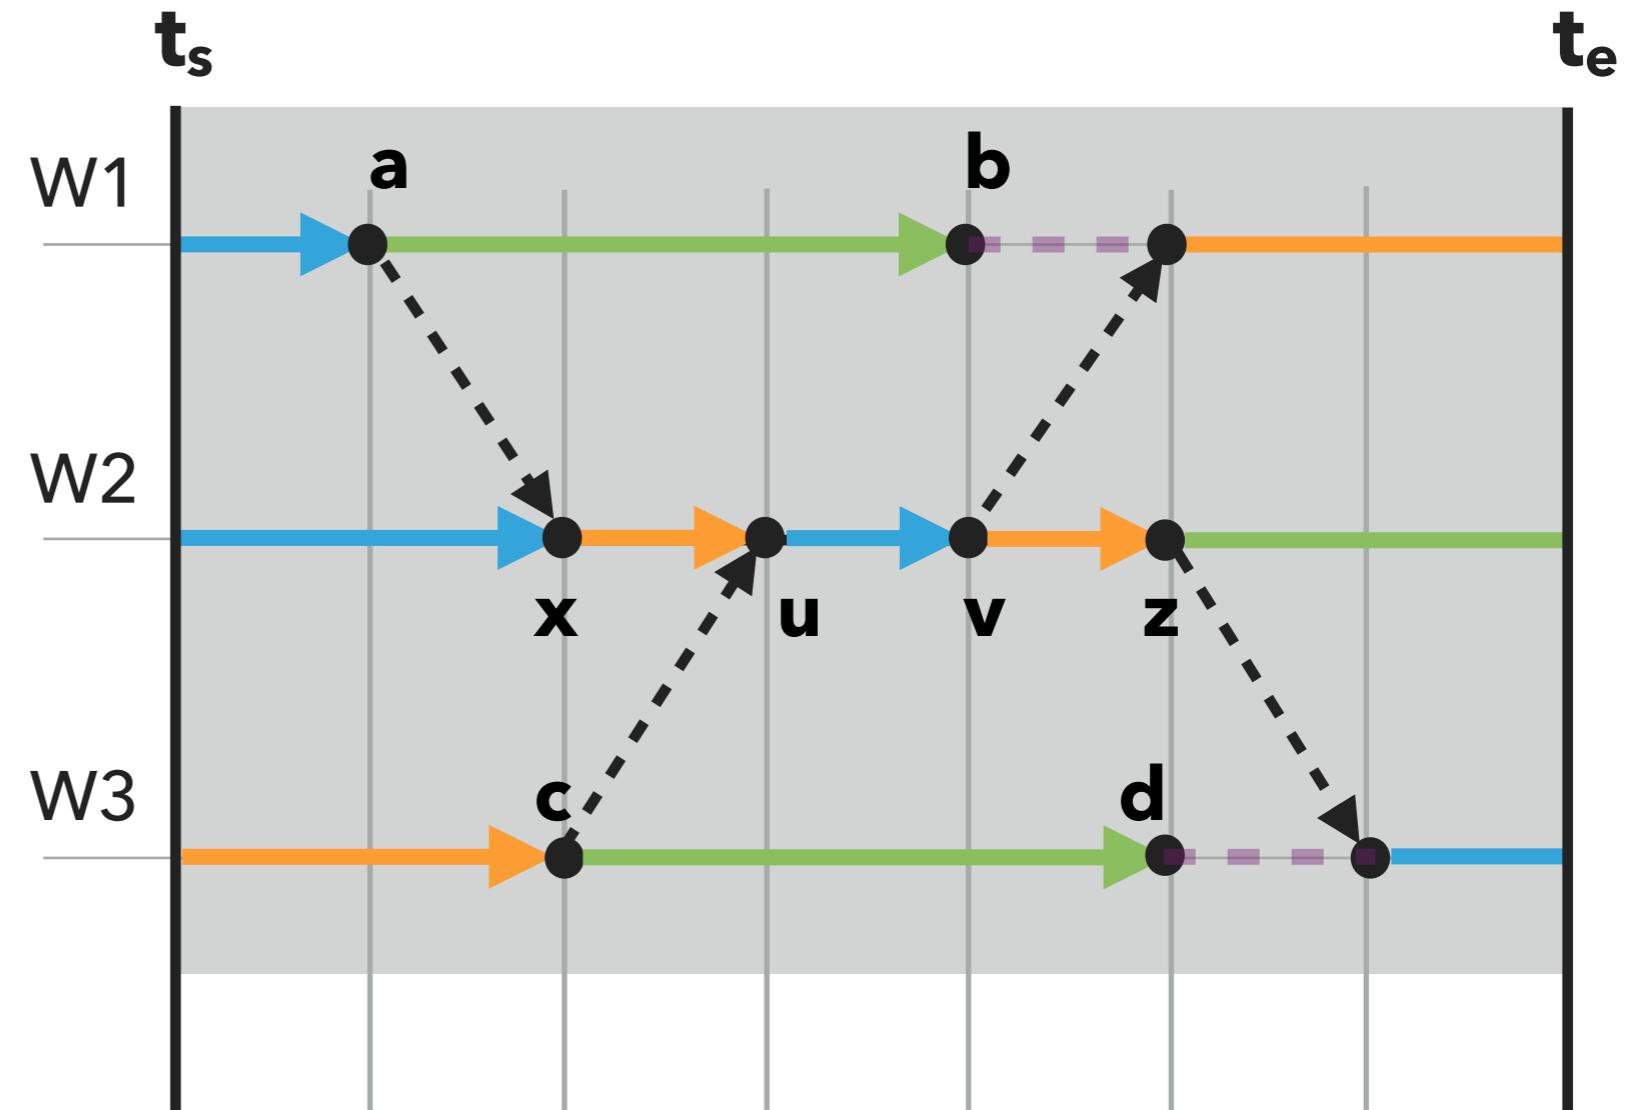


- ▶ All paths have the same length: $t_e - t_s$
- ▶ Choosing a random path might miss critical activities



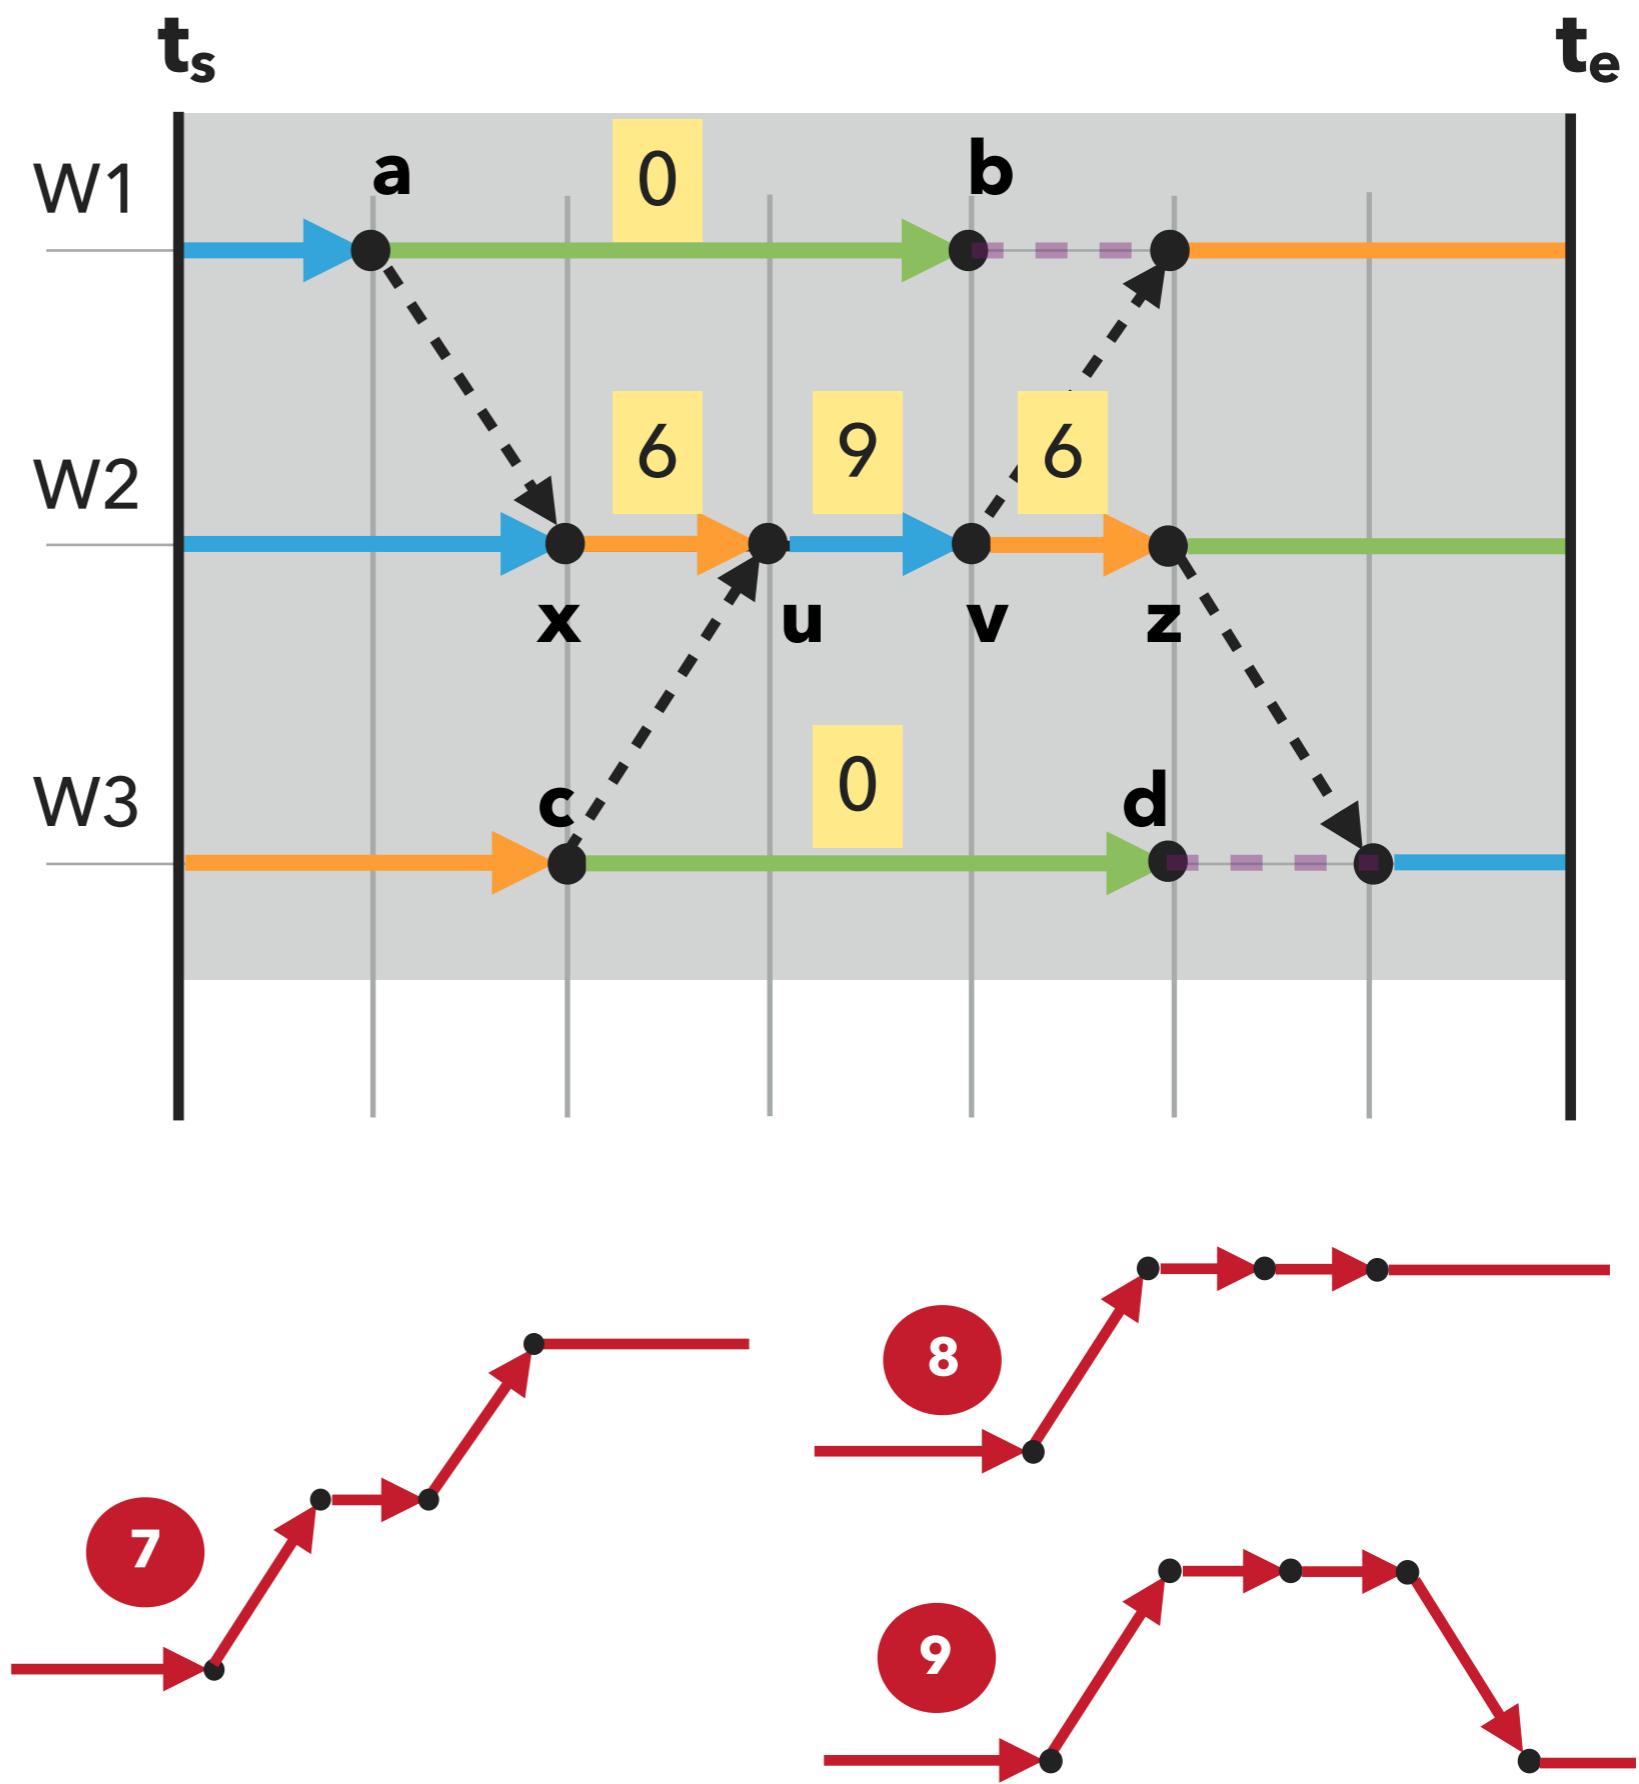
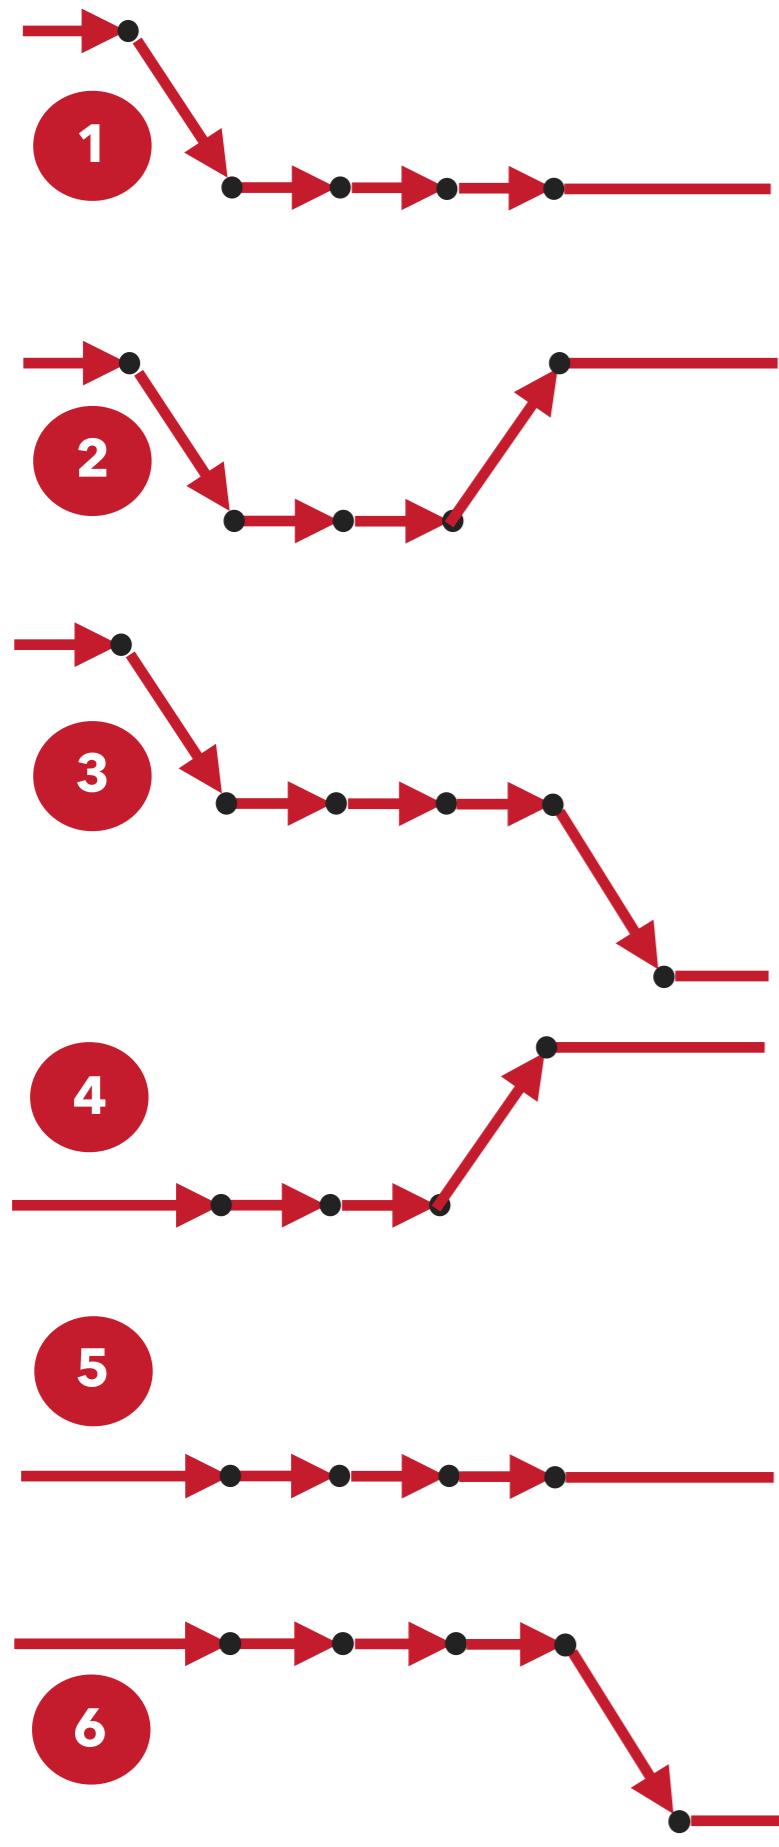
- ▶ All paths have the same length: $t_e - t_s$
- ▶ Choosing a random path might miss critical activities
- ▶ Enumerating all paths is impractical

All paths are potentially part of the **evolving critical path**



How to **rank** activities with regard to **criticality**?

Intuition: the more paths an activity appears on the more probable it is that this activity is critical



CRITICAL PARTICIPATION (CP METRIC)

An estimation of the activity's participation in the critical path

centrality: the number of paths this activity appears on

$$CP_a = \frac{C(a) \cdot a_w}{N(t_e - t_s)}$$

total number of paths in the snapshot

activity duration: edge weight

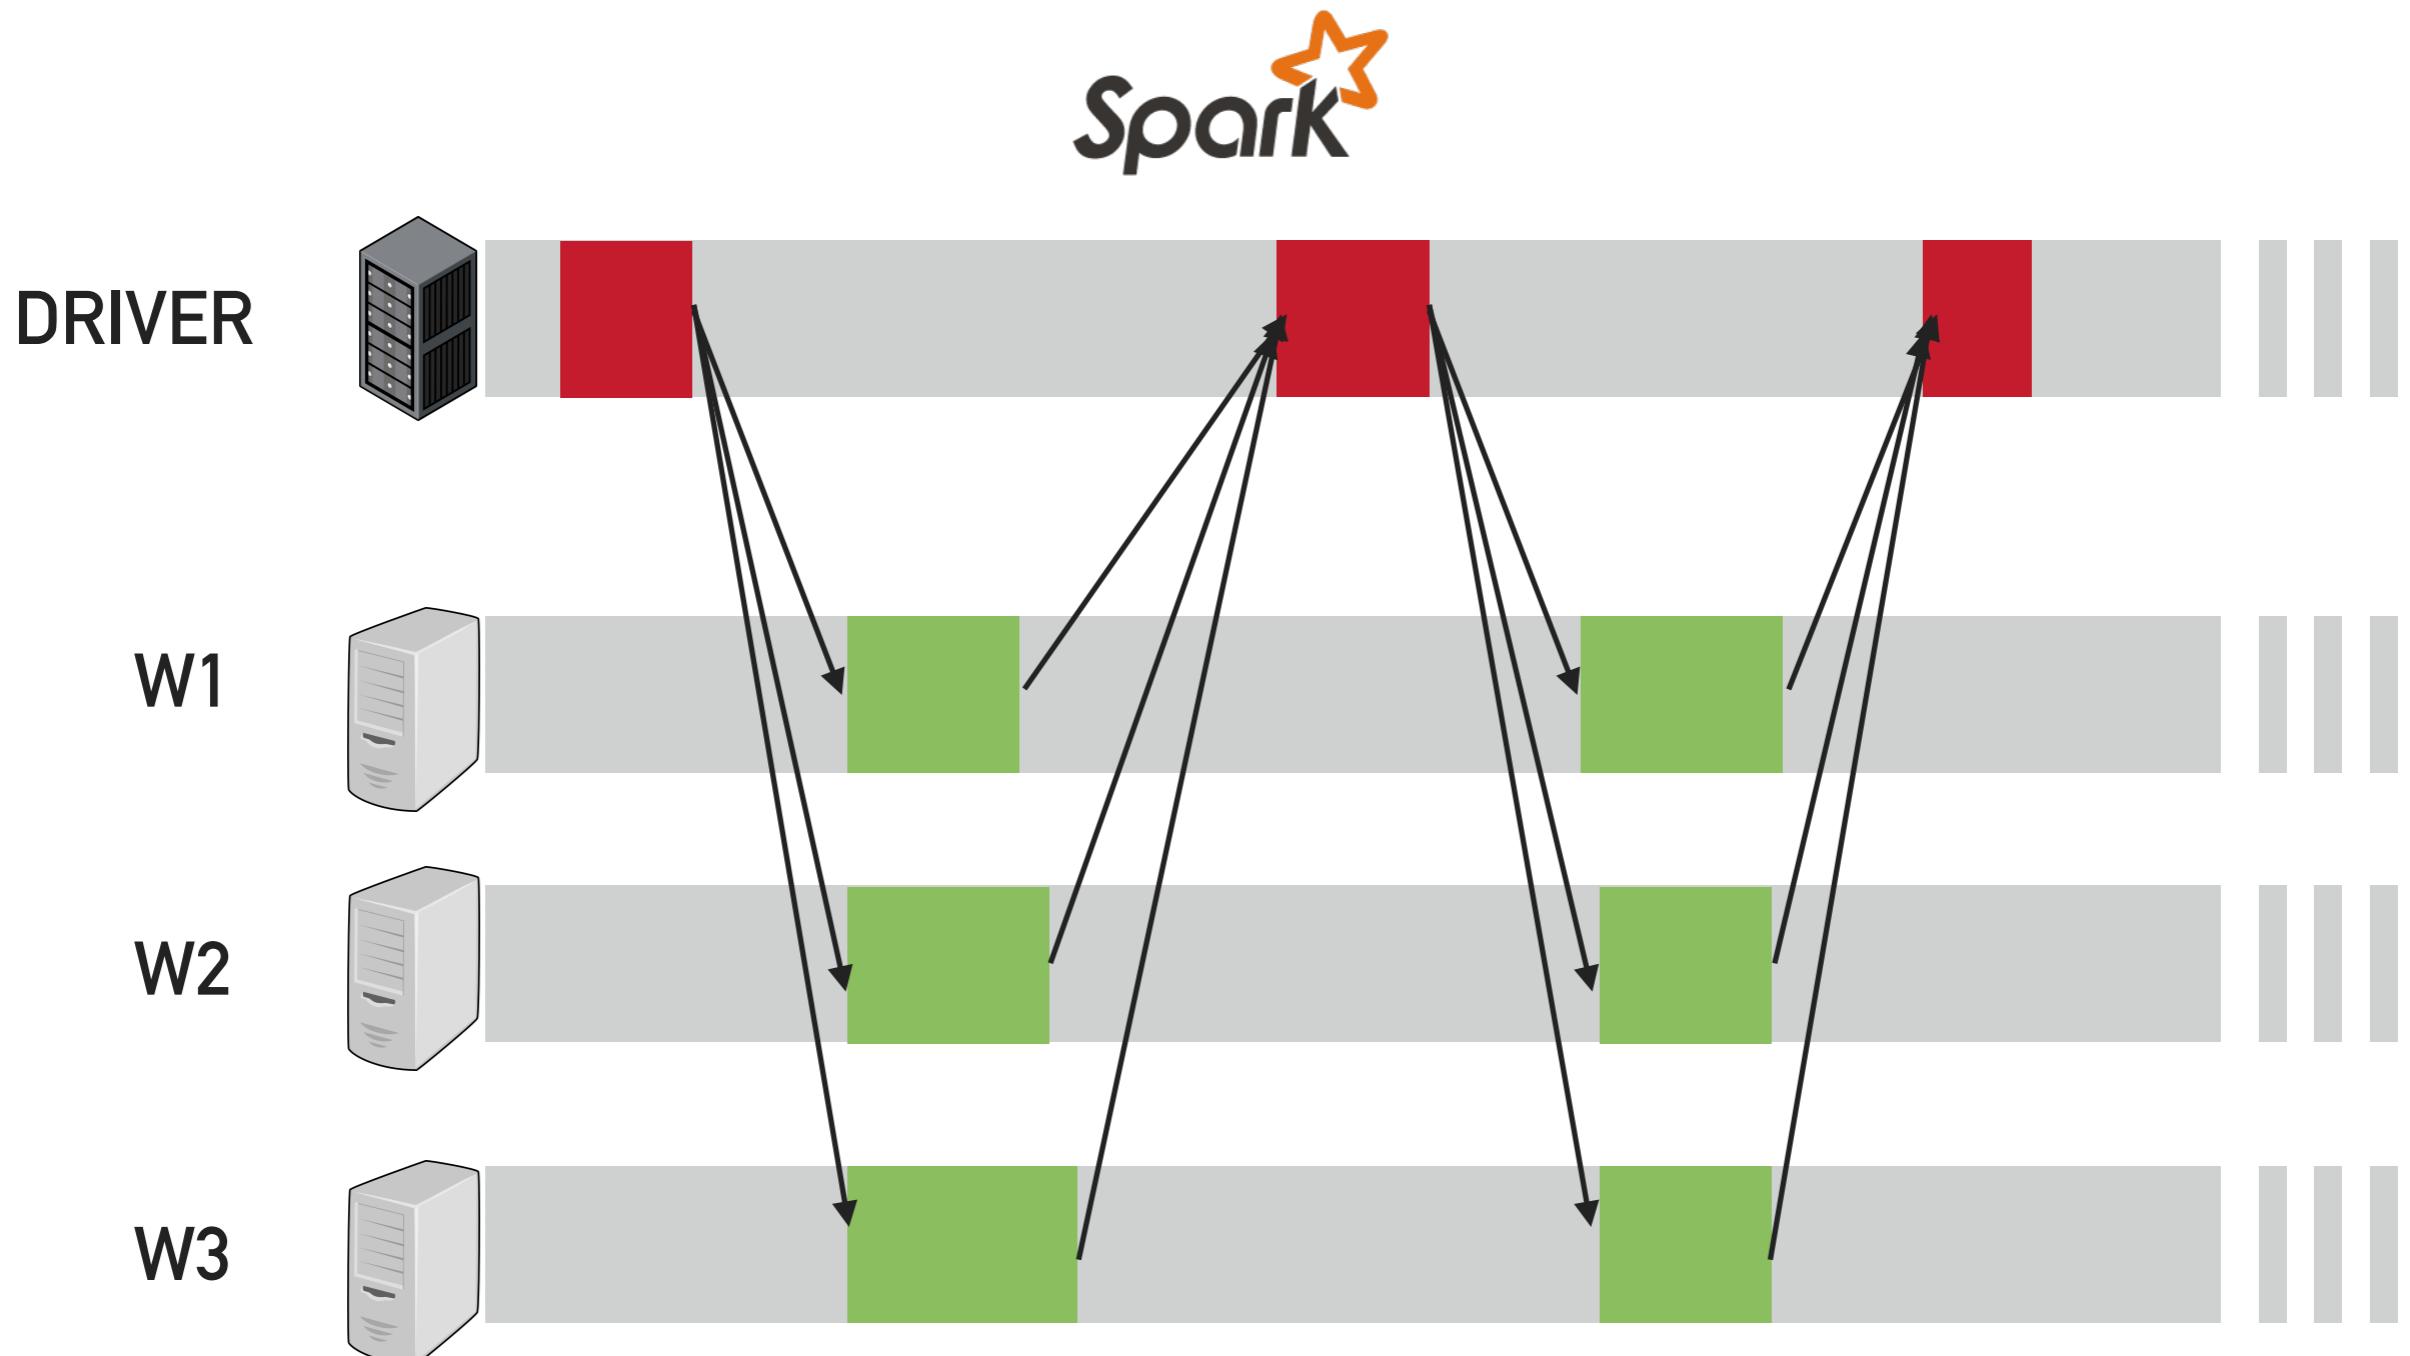
path length
(identical for all paths)

Can be computed without path enumeration!



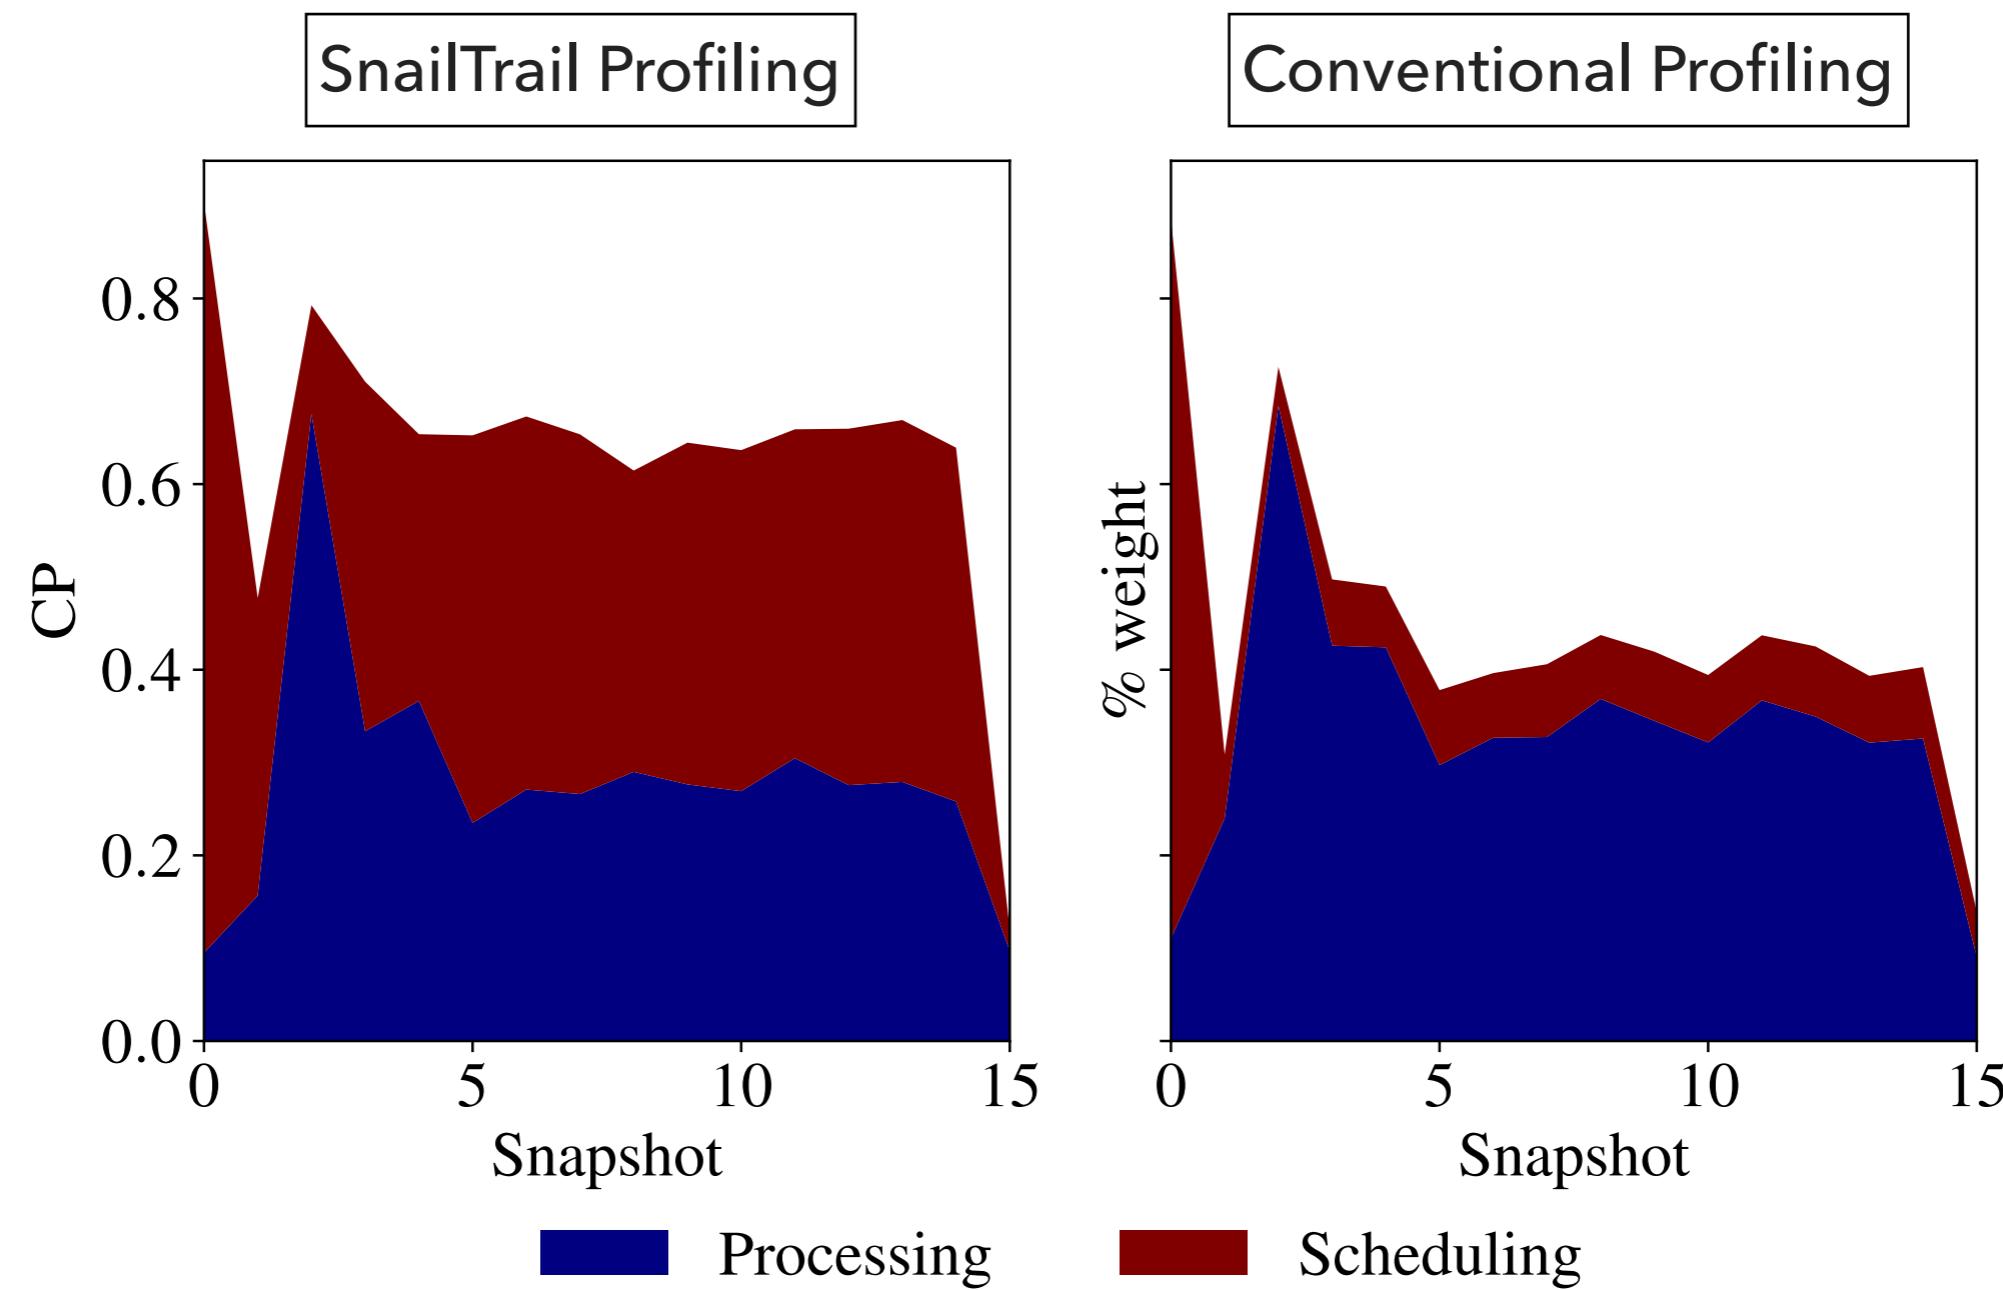
ONLINE PERFORMANCE ANALYSIS WITH **SAILTRAIL**

EXAMPLE: TASK SCHEDULING IN APACHE SPARK



Venkataraman, Shivaram, et al. "Drizzle: Fast and adaptable stream processing at scale." Spark Summit (2016).

SCEDULING BOTTLENECK IN APACHE SPARK

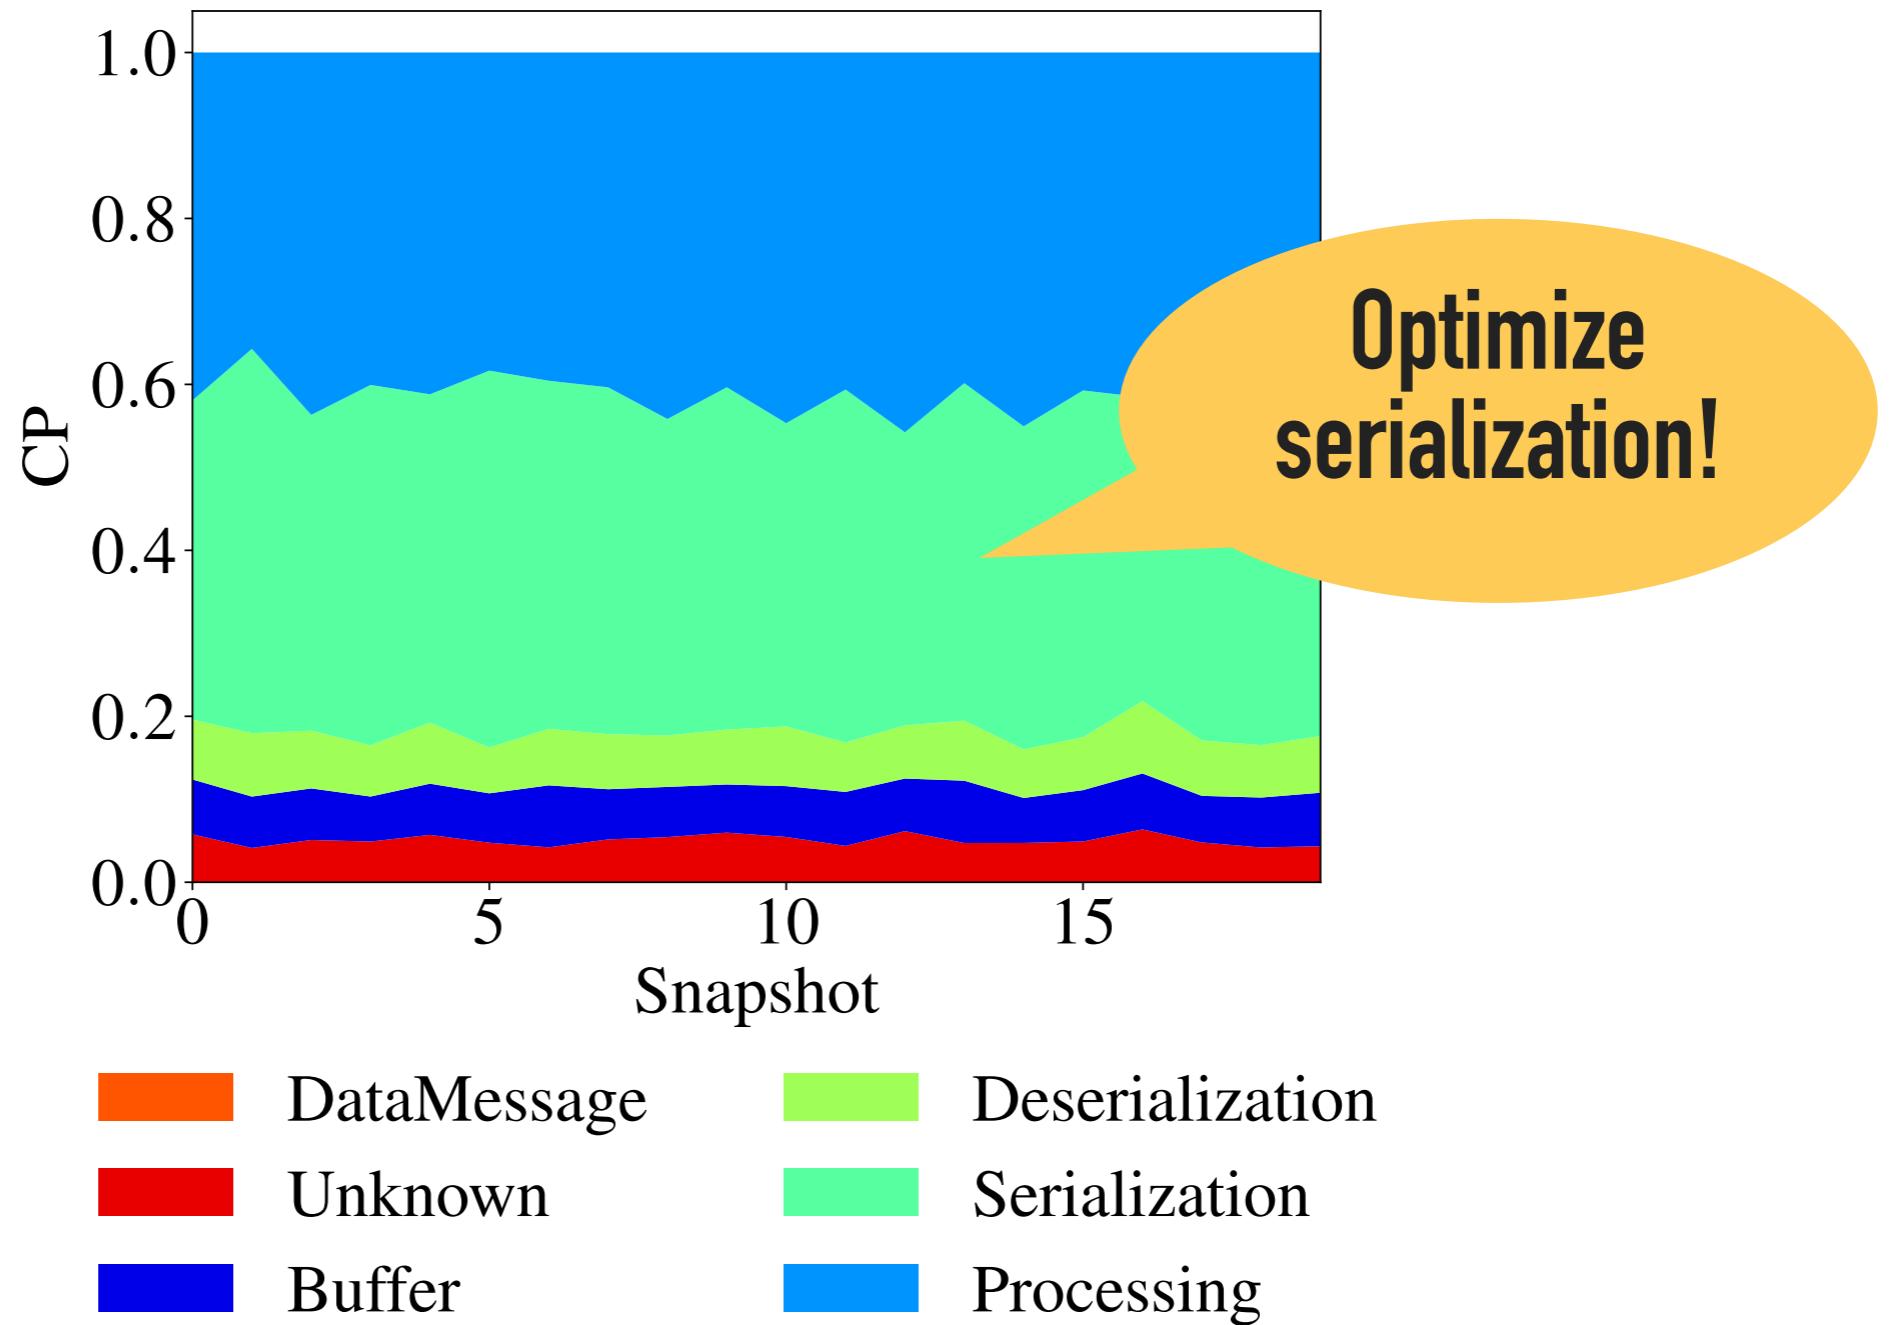


Apache Spark: Yahoo! Streaming Benchmark, 16 workers, 8s snapshots

SNAILTRAIL CP-BASED SUMMARIES

- ▶ **Activity** Summary
 - ▶ *which activity type is a bottleneck?*

Activity Summary

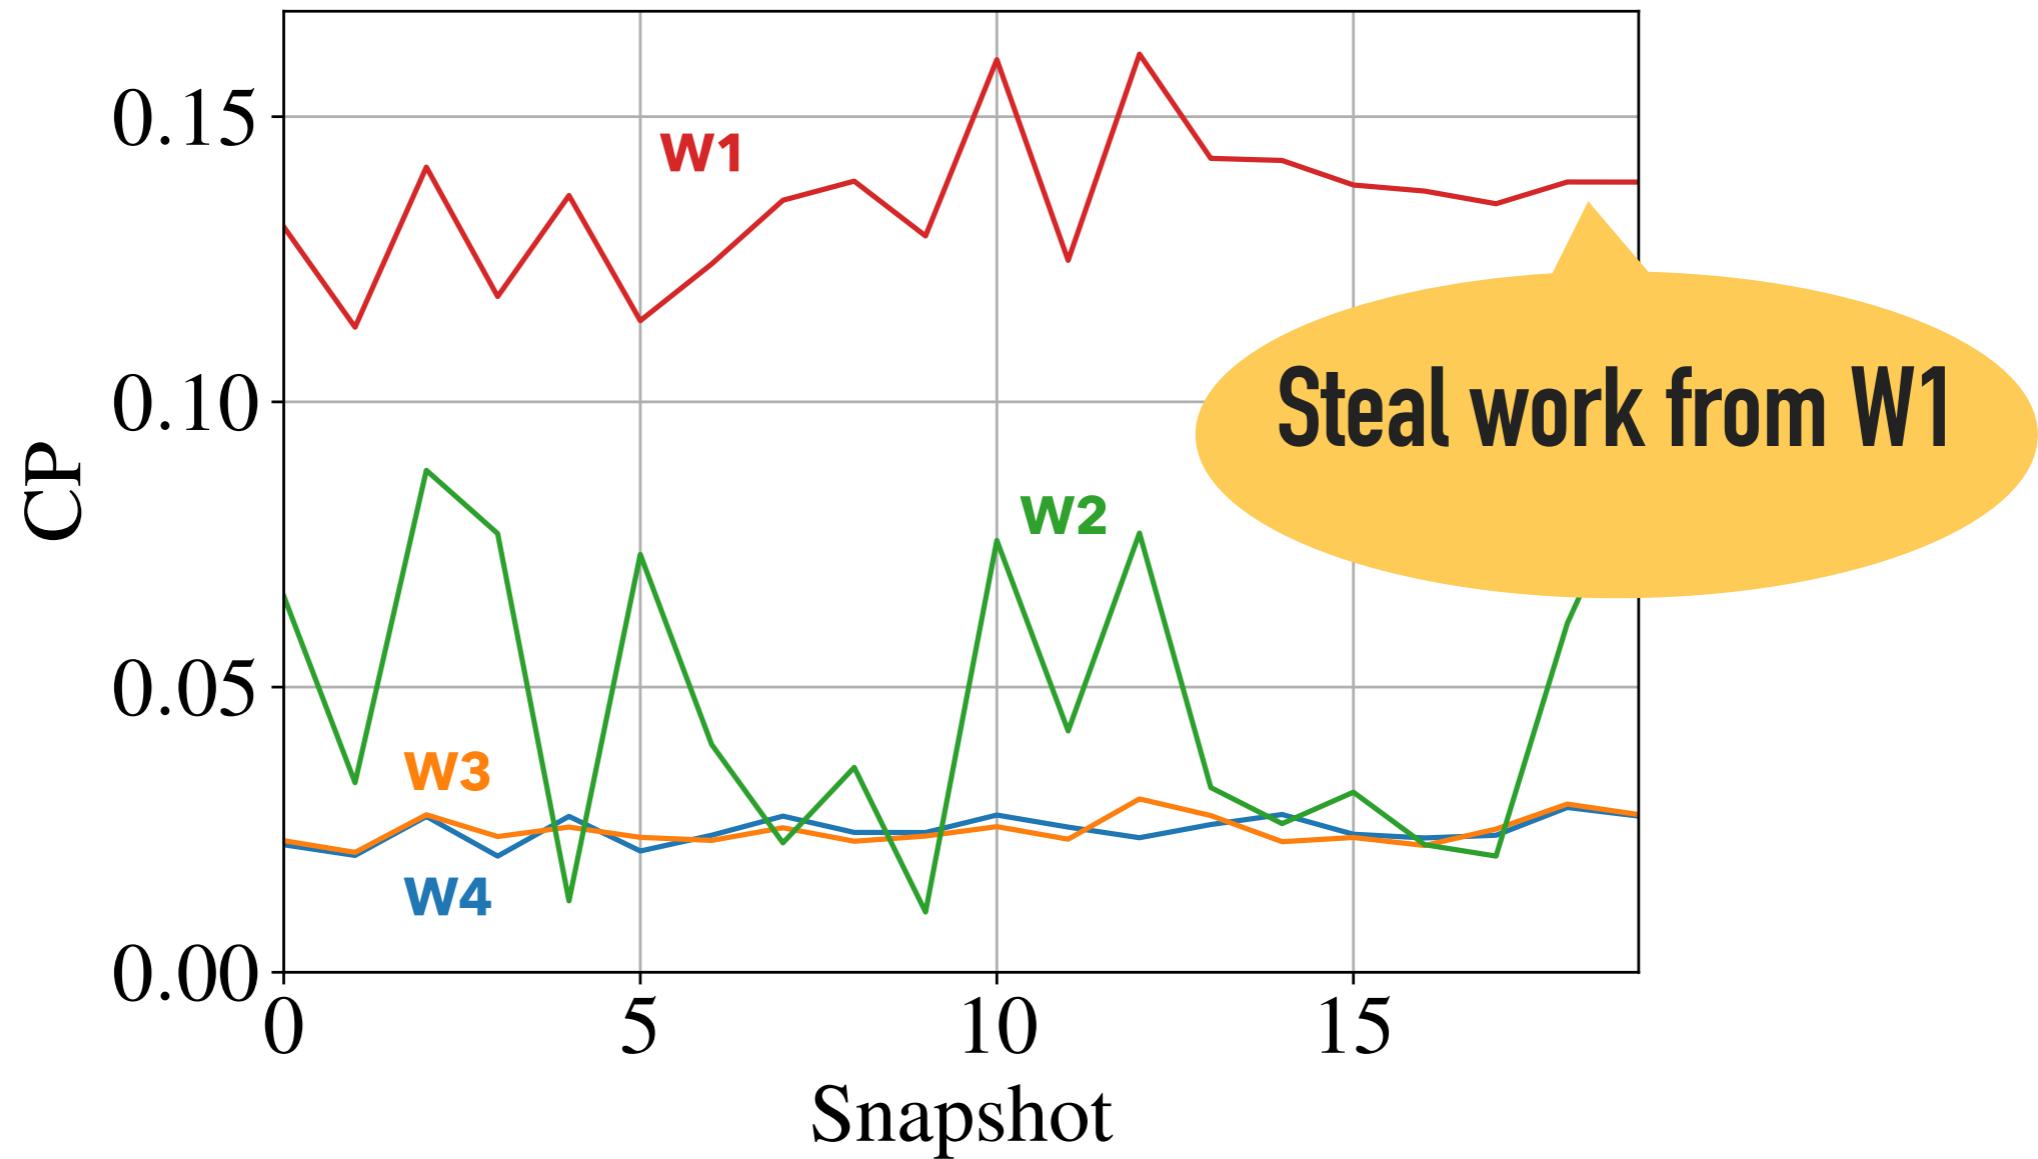


Apache Flink: Dhalion WordCount Benchmark, 4 workers, 1s snapshots

SNAILTRAIL CP-BASED SUMMARIES

- ▶ **Activity** Summary
 - ▶ *which activity type is a bottleneck?*
- ▶ **Straggler** Summary
 - ▶ *which worker is a bottleneck?*

Straggler Summary

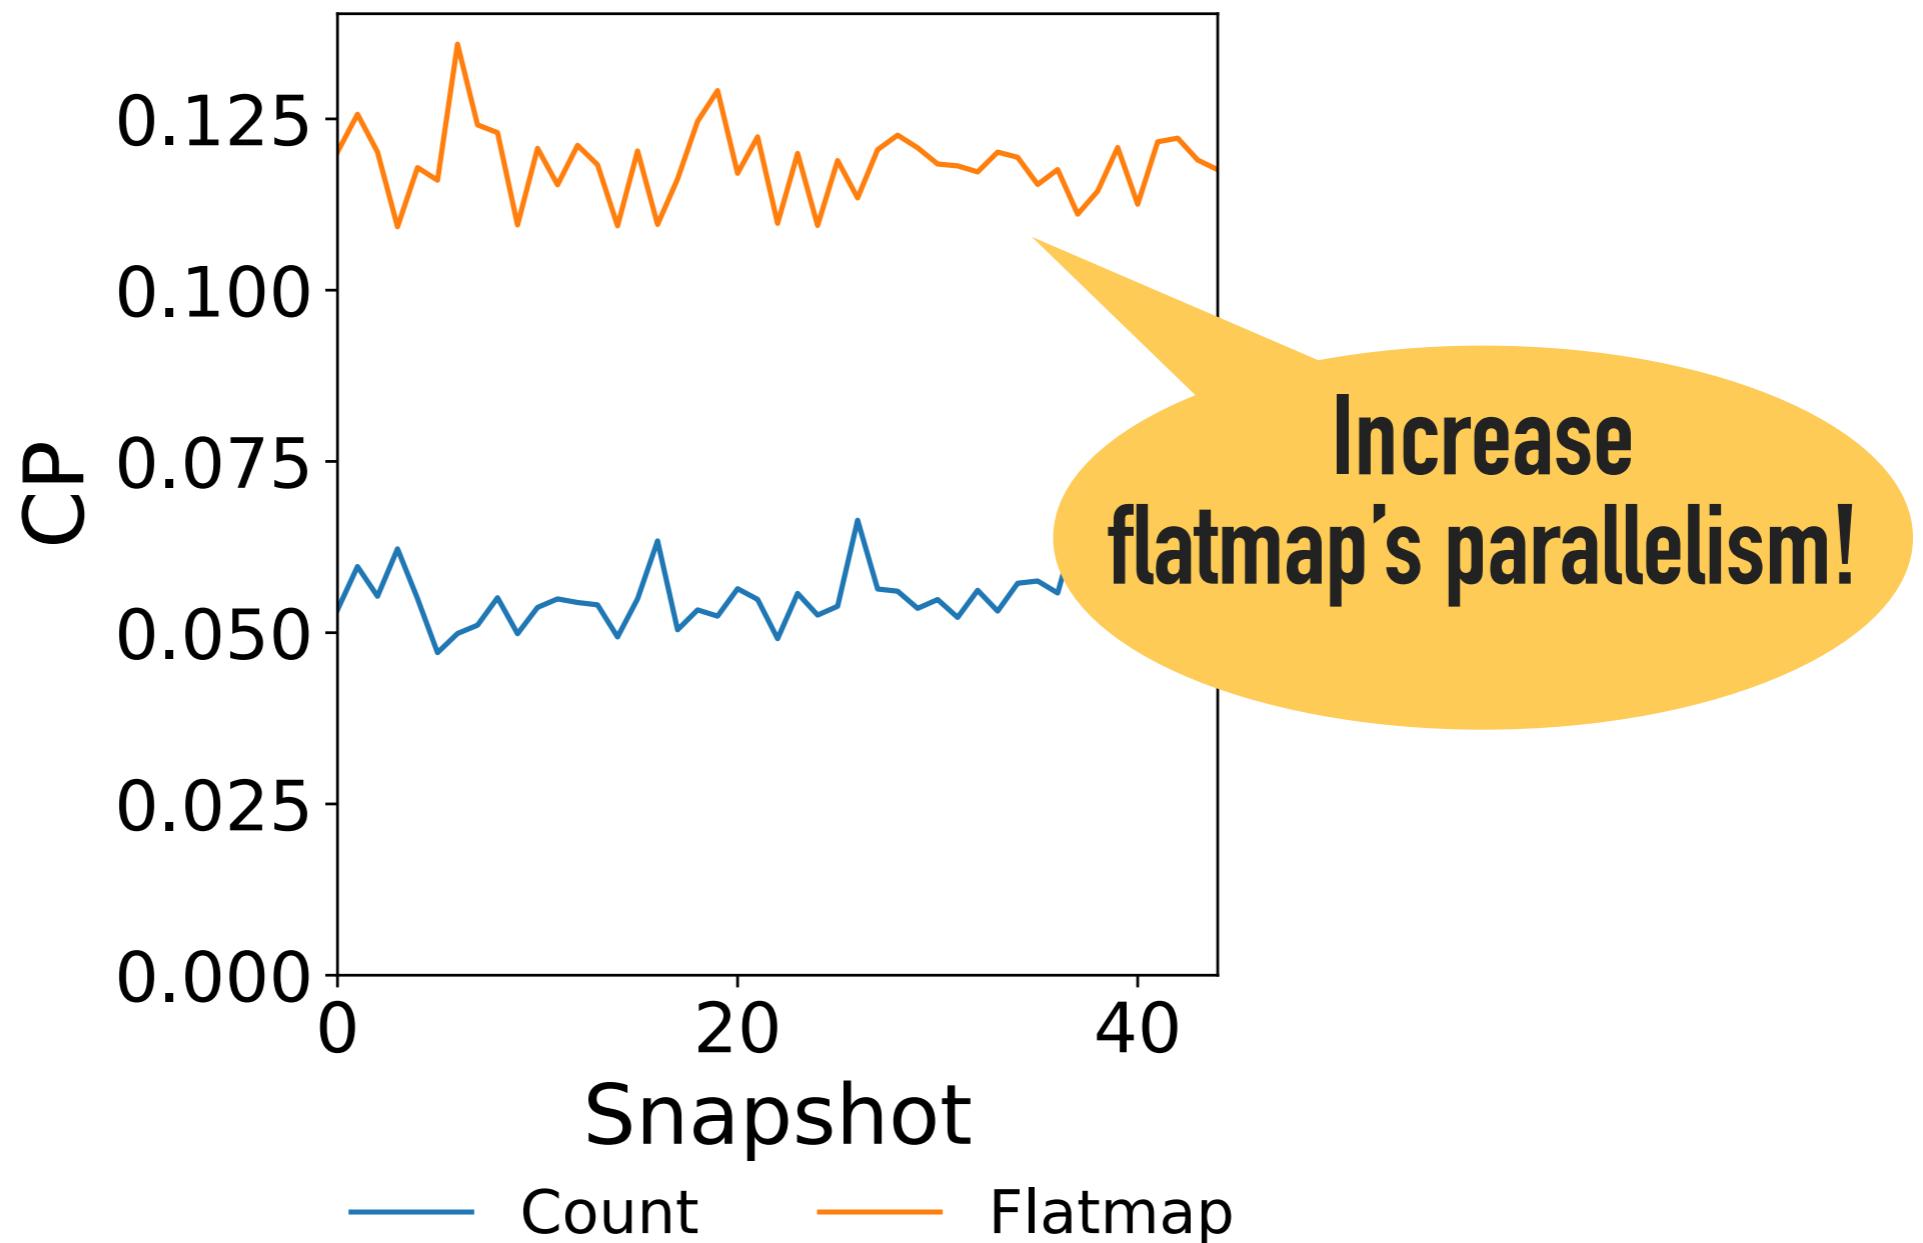


Apache Flink: Dhalion WordCount Benchmark, 4 workers, 1s snapshots

SNAILTRAIL CP-BASED SUMMARIES

- ▶ **Activity** Summary
 - ▶ *which activity type is a bottleneck?*
- ▶ **Straggler** Summary
 - ▶ *which worker is a bottleneck?*
- ▶ **Operator** Summary
 - ▶ *which operator is a bottleneck?*

Operator Summary

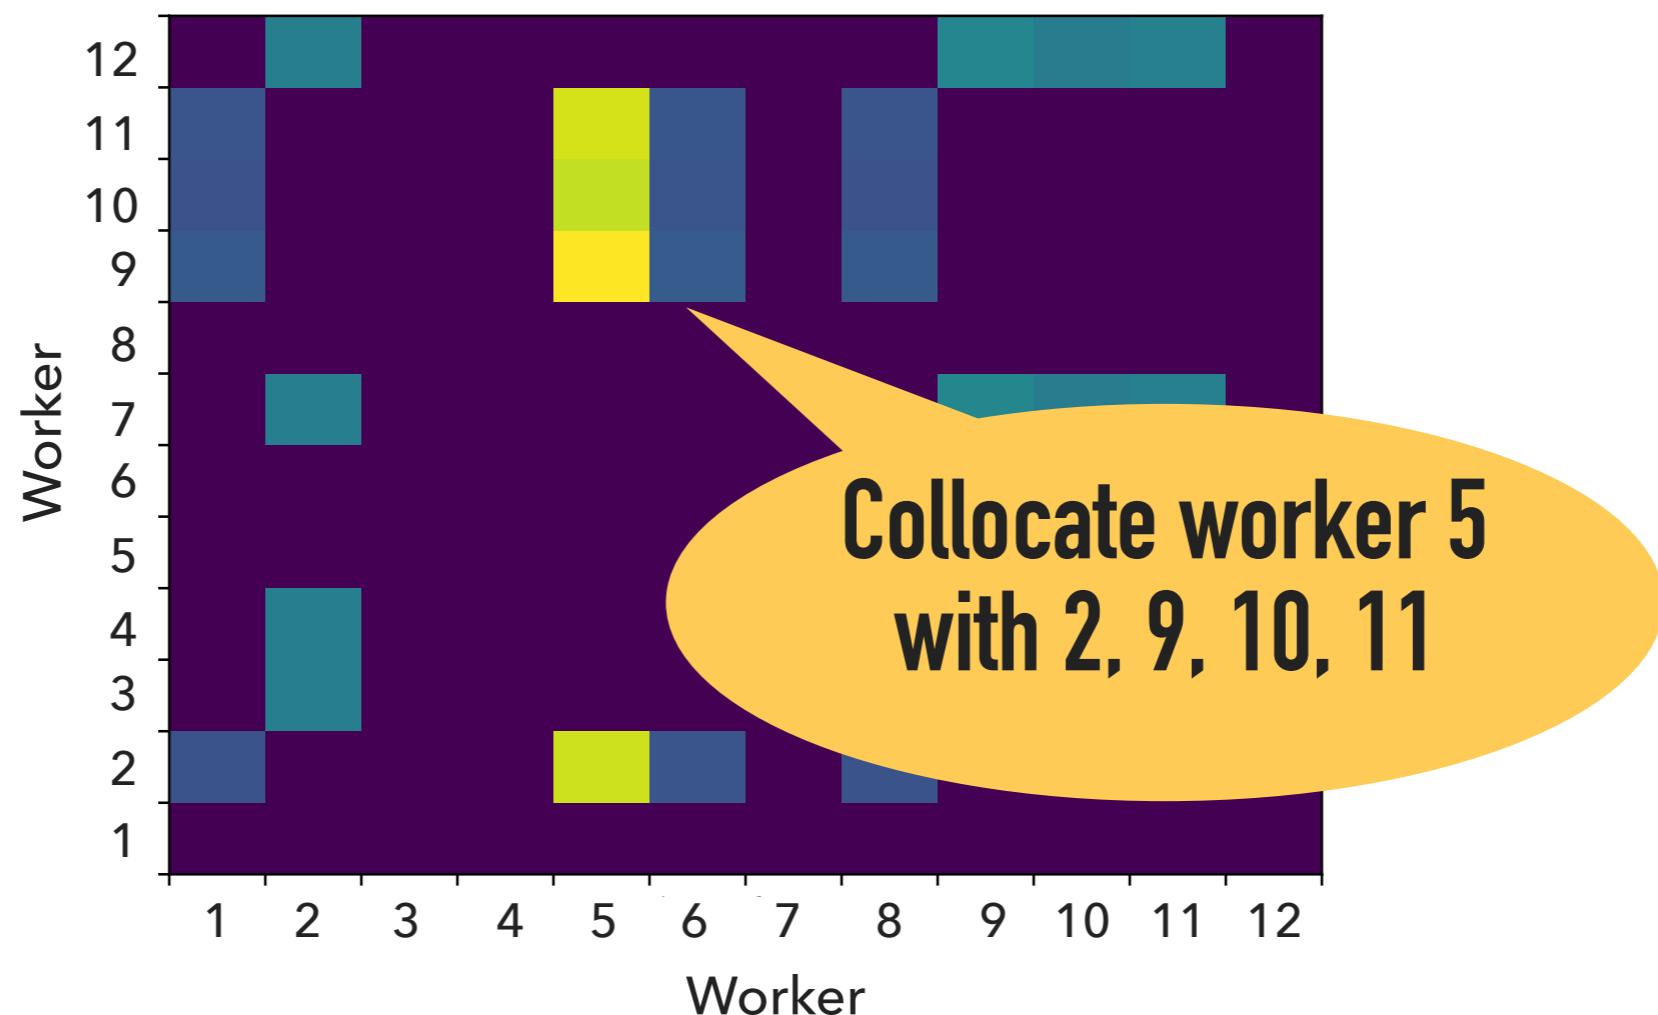


Apache Flink: Dhalion WordCount Benchmark, 10 workers, 1s snapshots

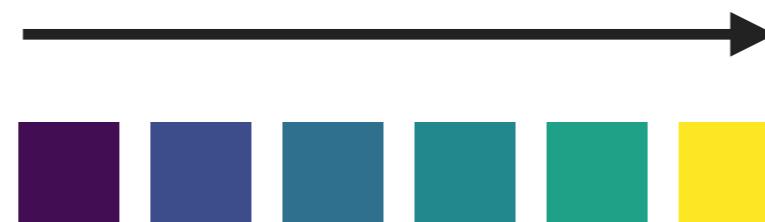
SNAILTRAIL CP-BASED SUMMARIES

- ▶ **Activity** Summary
 - ▶ *which activity type is a bottleneck?*
- ▶ **Straggler** Summary
 - ▶ *which worker is a bottleneck?*
- ▶ **Operator** Summary
 - ▶ *which operator is a bottleneck?*
- ▶ **Communication** Summary
 - ▶ *which communication channels are bottlenecks?*

Communication Summary



Communication Criticality

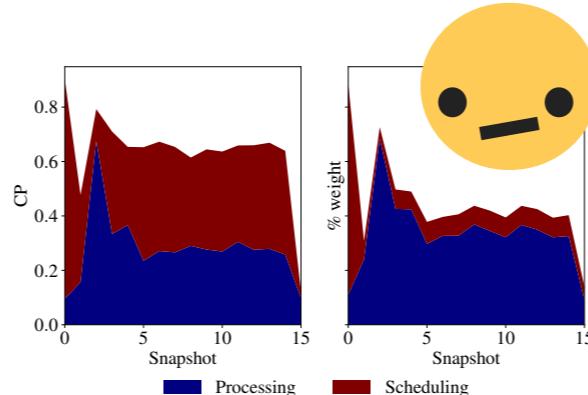


SNAILTRAIL PERFORMANCE

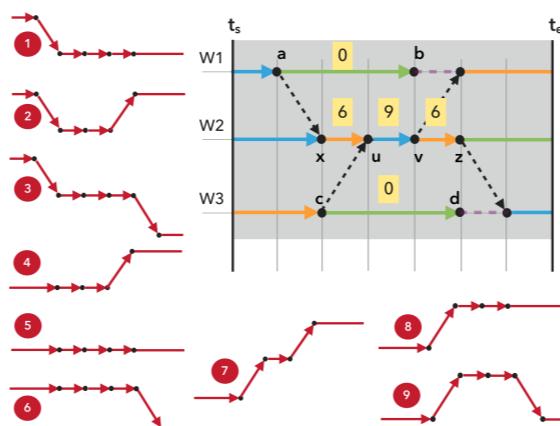
- ▶ Low instrumentation overhead
 - ▶ < 10% for all reference systems
- ▶ High throughput
 - ▶ 1.2 million events per second
- ▶ Always online
 - ▶ 1s of traces in 6ms, 256s of traces in < 25s

**SnailTrail on Intel Xeon E5-4640, 2.40GHz, 32 cores, 512GB RAM
Traces from Apache Flink Sessionization, 48 workers, 1s-256s snapshots**

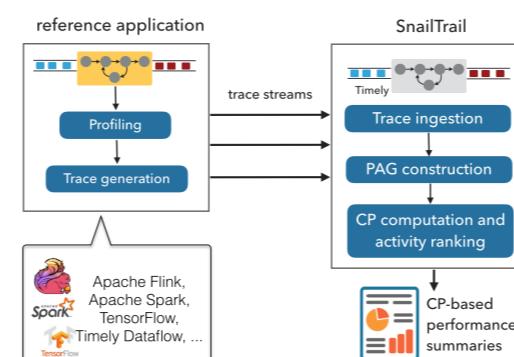
RECAP



Conventional profiling is misleading

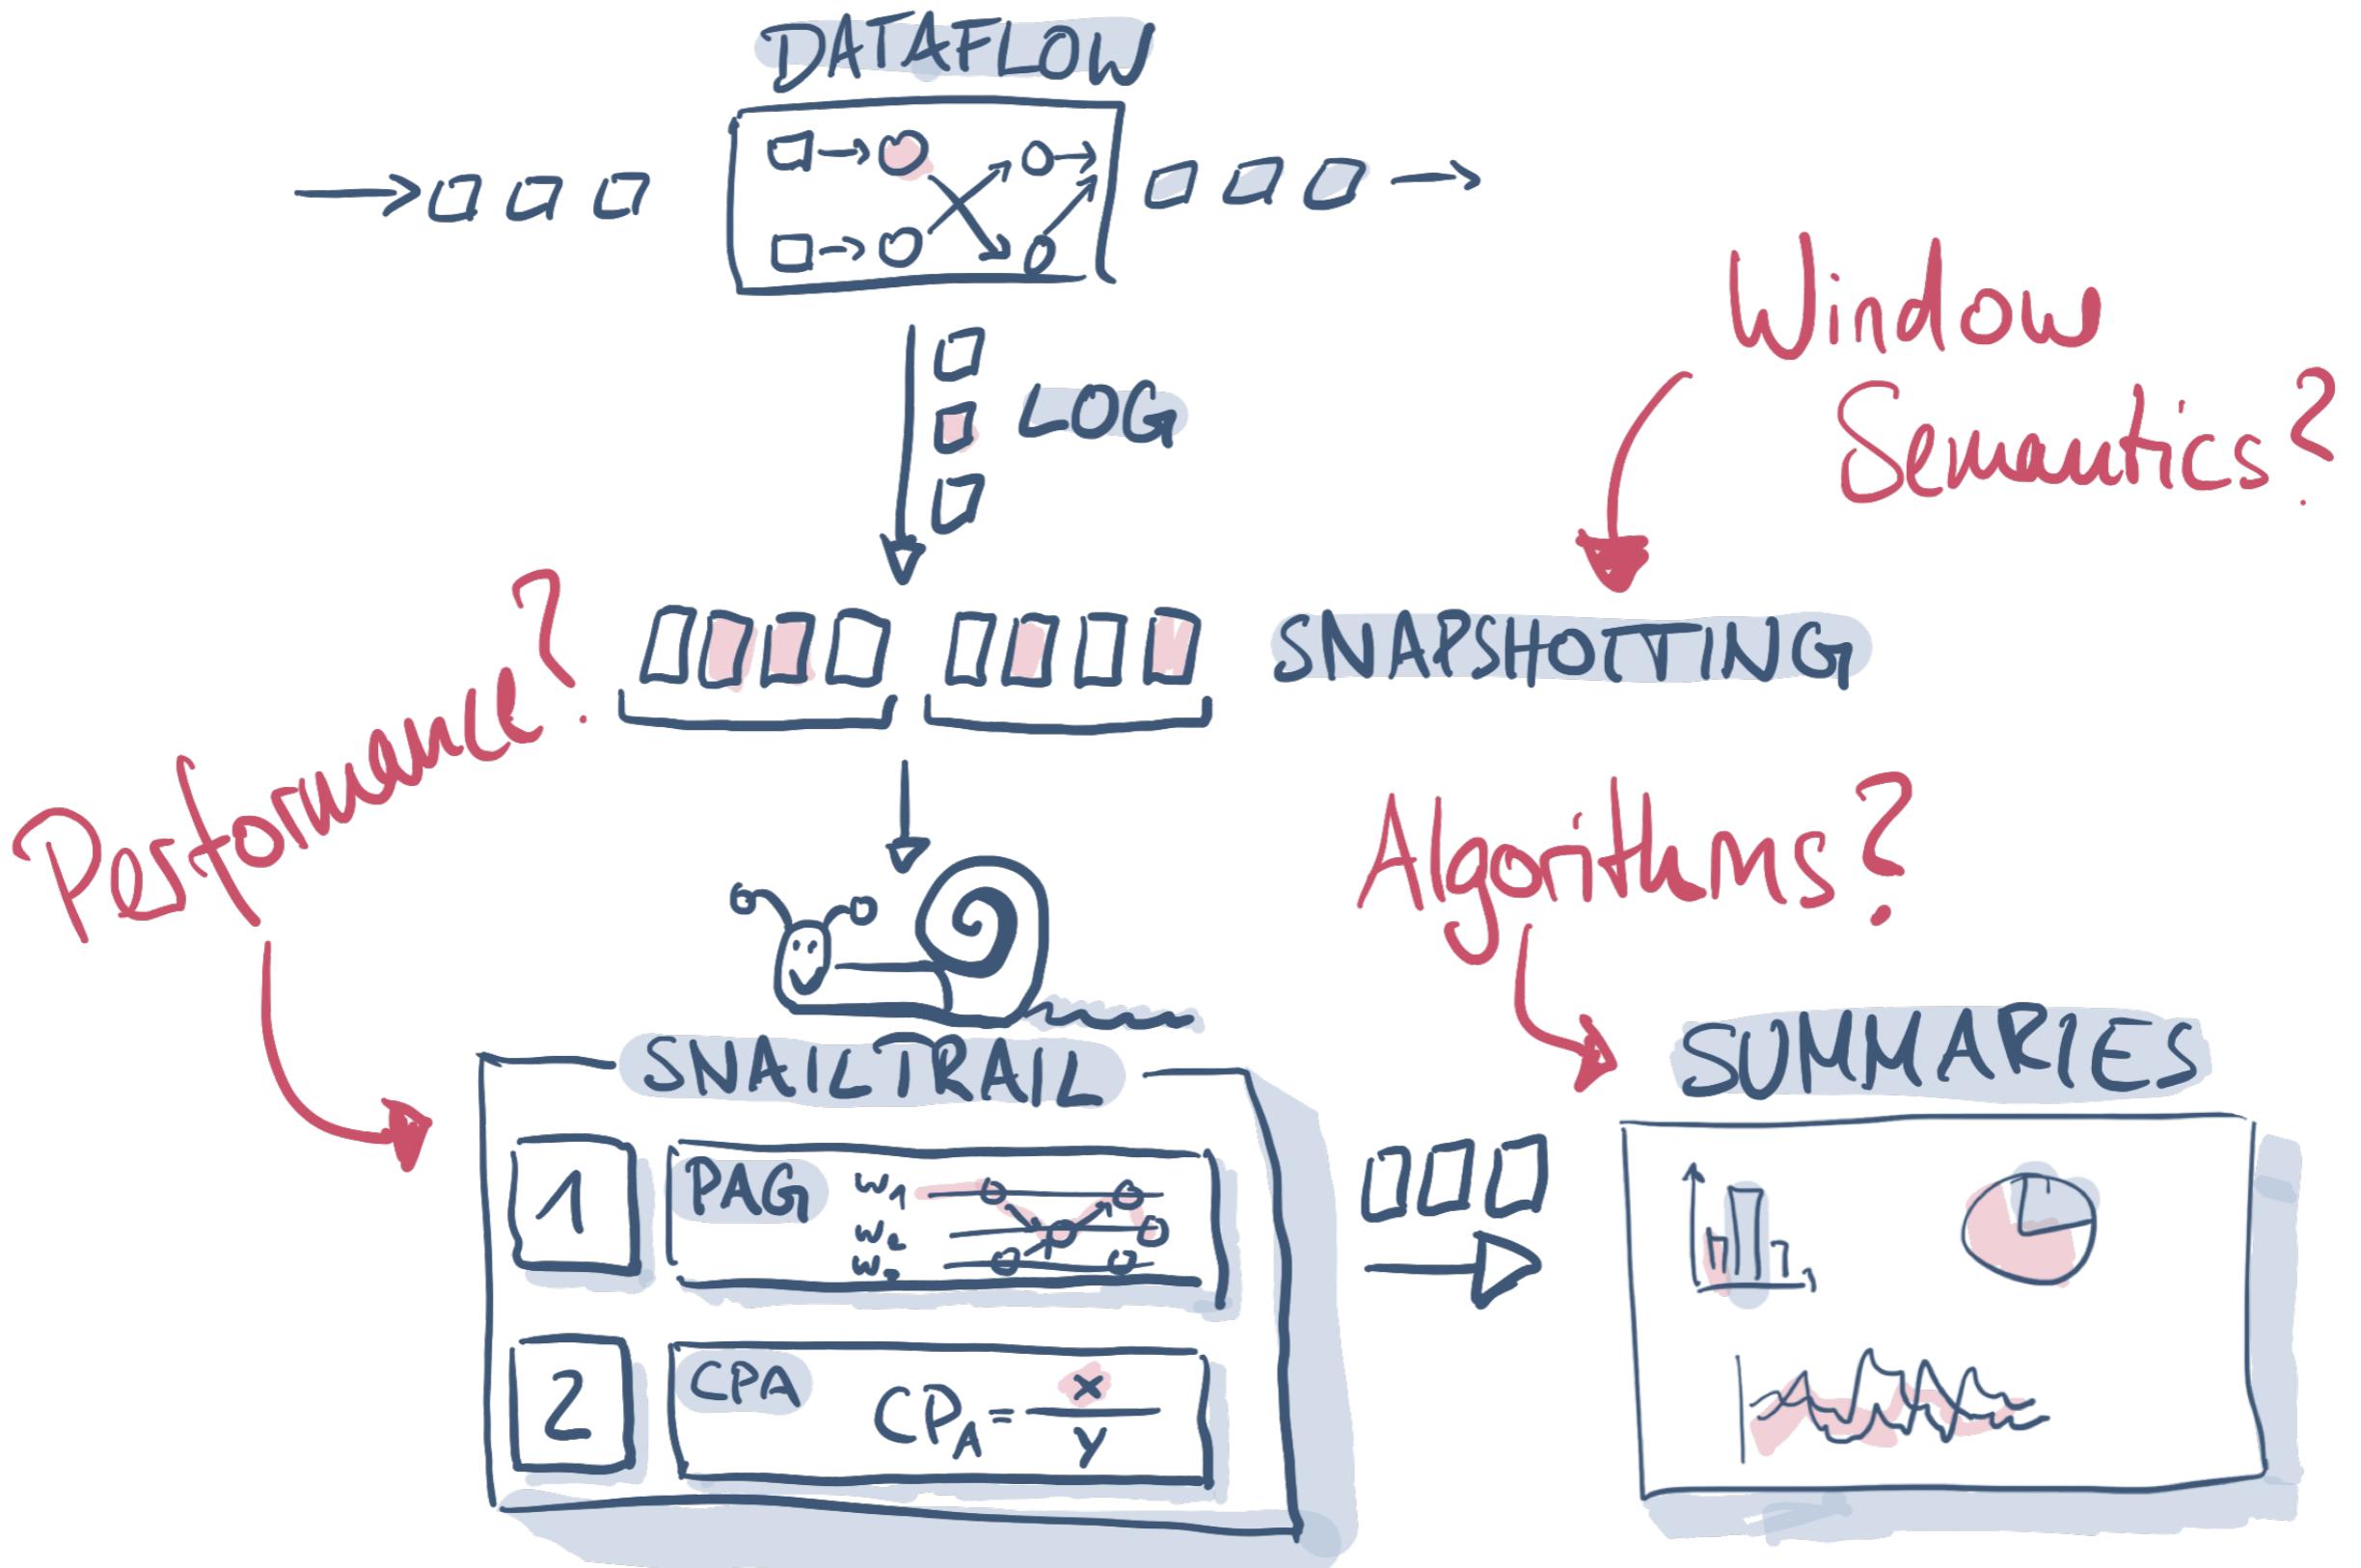


CP-metric: online critical path analysis



SnailTrail: online CP-based summaries

FUTURE WORK (ALSO, MY THESIS)



SnailTrail: Generalizing Critical Paths for Online Analysis of Distributed Dataflows*

strymon.systems.ethz.ch

*Moritz Hoffmann, Andrea Lattuada, John Liagouris, Vasiliki Kalavri,
Desislava Dimitrova, Sebastian Wicki, Zaheer Chothia, Timothy Roscoe*

Systems Group, Department of Computer Science, ETH Zürich

firstname.lastname@inf.ethz.ch

SnailTrail (stable):

github.com/strymon-system/snailtrail

SnailTrail (very much WIP):

github.com/li1/snailtrail

Strymon Ecosystem:

github.com/strymon-system

Timely, Differential:

github.com/timelydataflow

Malte Sandstede (samalt@ethz.ch)